

COMP 414/514:
Optimization – Algorithms, Complexity
and Approximations

Lecture 2

Overview

\min_x

s.t.

$$f(x)$$
$$x \in C$$

- Different objective classes
- Different strategies within each problem
- Different approaches based on computational capabilities
- Different approaches based on constraints

And, always having in mind applications in machine learning,
AI and signal processing

Overview

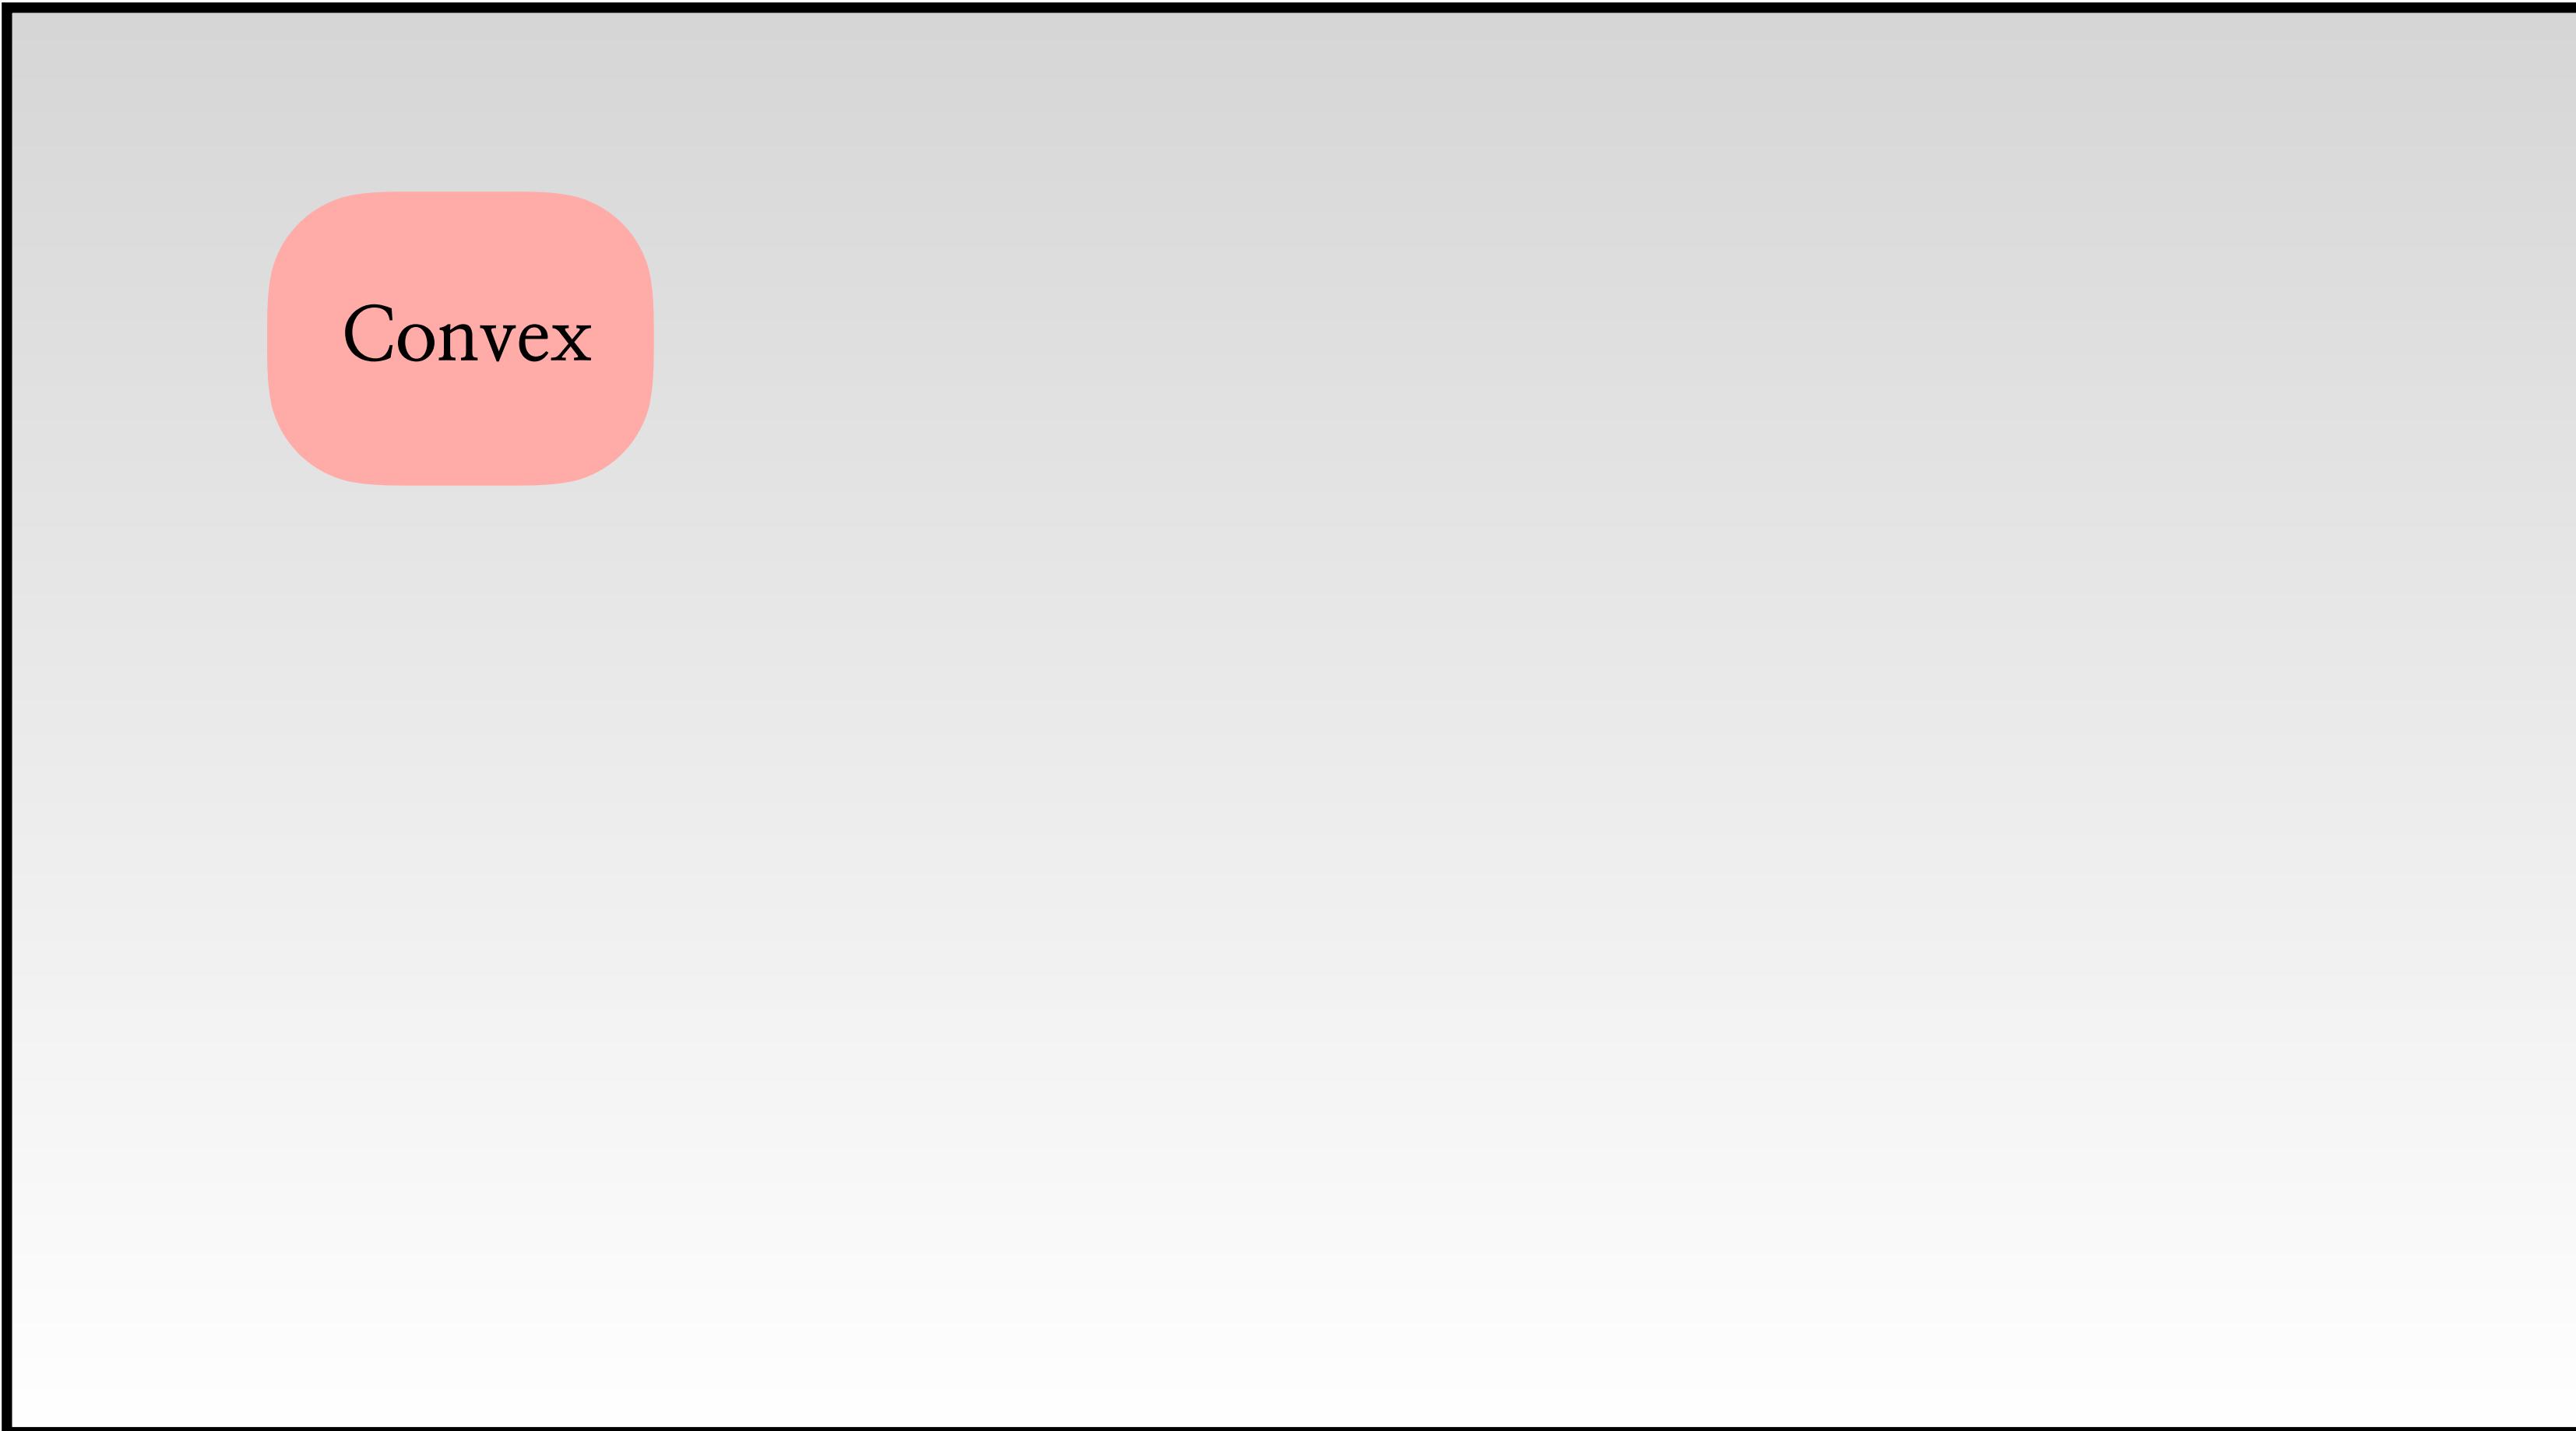
- In the last lecture, we:
 - Introduced some very basic ideas from linear algebra
- In this lecture, we will:
 - Discuss briefly **smooth continuous optimization**
 - Introduce the important class of **convex optimization**
 - Discuss about **convergence rates** and some **lower bounds** on such rates

Convex vs. non-convex optimization



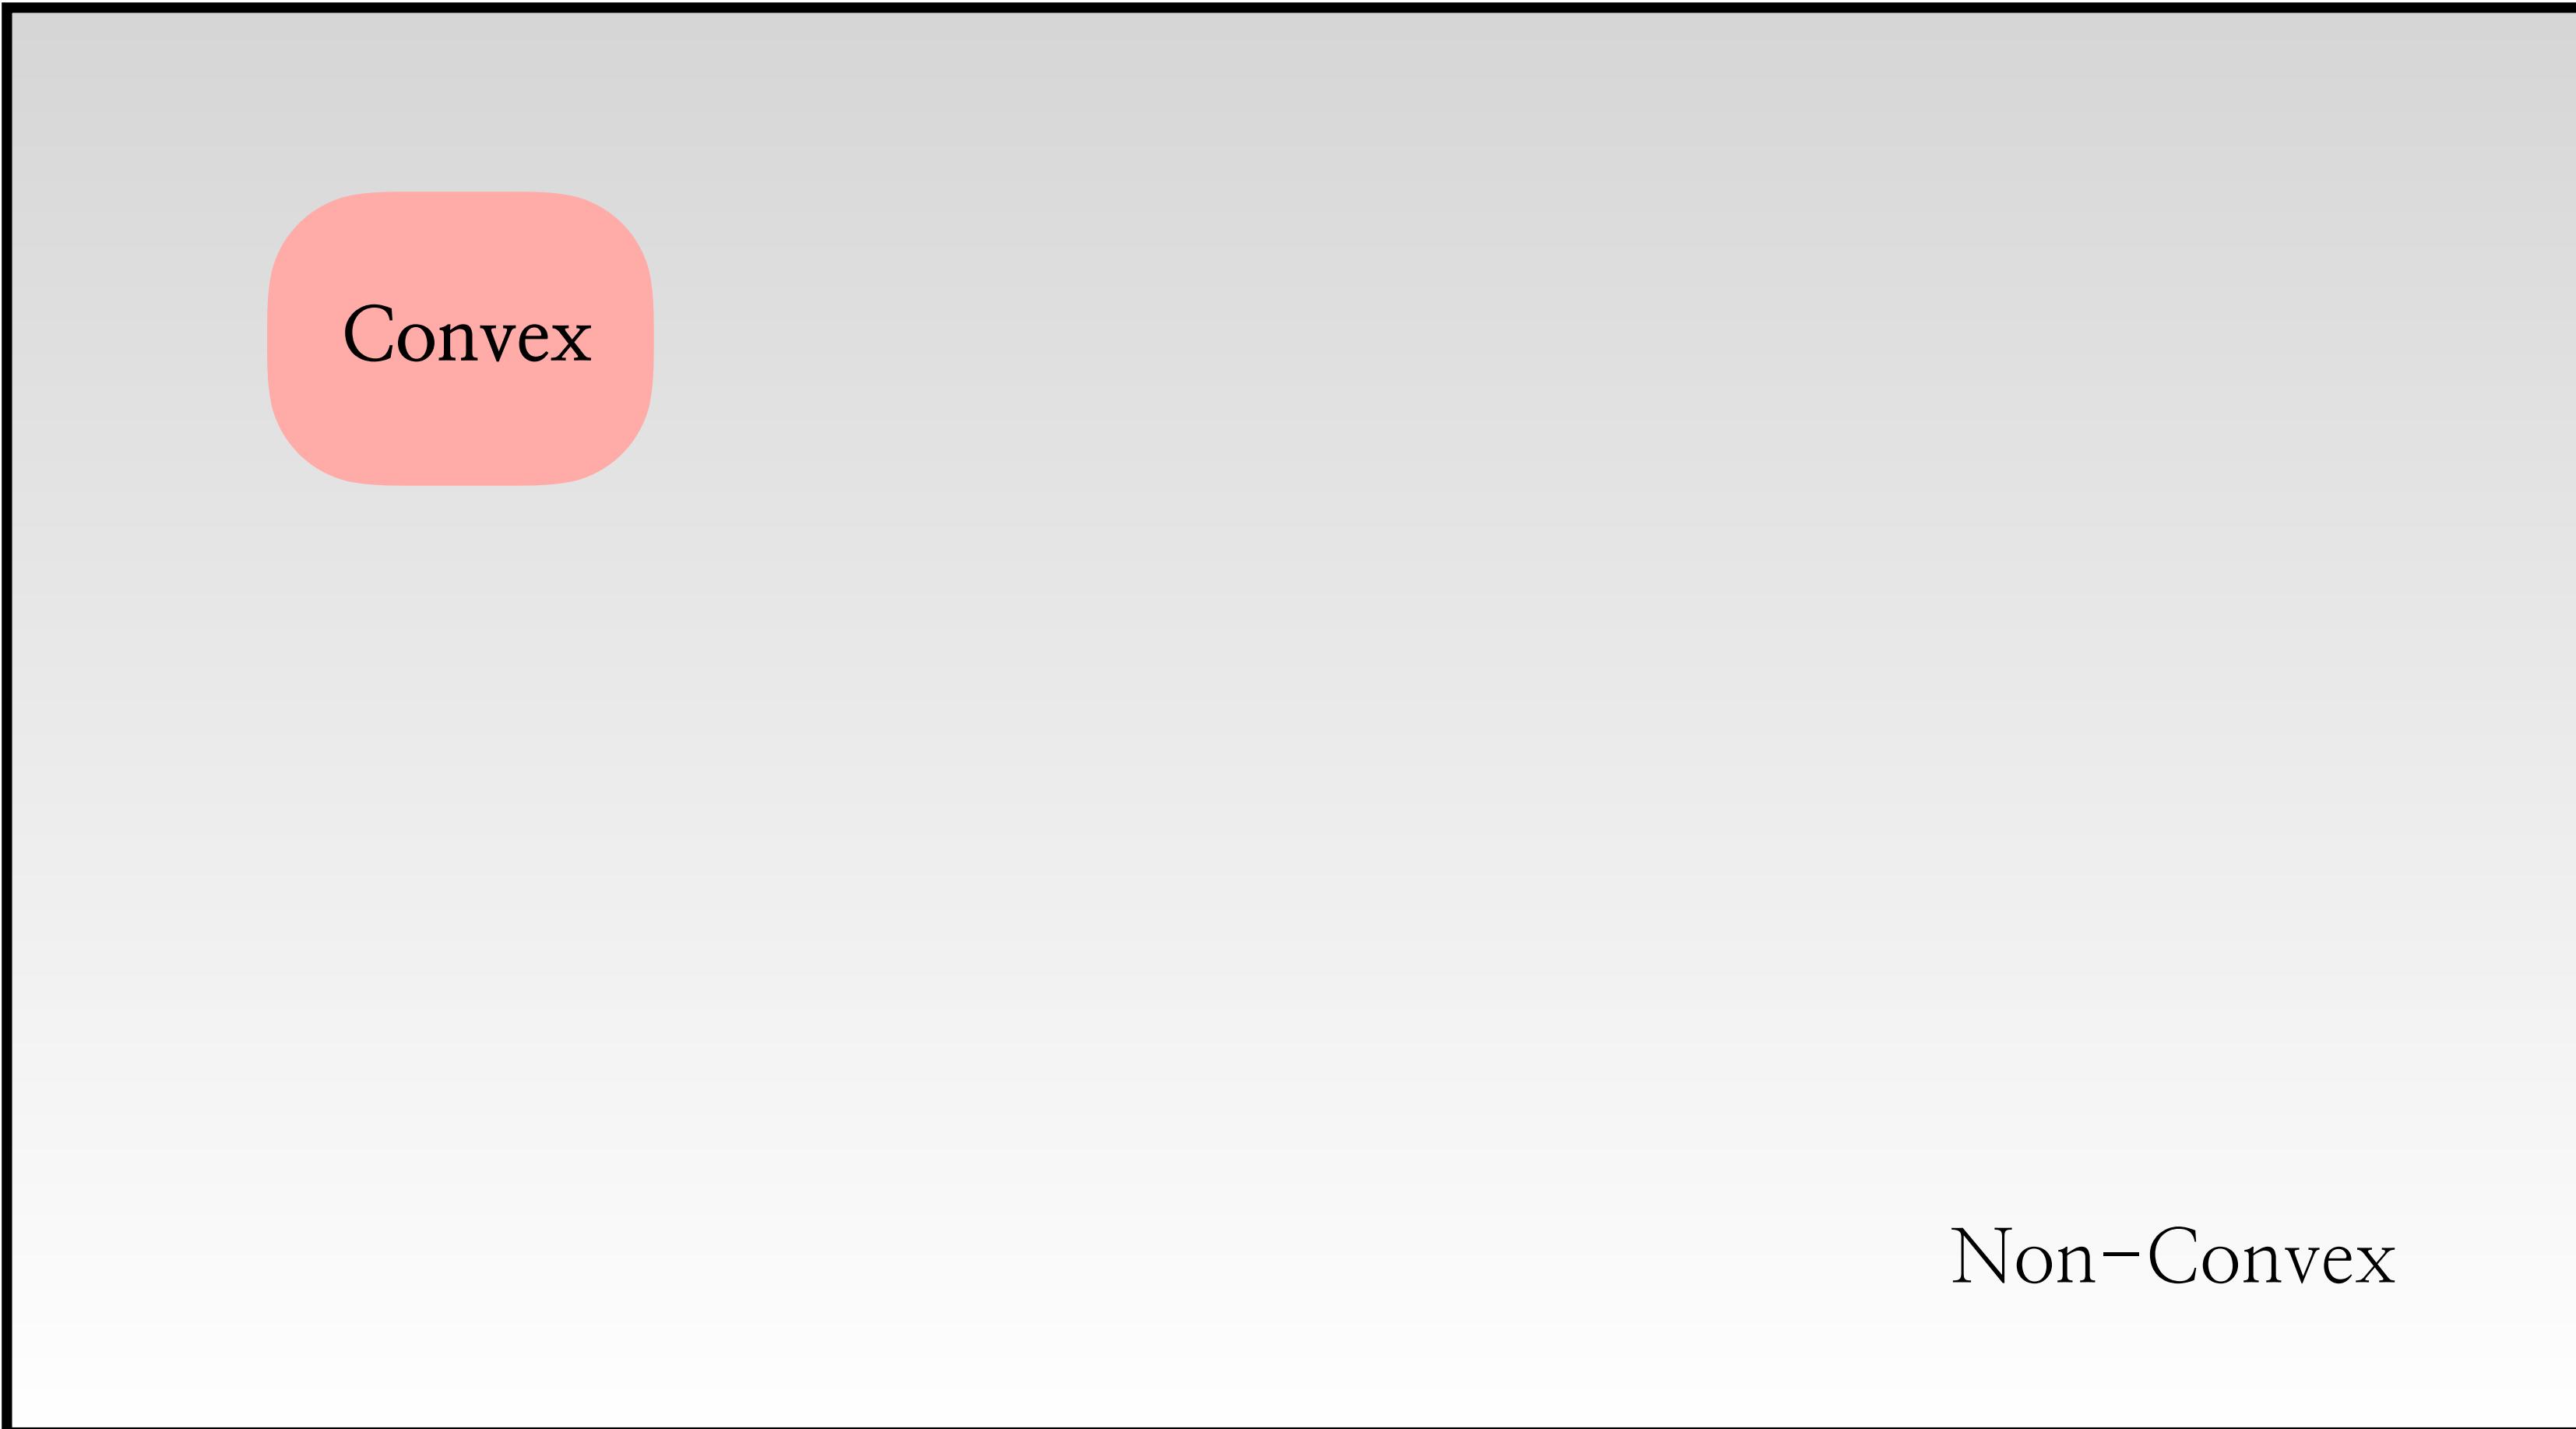
(Naive interpretation of) Space of optimization problems

Convex vs. non-convex optimization



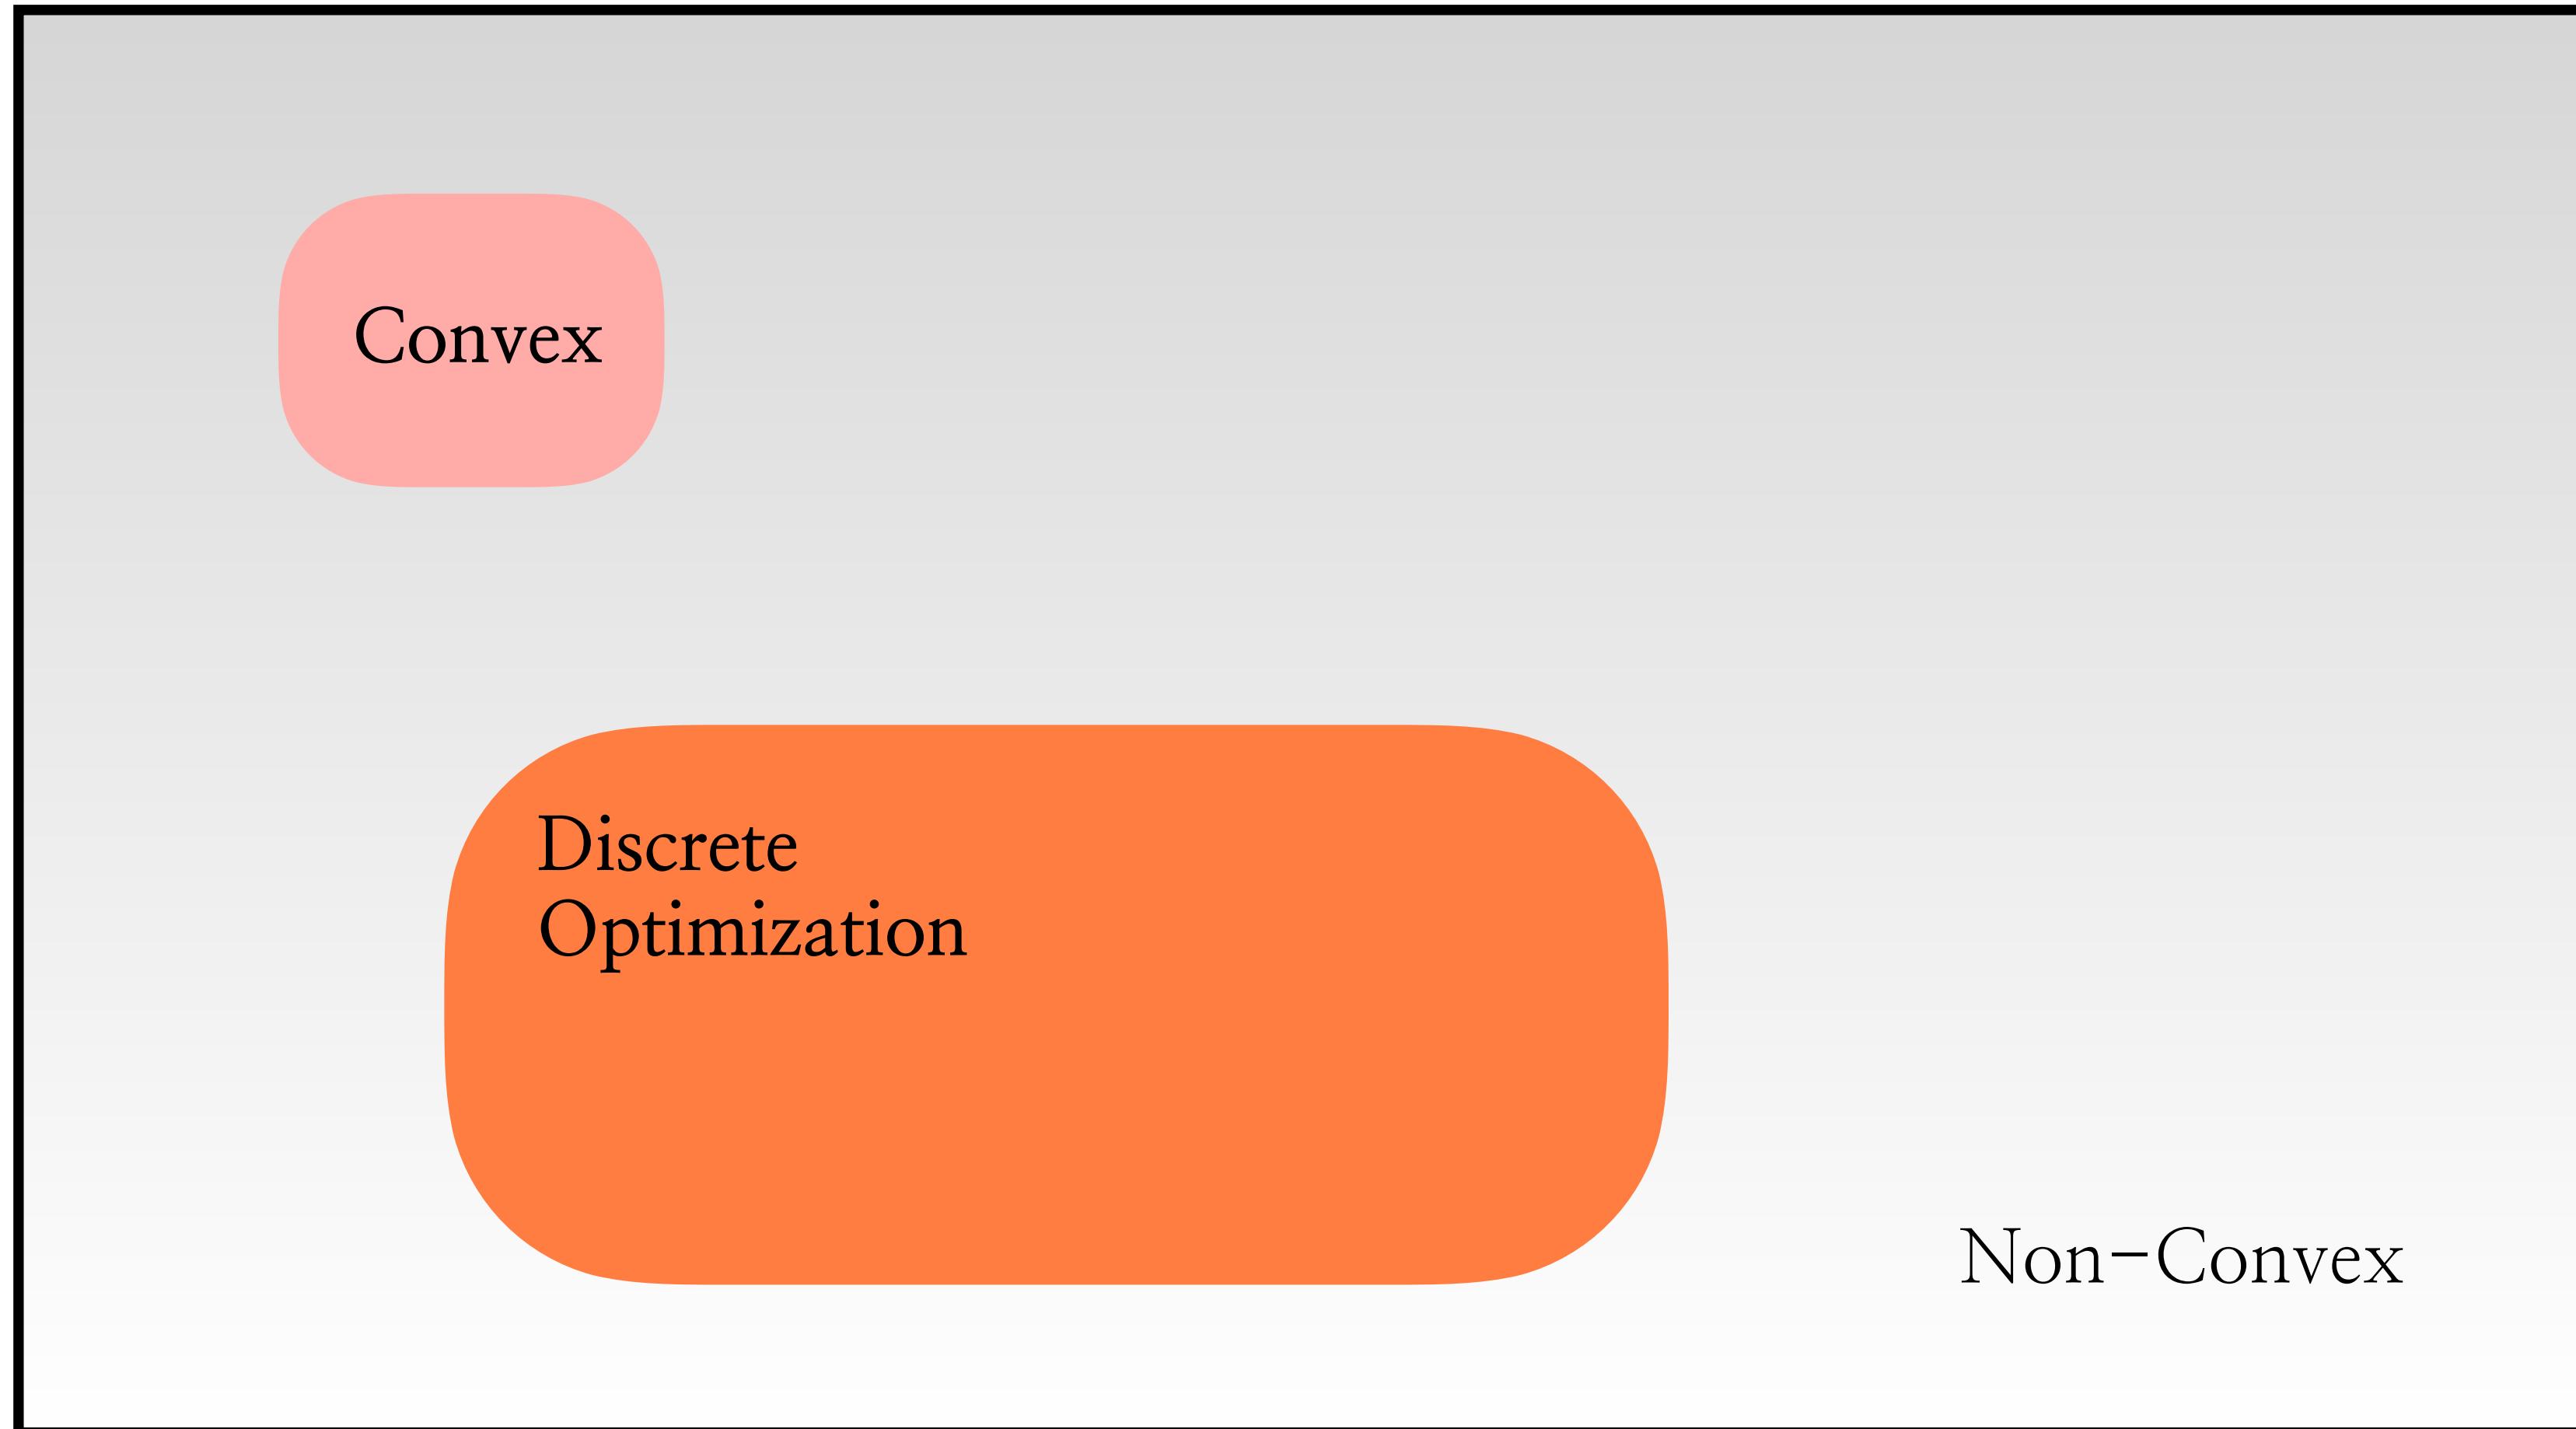
(Naive interpretation of) Space of optimization problems

Convex vs. non-convex optimization



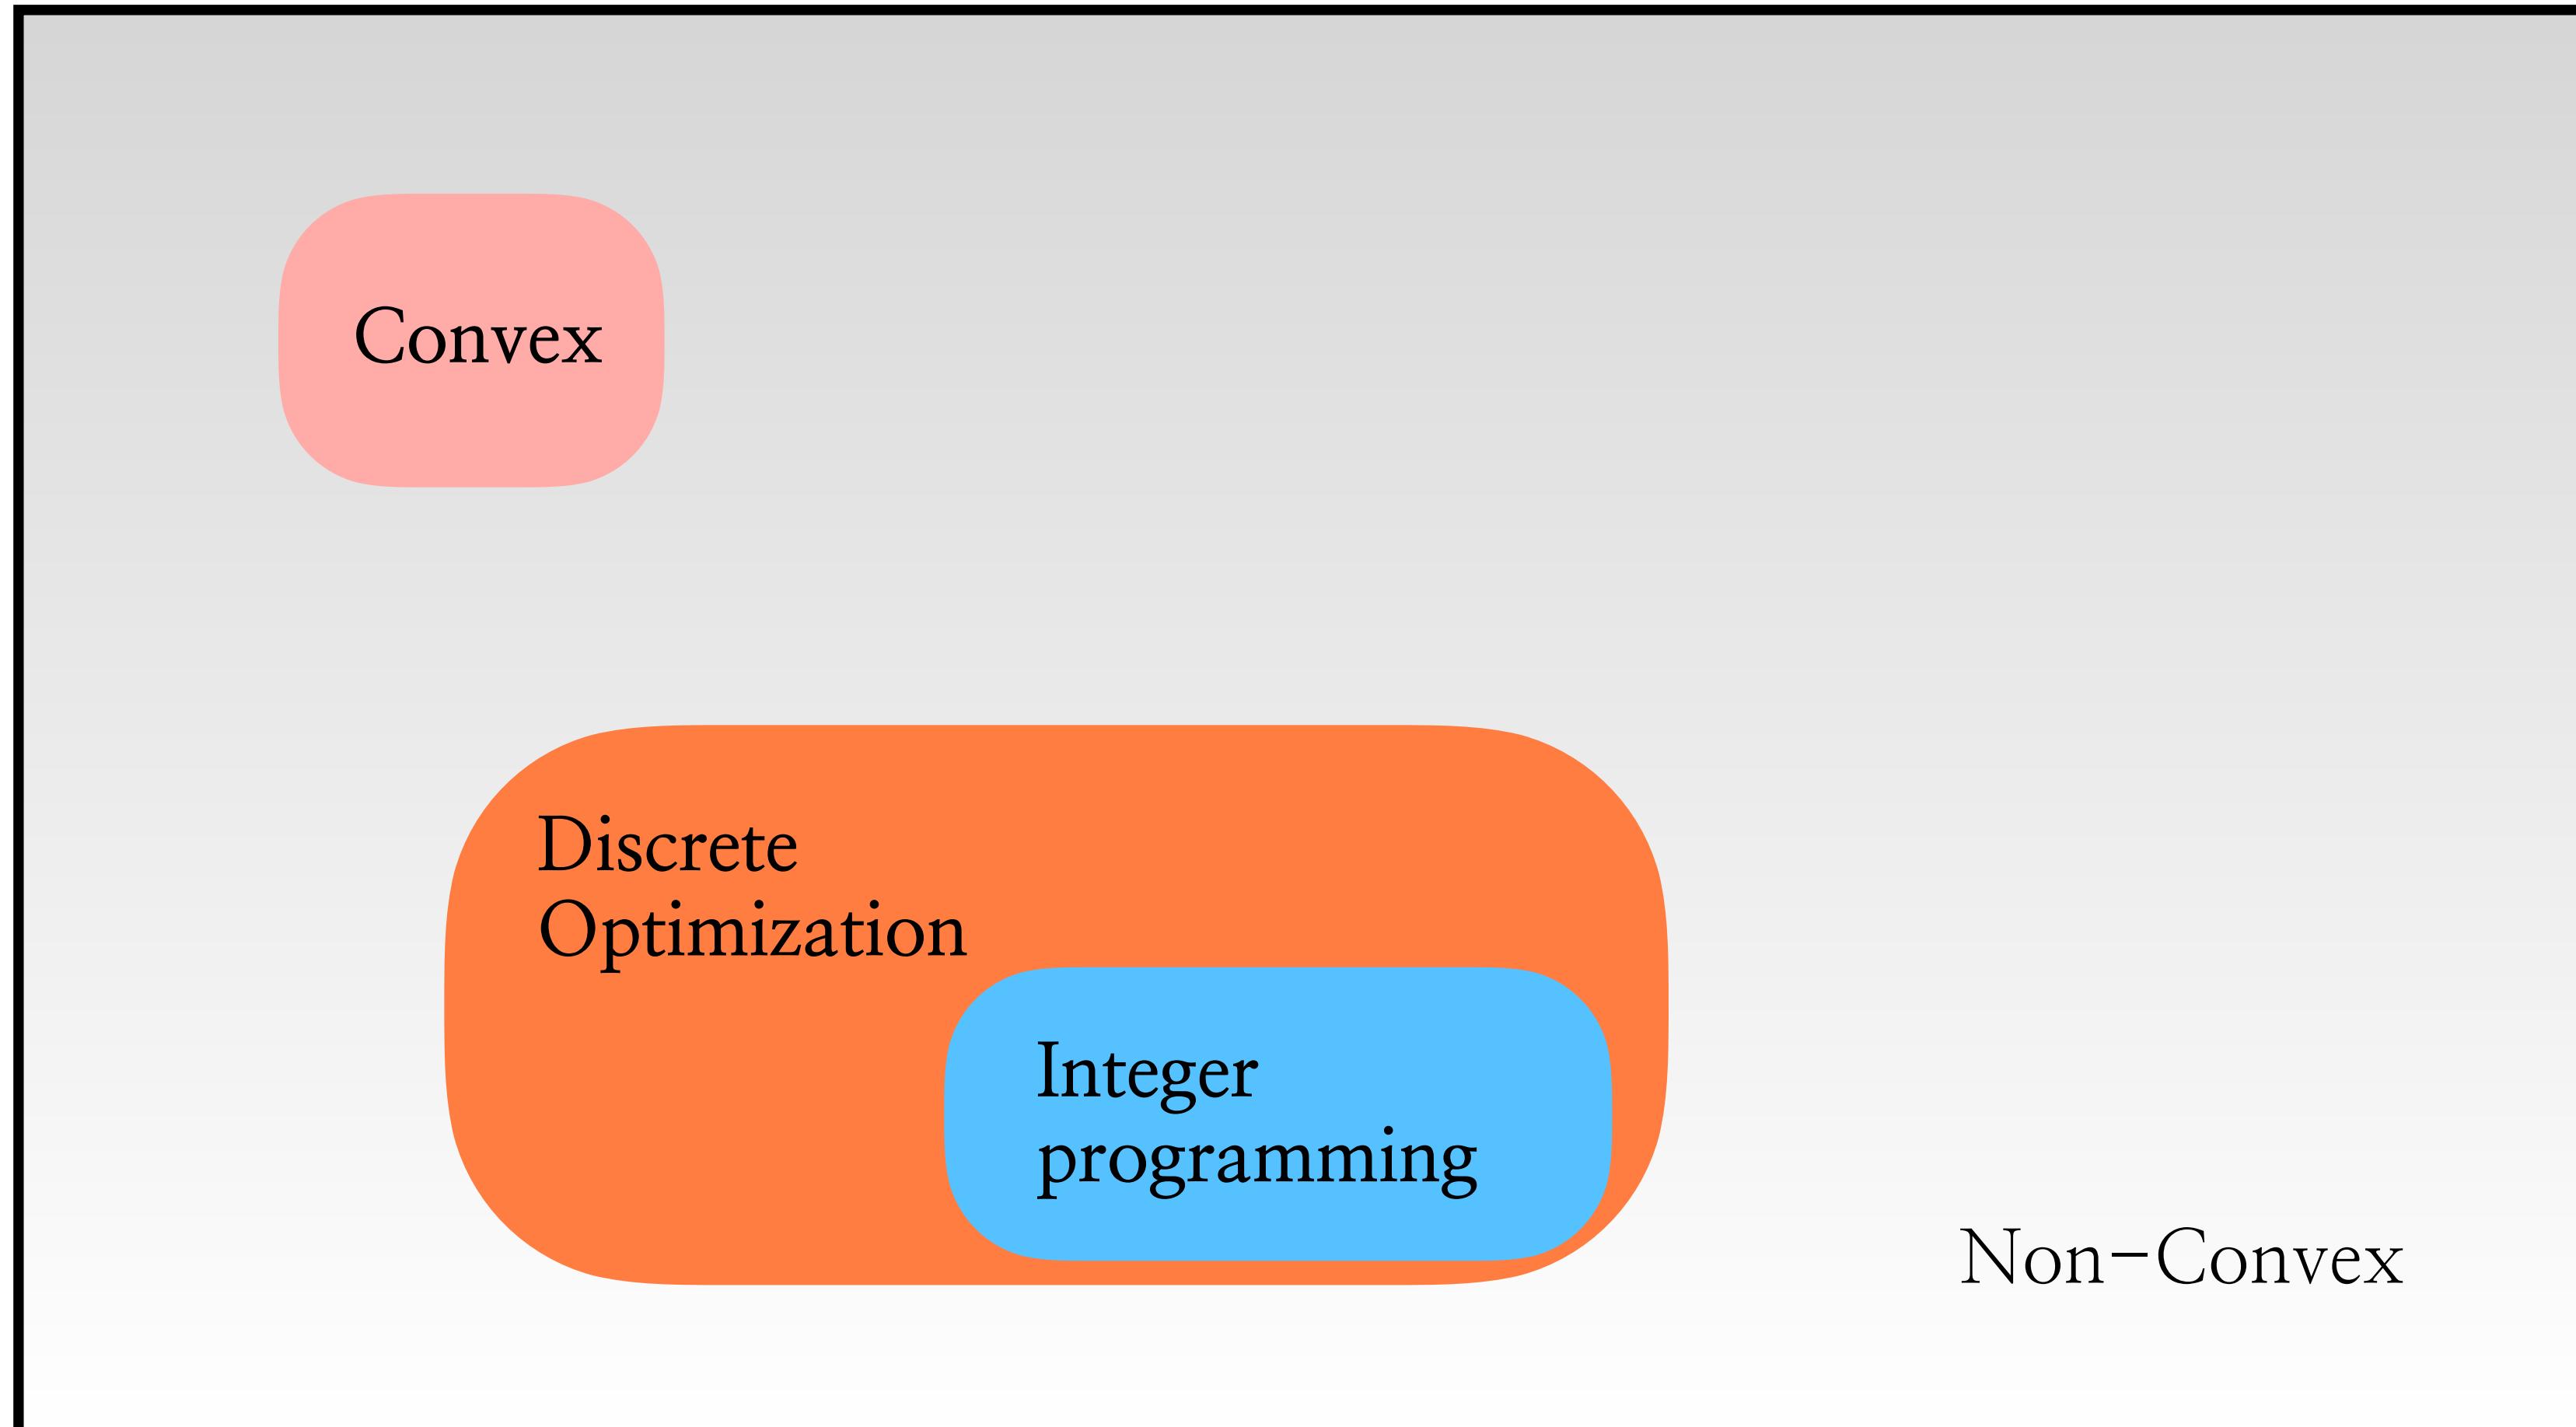
(Naive interpretation of) Space of optimization problems

Convex vs. non-convex optimization



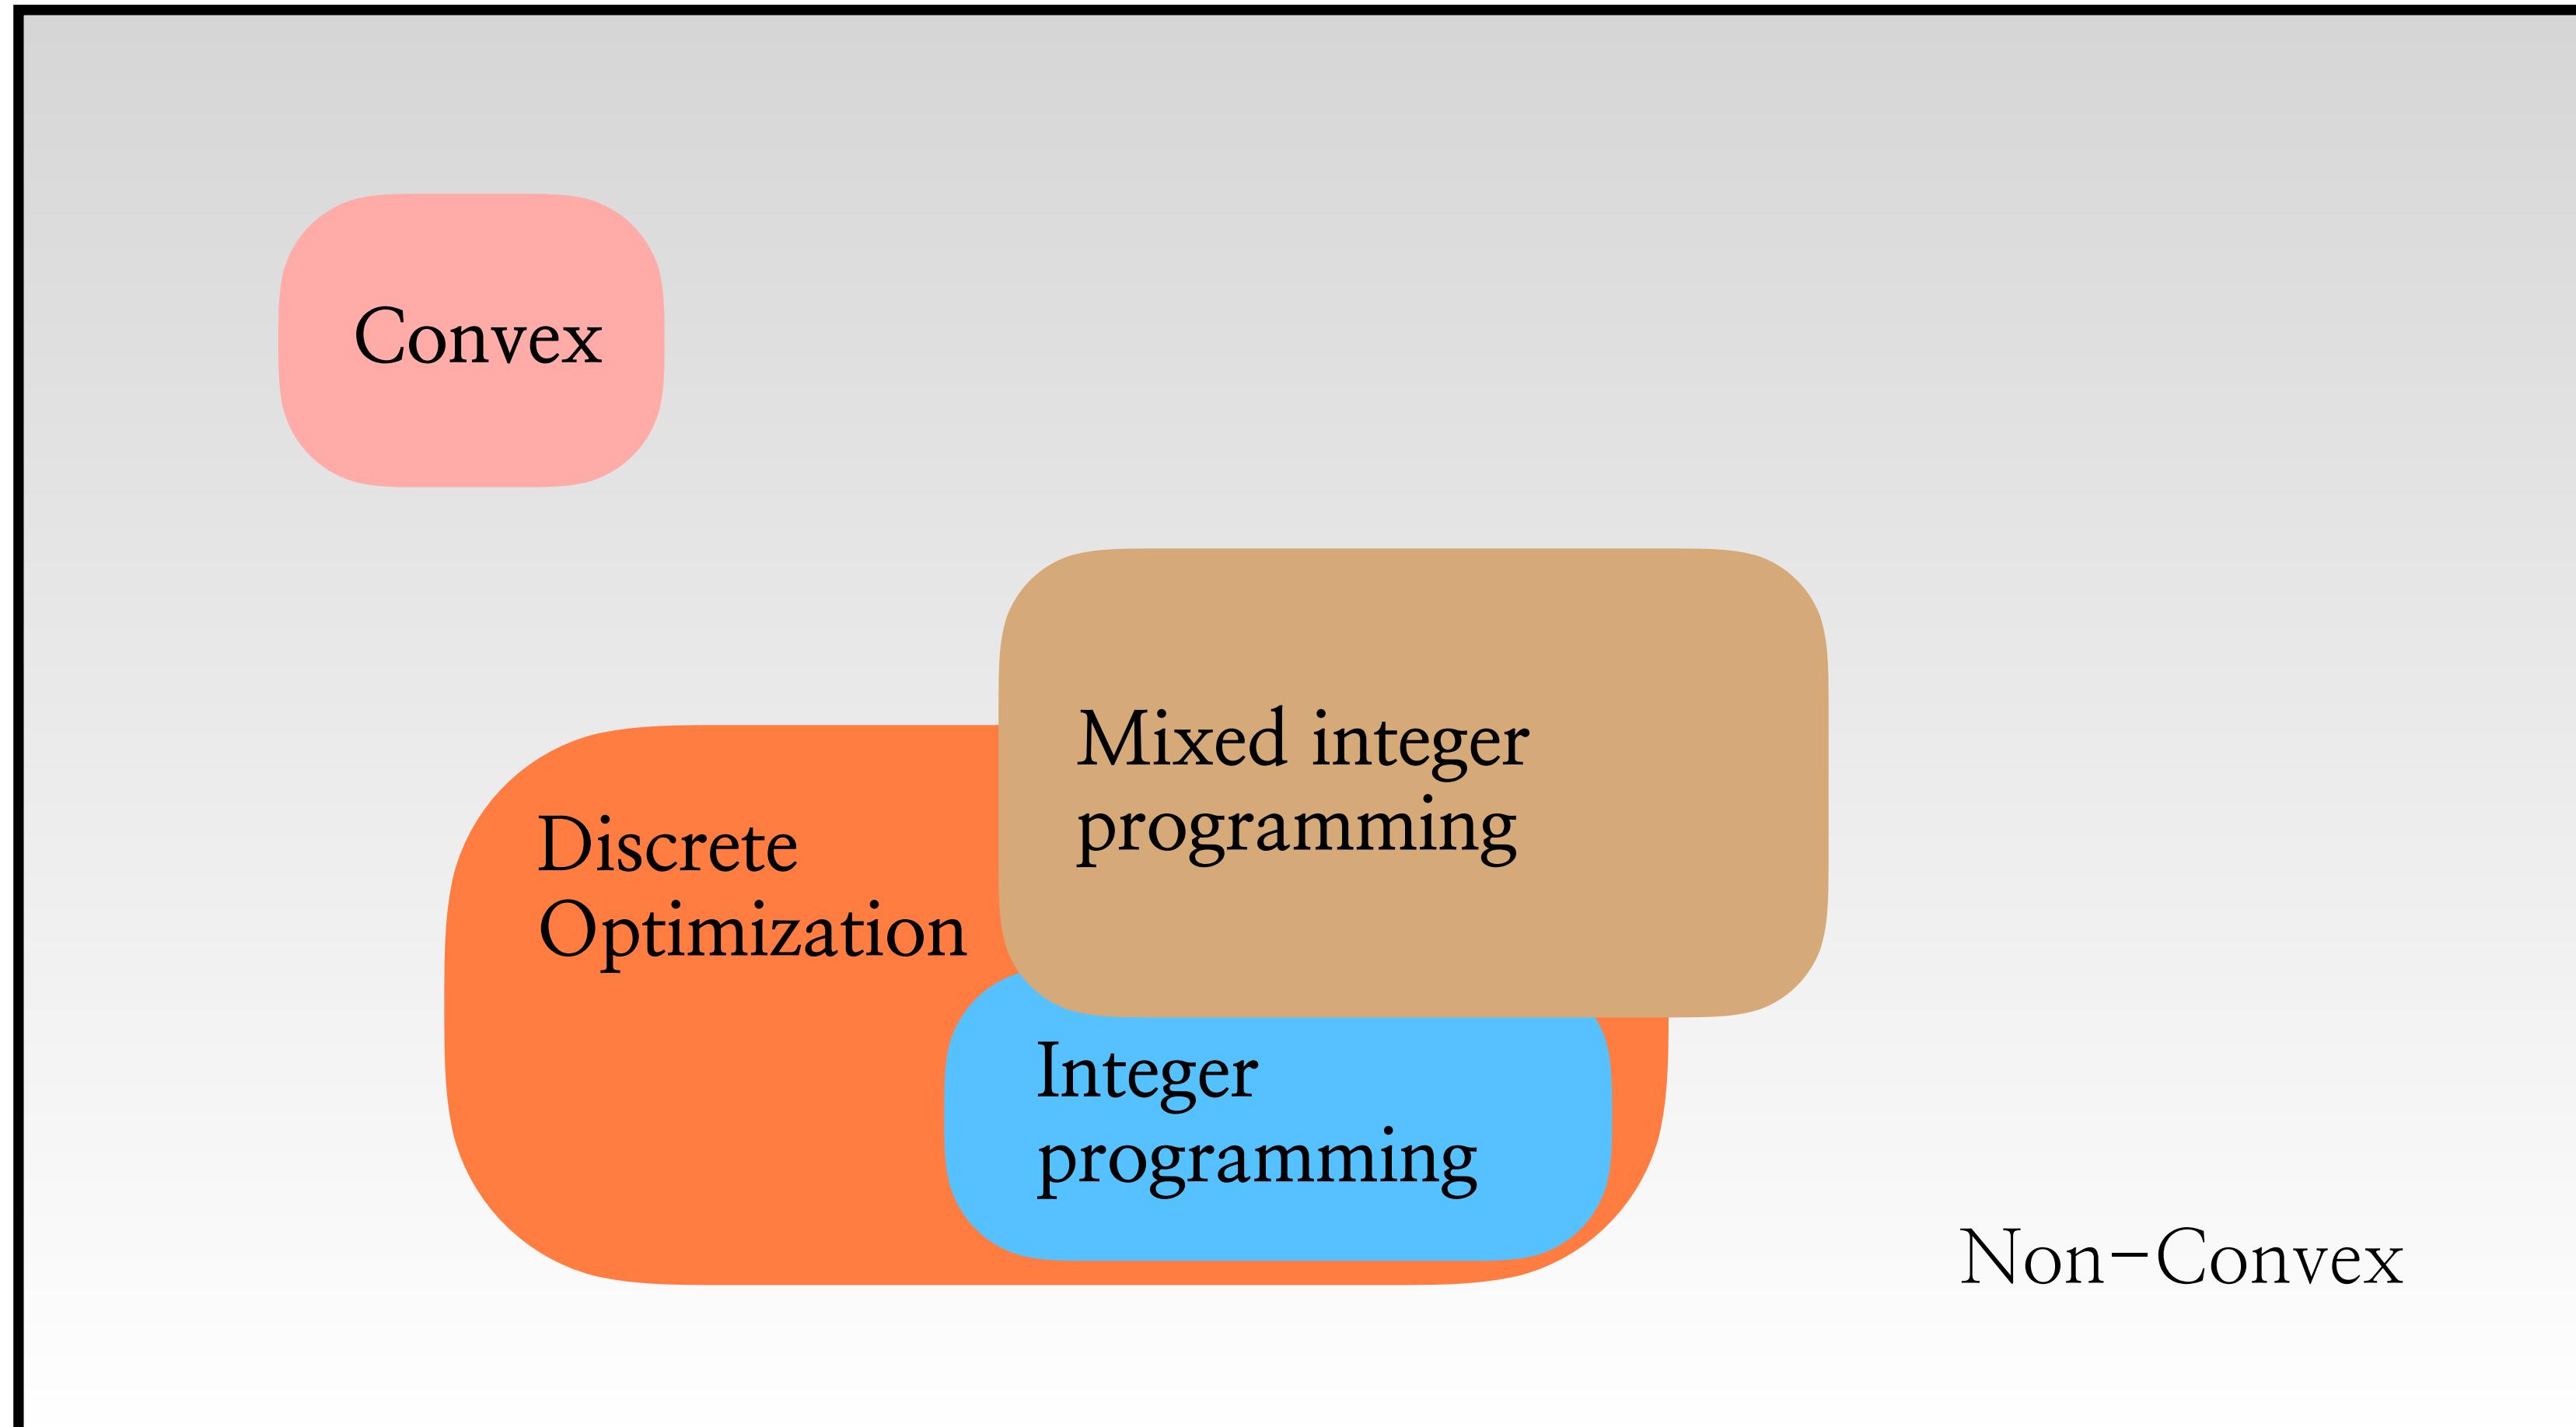
(Naive interpretation of) Space of optimization problems

Convex vs. non-convex optimization



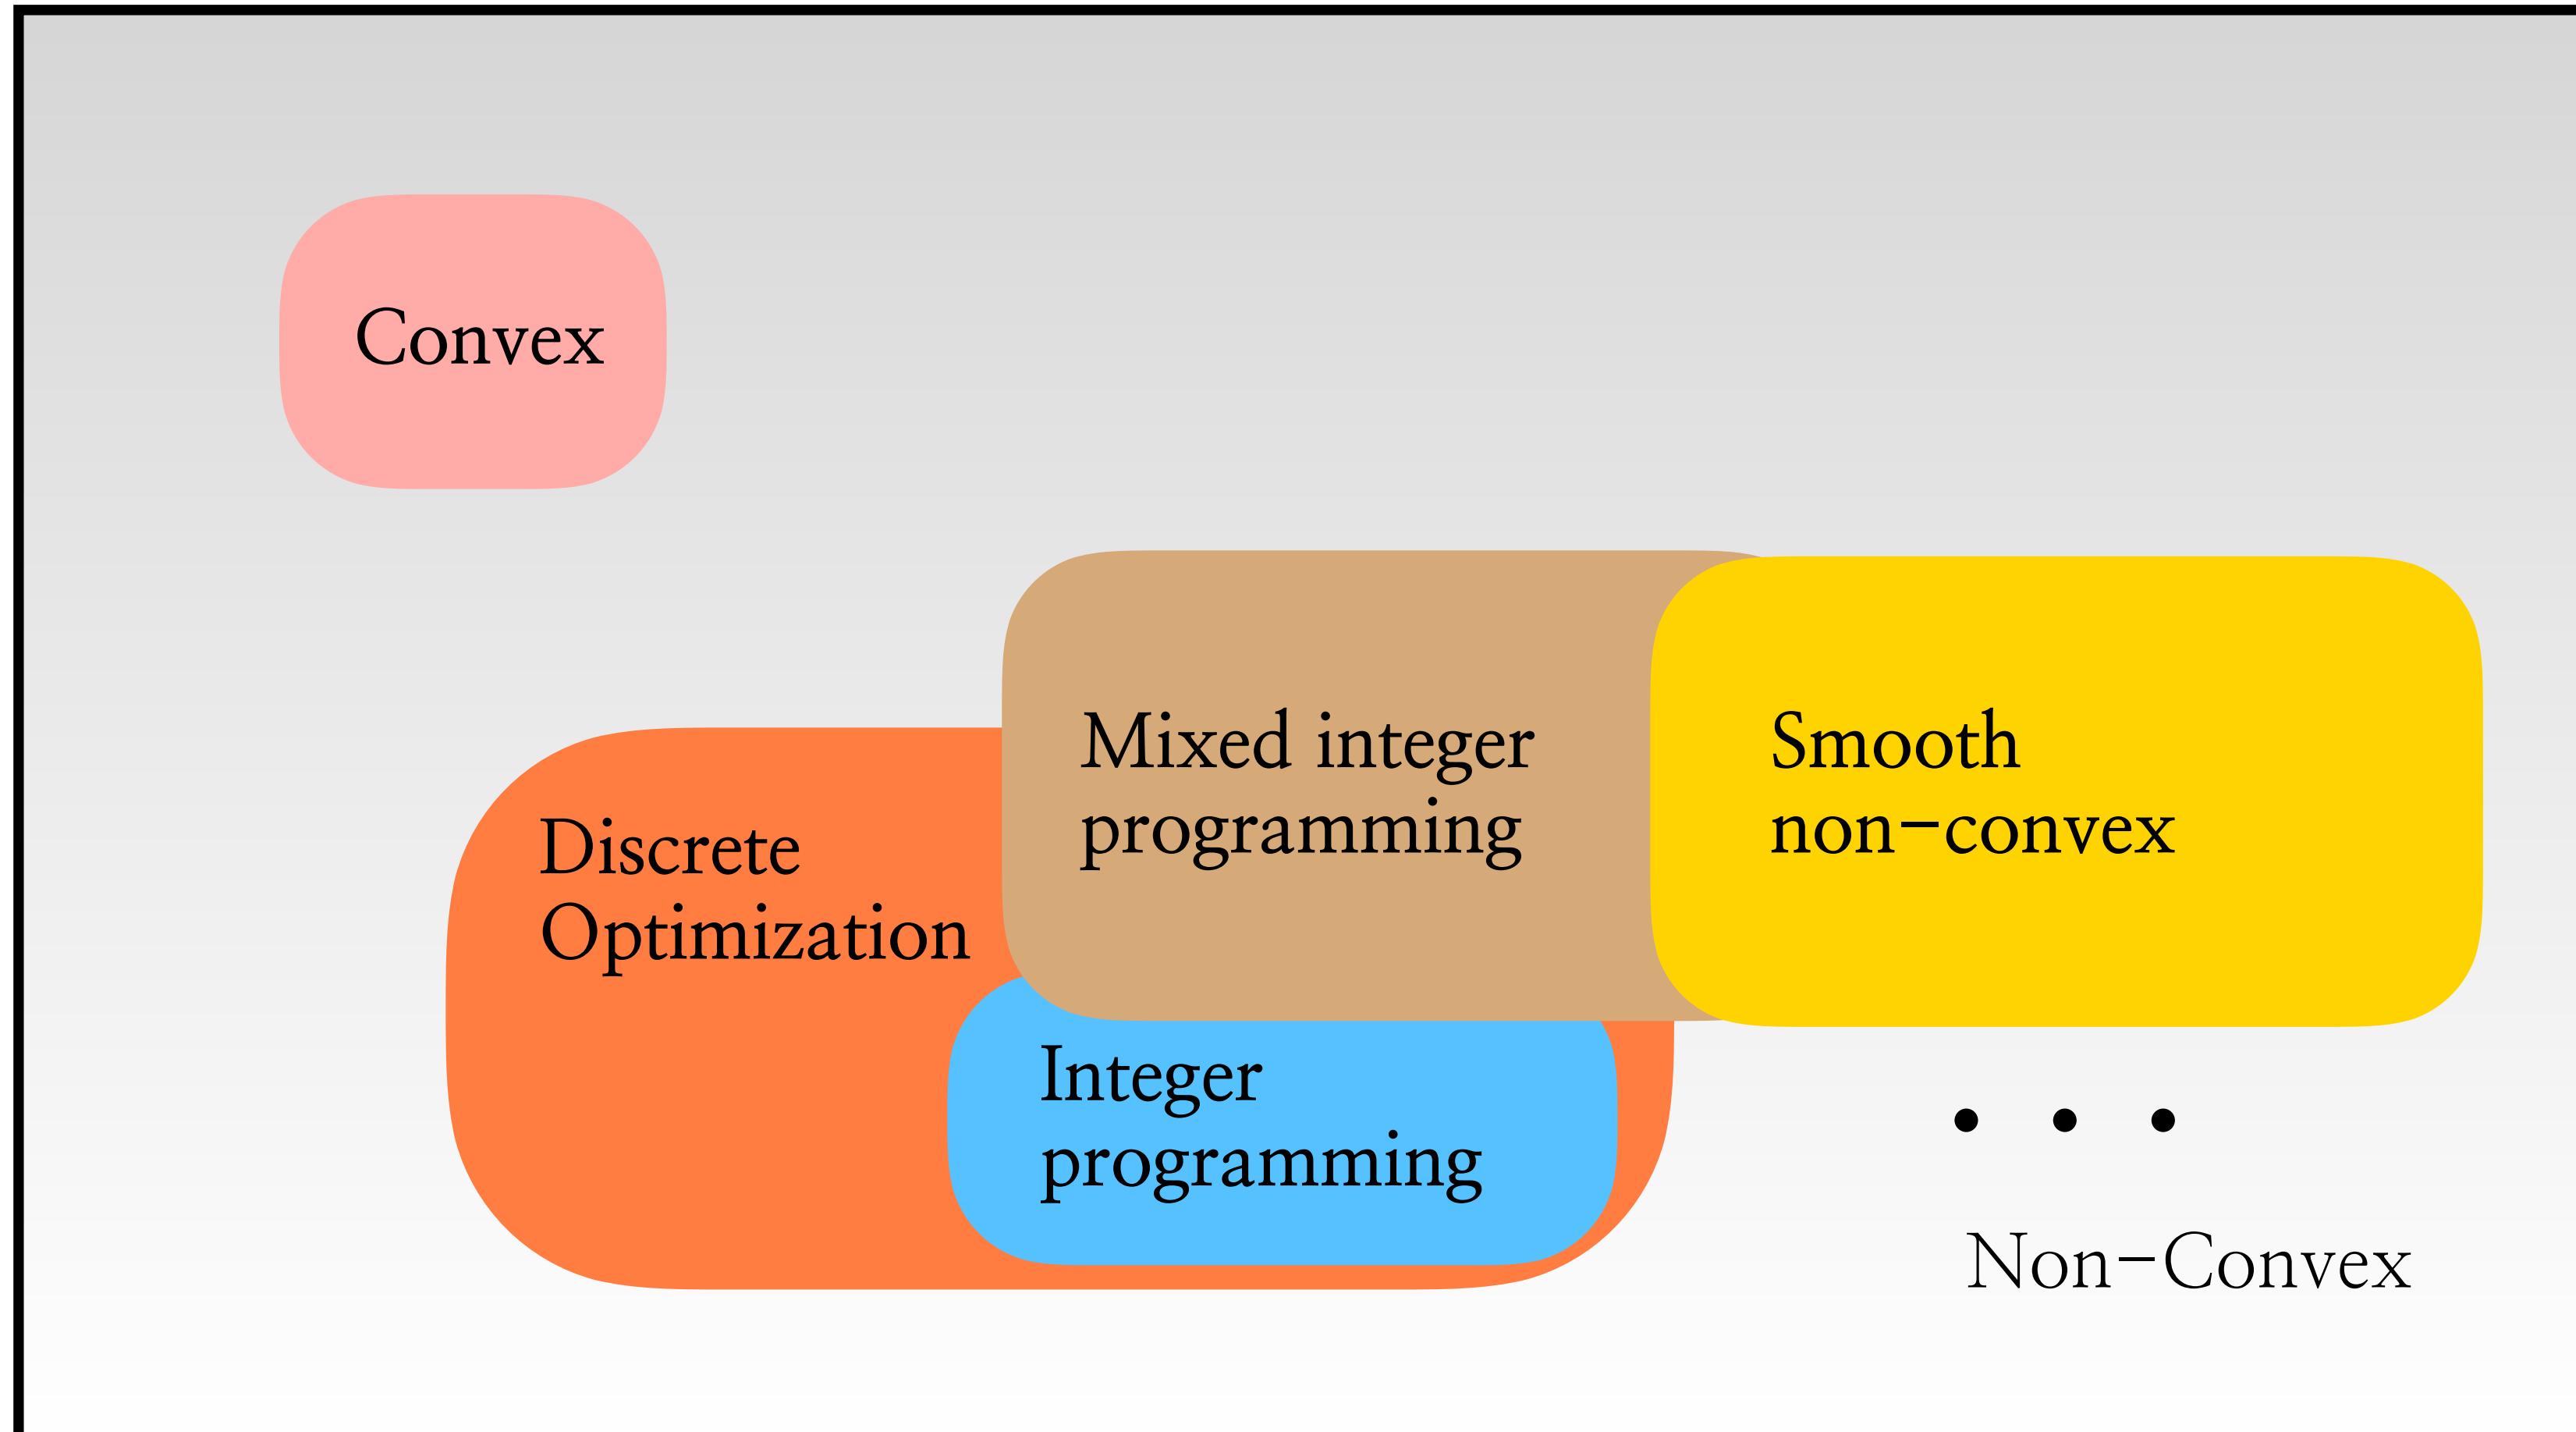
(Naive interpretation of) Space of optimization problems

Convex vs. non-convex optimization



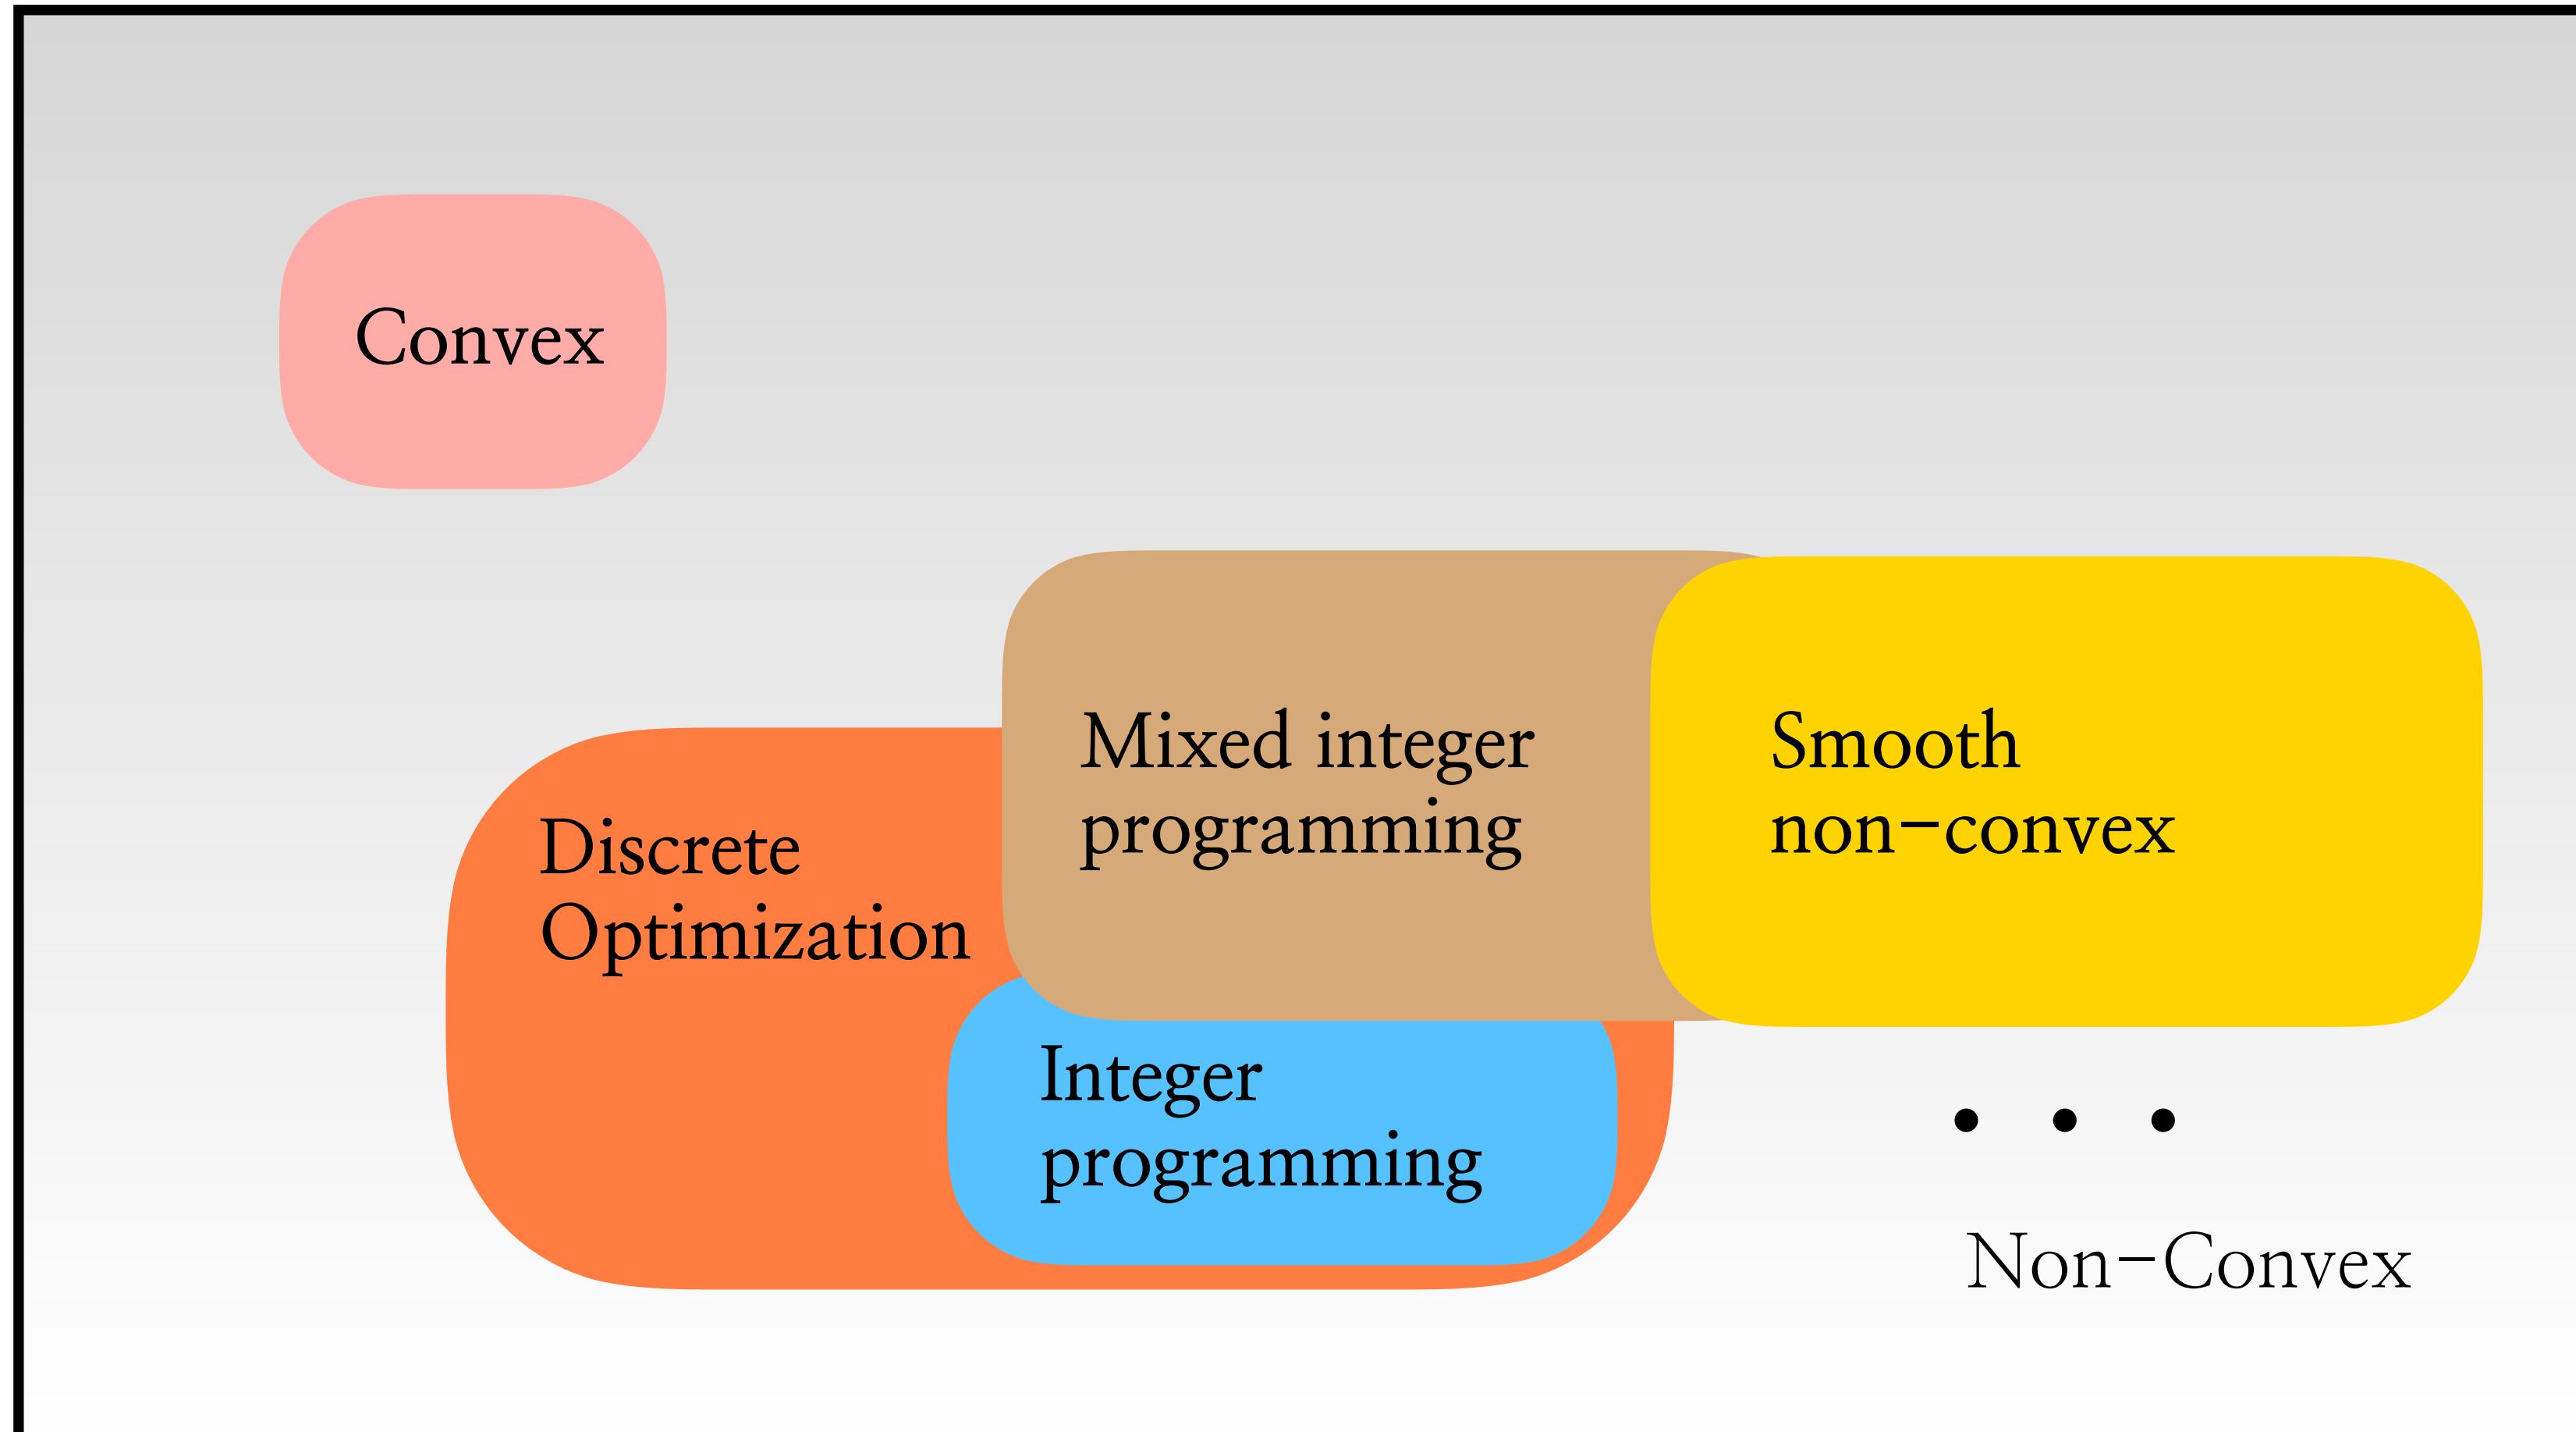
(Naive interpretation of) Space of optimization problems

Convex vs. non-convex optimization



(Naive interpretation of) Space of optimization problems

Convex vs. non-convex optimization



(Naive interpretation of) Space of optimization problems

Derivatives and gradients

- Definition of a **derivative**

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad | \quad \frac{\partial f}{\partial x} = f'(x) = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x)}{\epsilon}$$

Derivatives and gradients

- Definition of a **derivative**

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad | \quad \frac{\partial f}{\partial x} = f'(x) = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x)}{\epsilon}$$

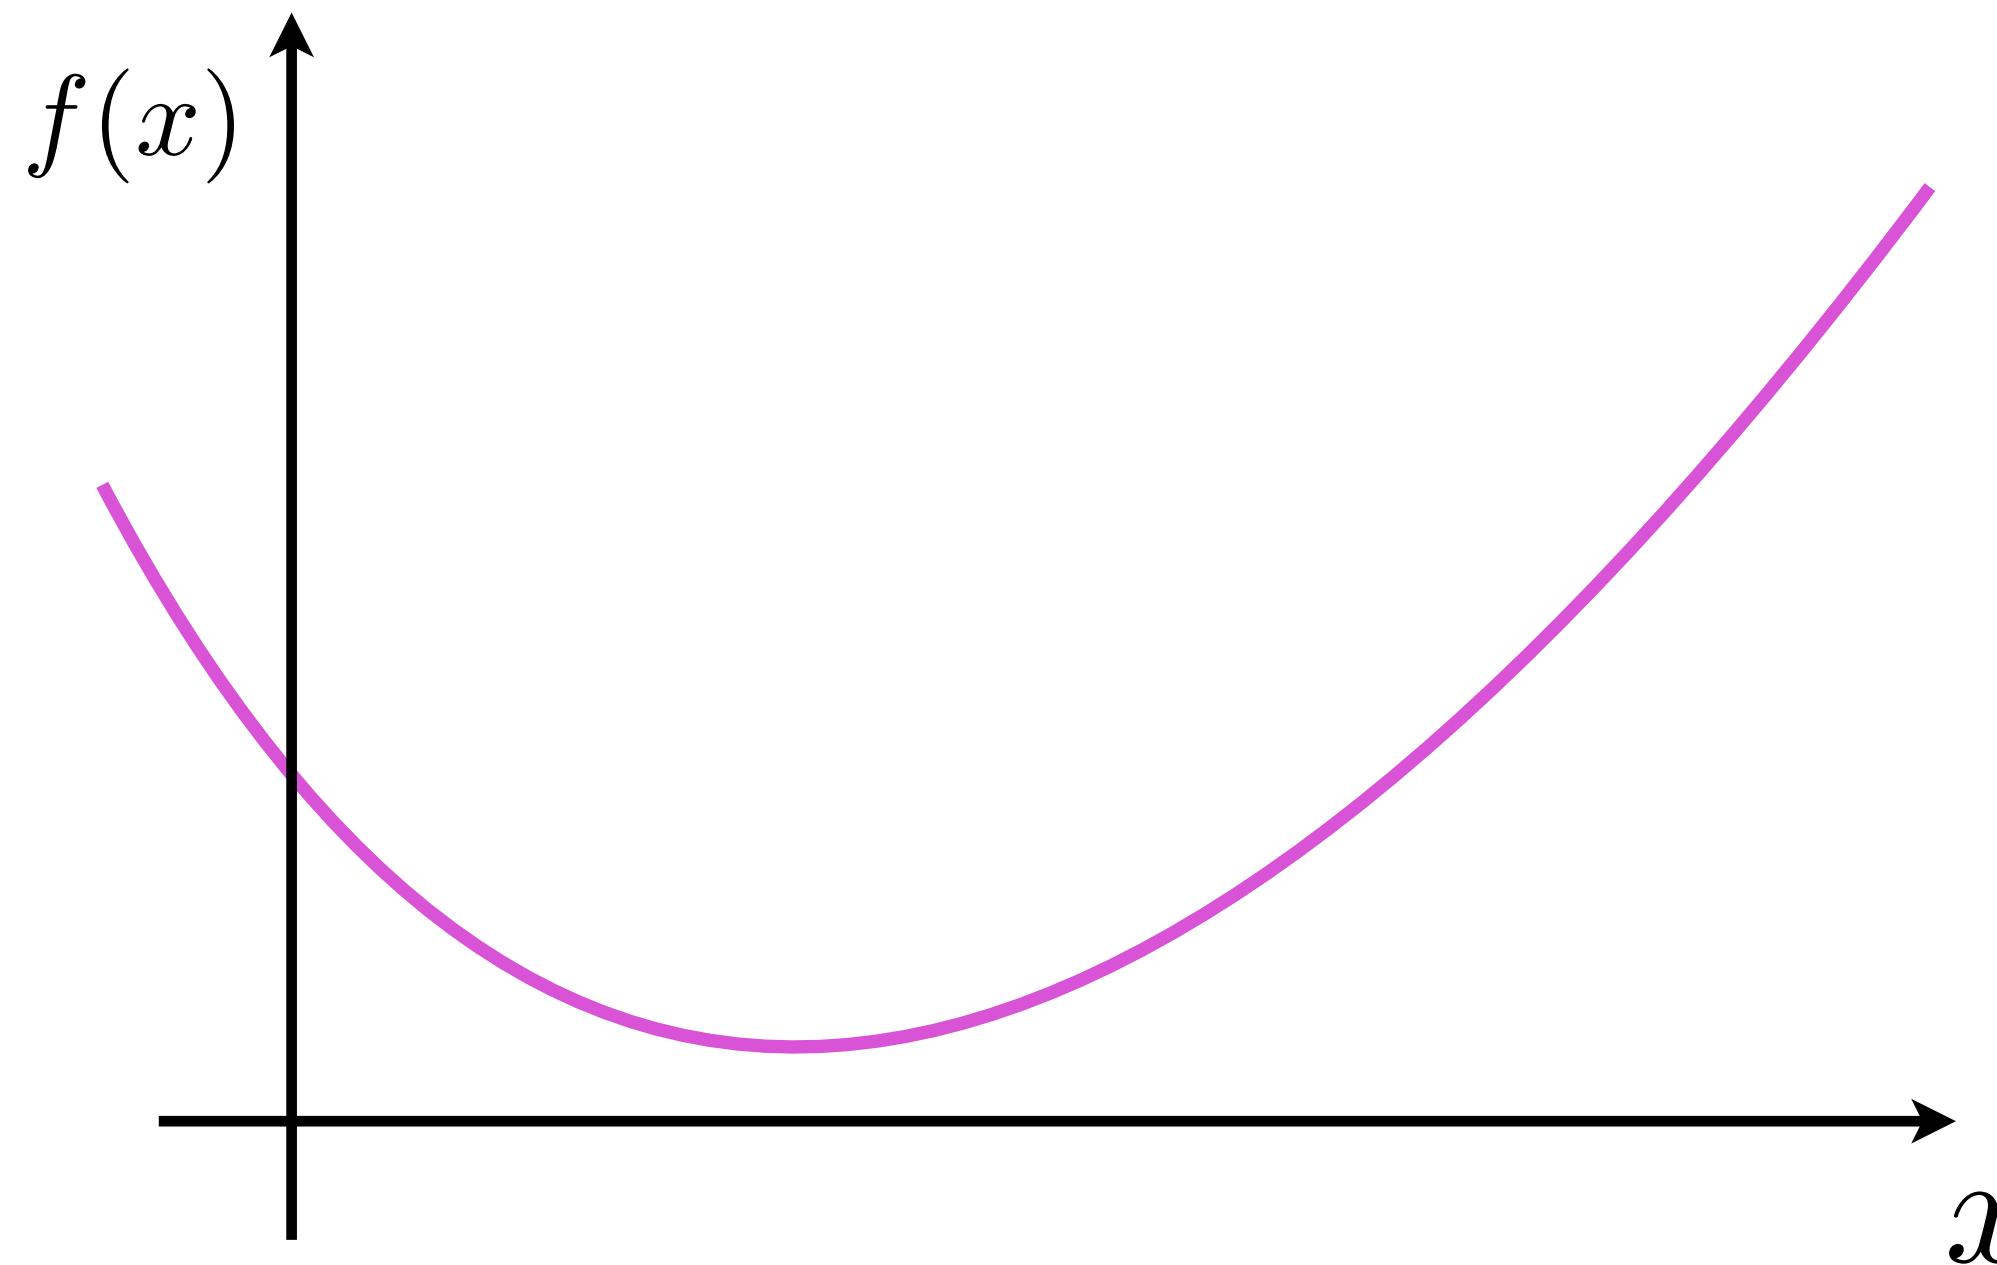
- Intuition: generate a sequence of points $\{f(x + \epsilon_i), \epsilon_i\}$, and compute the limit as $\epsilon_i \rightarrow 0$

Derivatives and gradients

- Definition of a **derivative**

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad | \quad \frac{\partial f}{\partial x} = f'(x) = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x)}{\epsilon}$$

- Intuition: generate a sequence of points $\{f(x + \epsilon_i), \epsilon_i\}$, and compute the limit as $\epsilon_i \rightarrow 0$

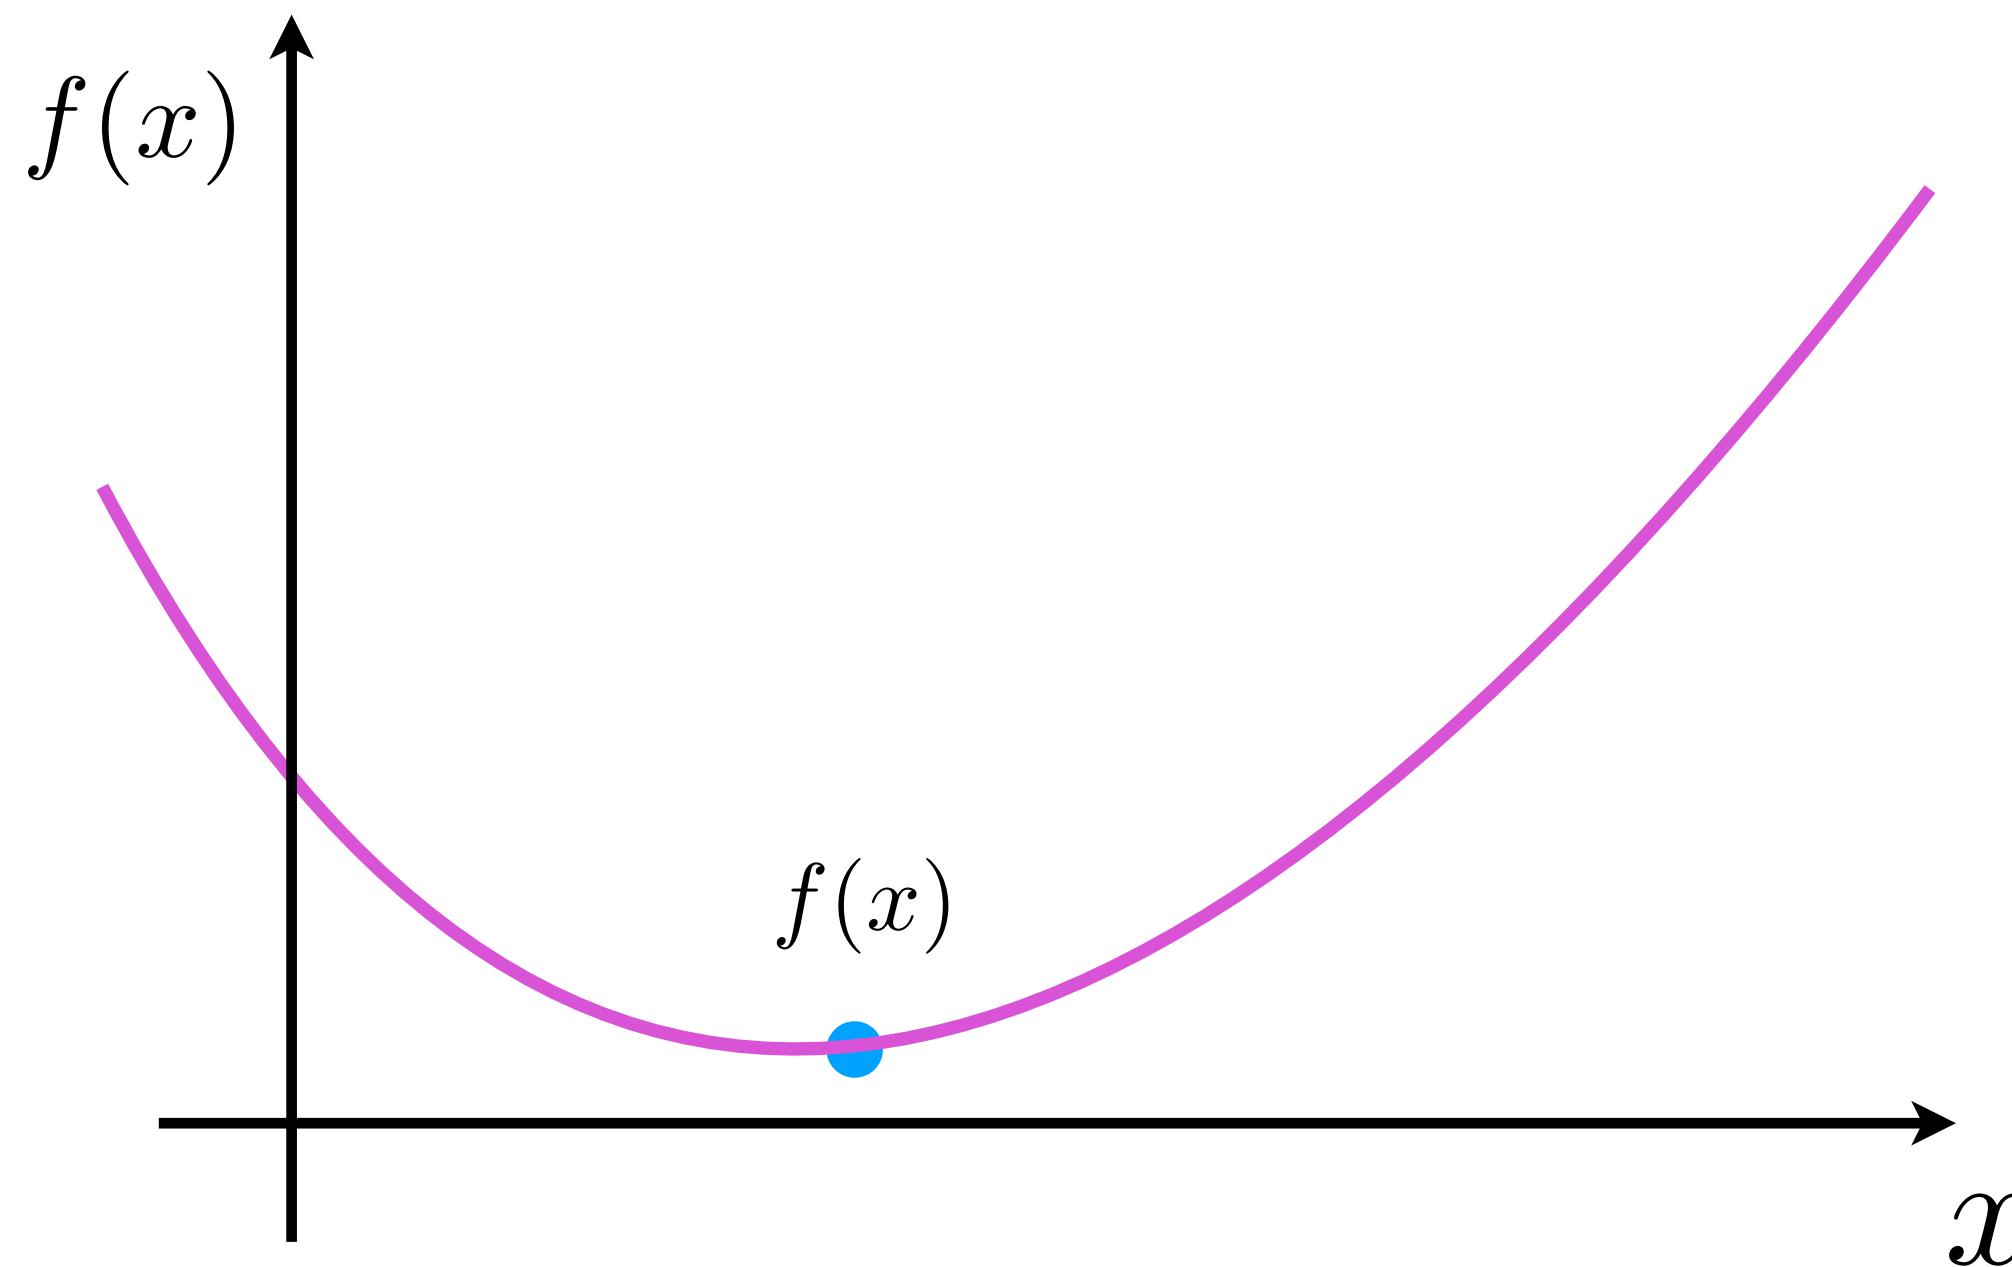


Derivatives and gradients

- Definition of a **derivative**

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad | \quad \frac{\partial f}{\partial x} = f'(x) = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x)}{\epsilon}$$

- Intuition: generate a sequence of points $\{f(x + \epsilon_i), \epsilon_i\}$, and compute the limit as $\epsilon_i \rightarrow 0$

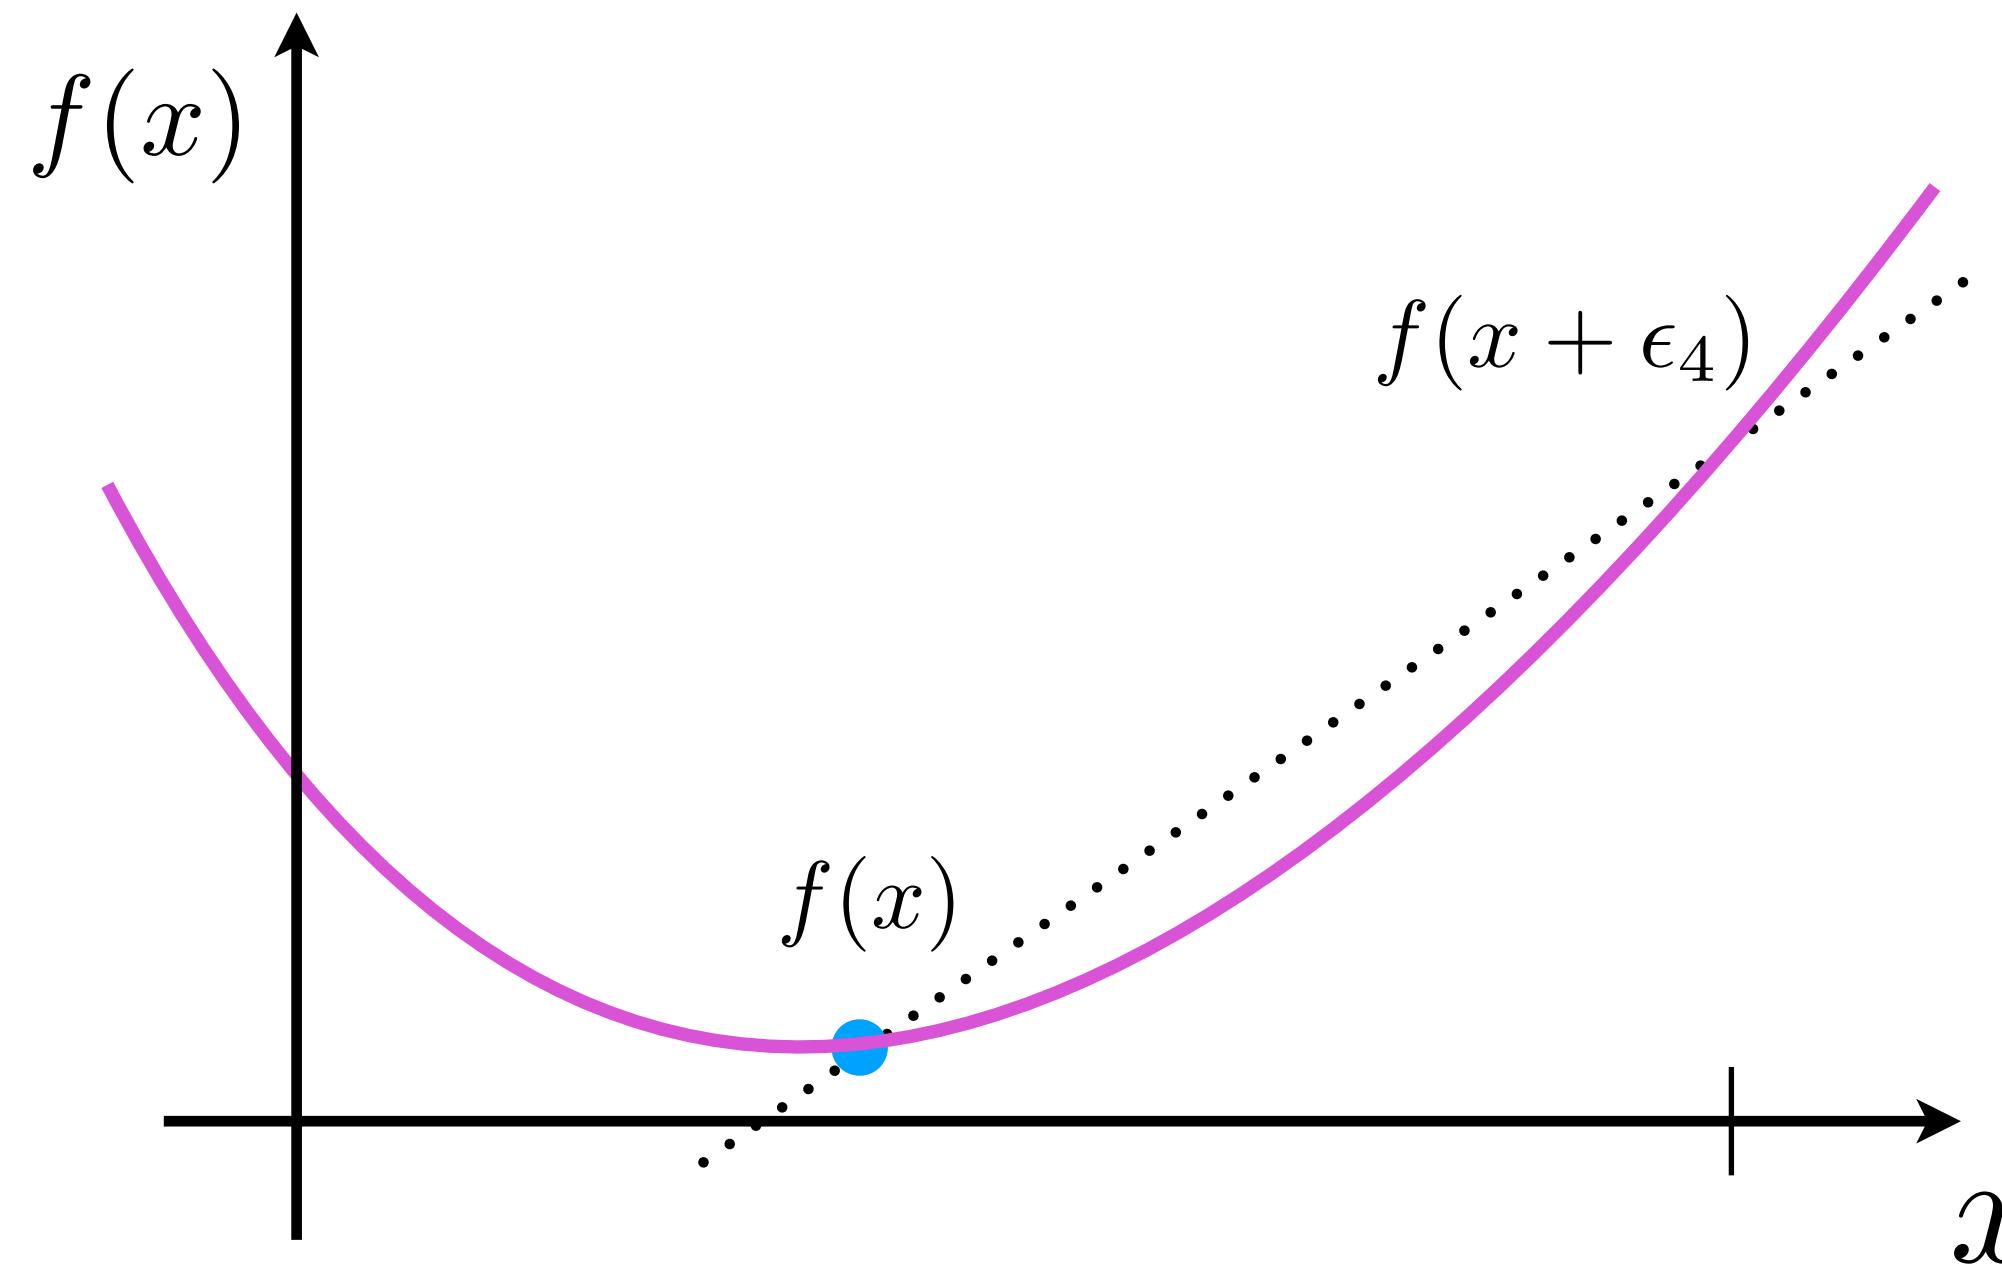


Derivatives and gradients

- Definition of a **derivative**

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad | \quad \frac{\partial f}{\partial x} = f'(x) = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x)}{\epsilon}$$

- Intuition: generate a sequence of points $\{f(x + \epsilon_i), \epsilon_i\}$, and compute the limit as $\epsilon_i \rightarrow 0$

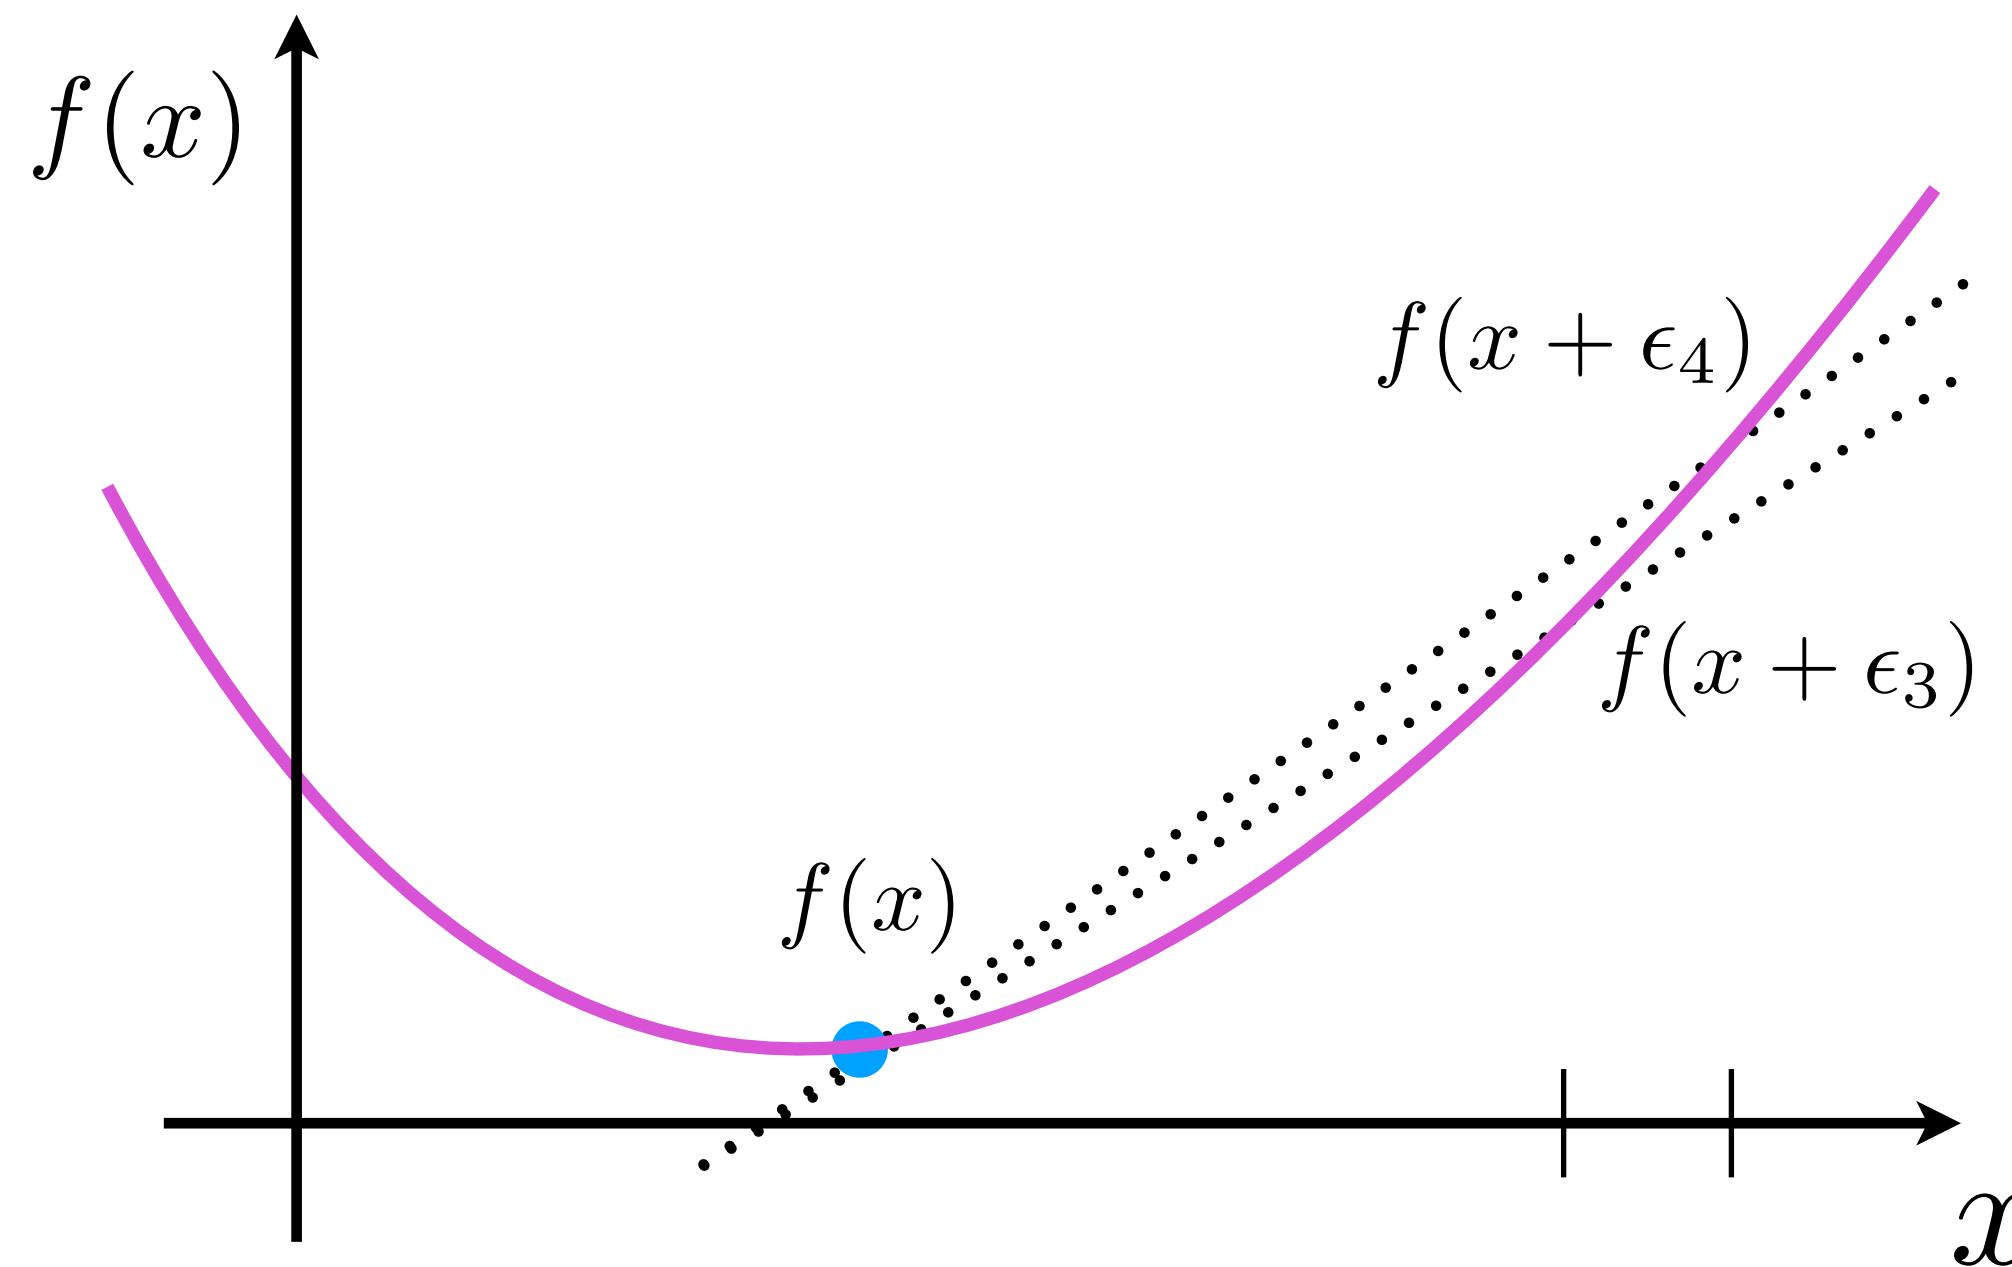


Derivatives and gradients

- Definition of a **derivative**

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad | \quad \frac{\partial f}{\partial x} = f'(x) = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x)}{\epsilon}$$

- Intuition: generate a sequence of points $\{f(x + \epsilon_i), \epsilon_i\}$, and compute the limit as $\epsilon_i \rightarrow 0$

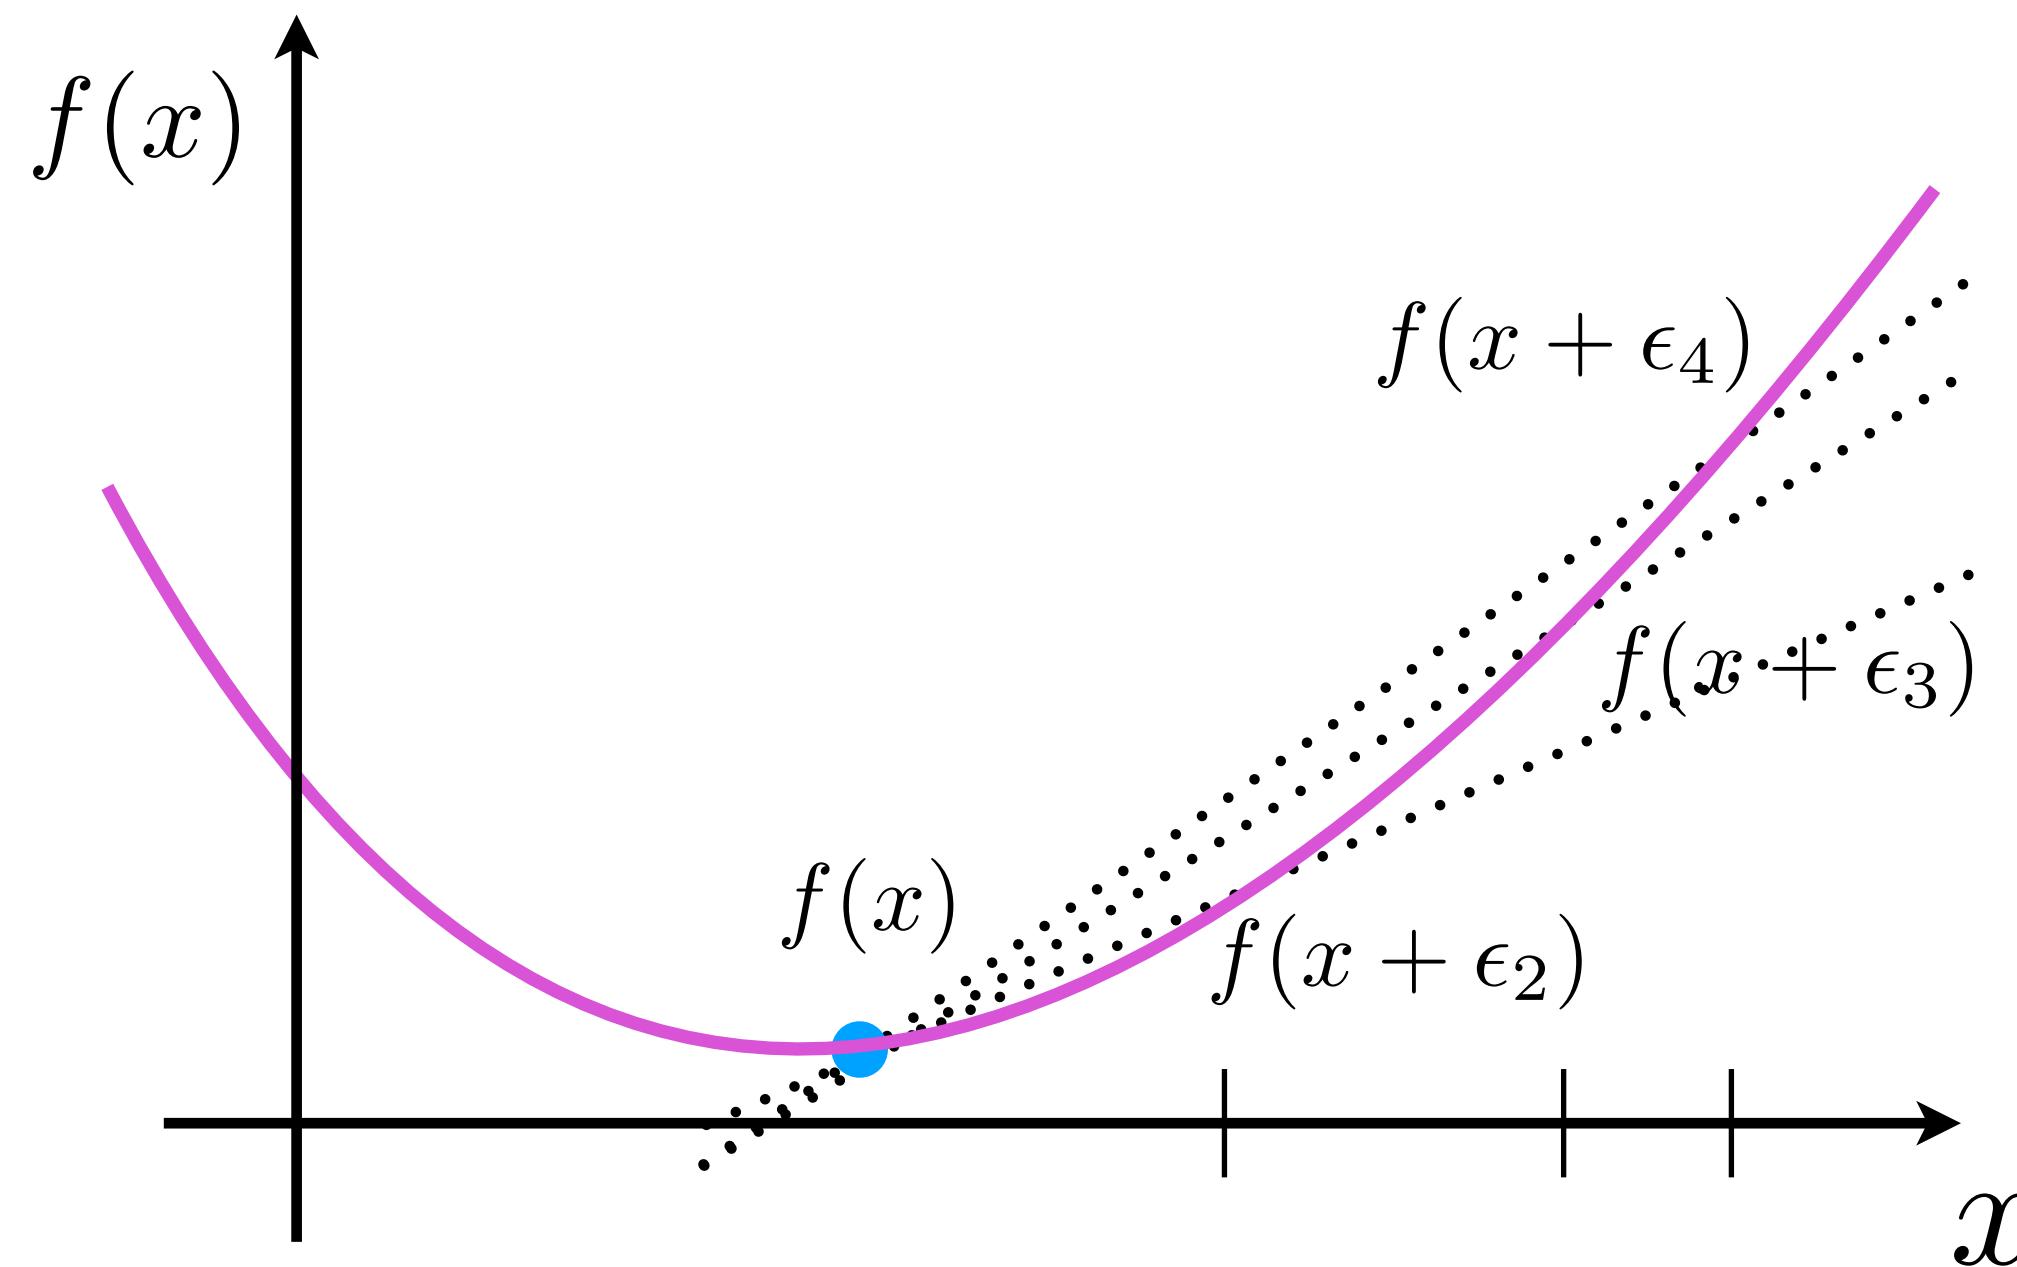


Derivatives and gradients

- Definition of a **derivative**

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad | \quad \frac{\partial f}{\partial x} = f'(x) = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x)}{\epsilon}$$

- Intuition: generate a sequence of points $\{f(x + \epsilon_i), \epsilon_i\}$, and compute the limit as $\epsilon_i \rightarrow 0$

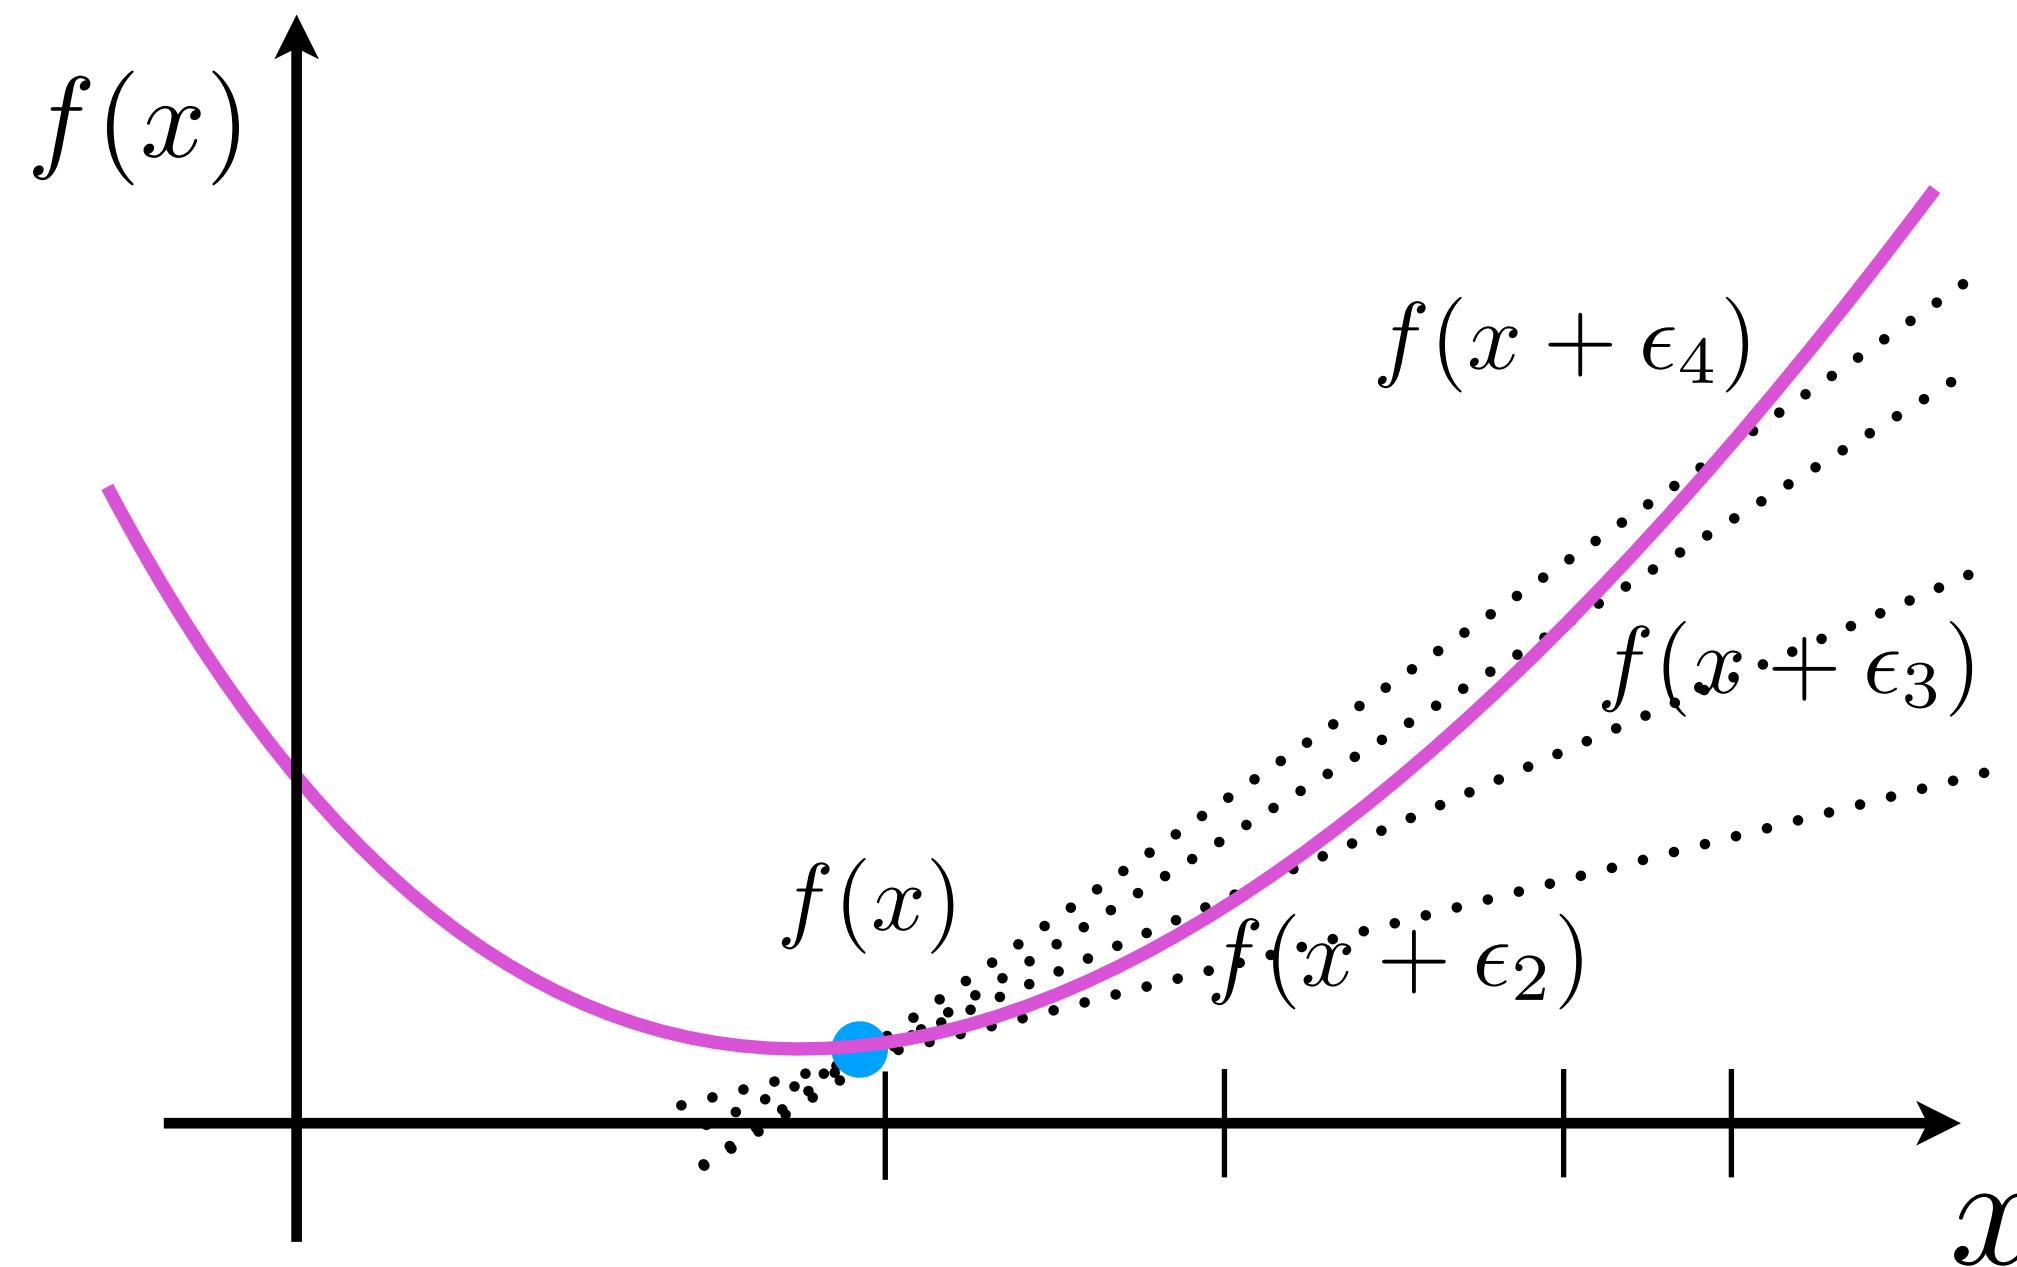


Derivatives and gradients

- Definition of a **derivative**

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad | \quad \frac{\partial f}{\partial x} = f'(x) = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x)}{\epsilon}$$

- Intuition: generate a sequence of points $\{f(x + \epsilon_i), \epsilon_i\}$, and compute the limit as $\epsilon_i \rightarrow 0$



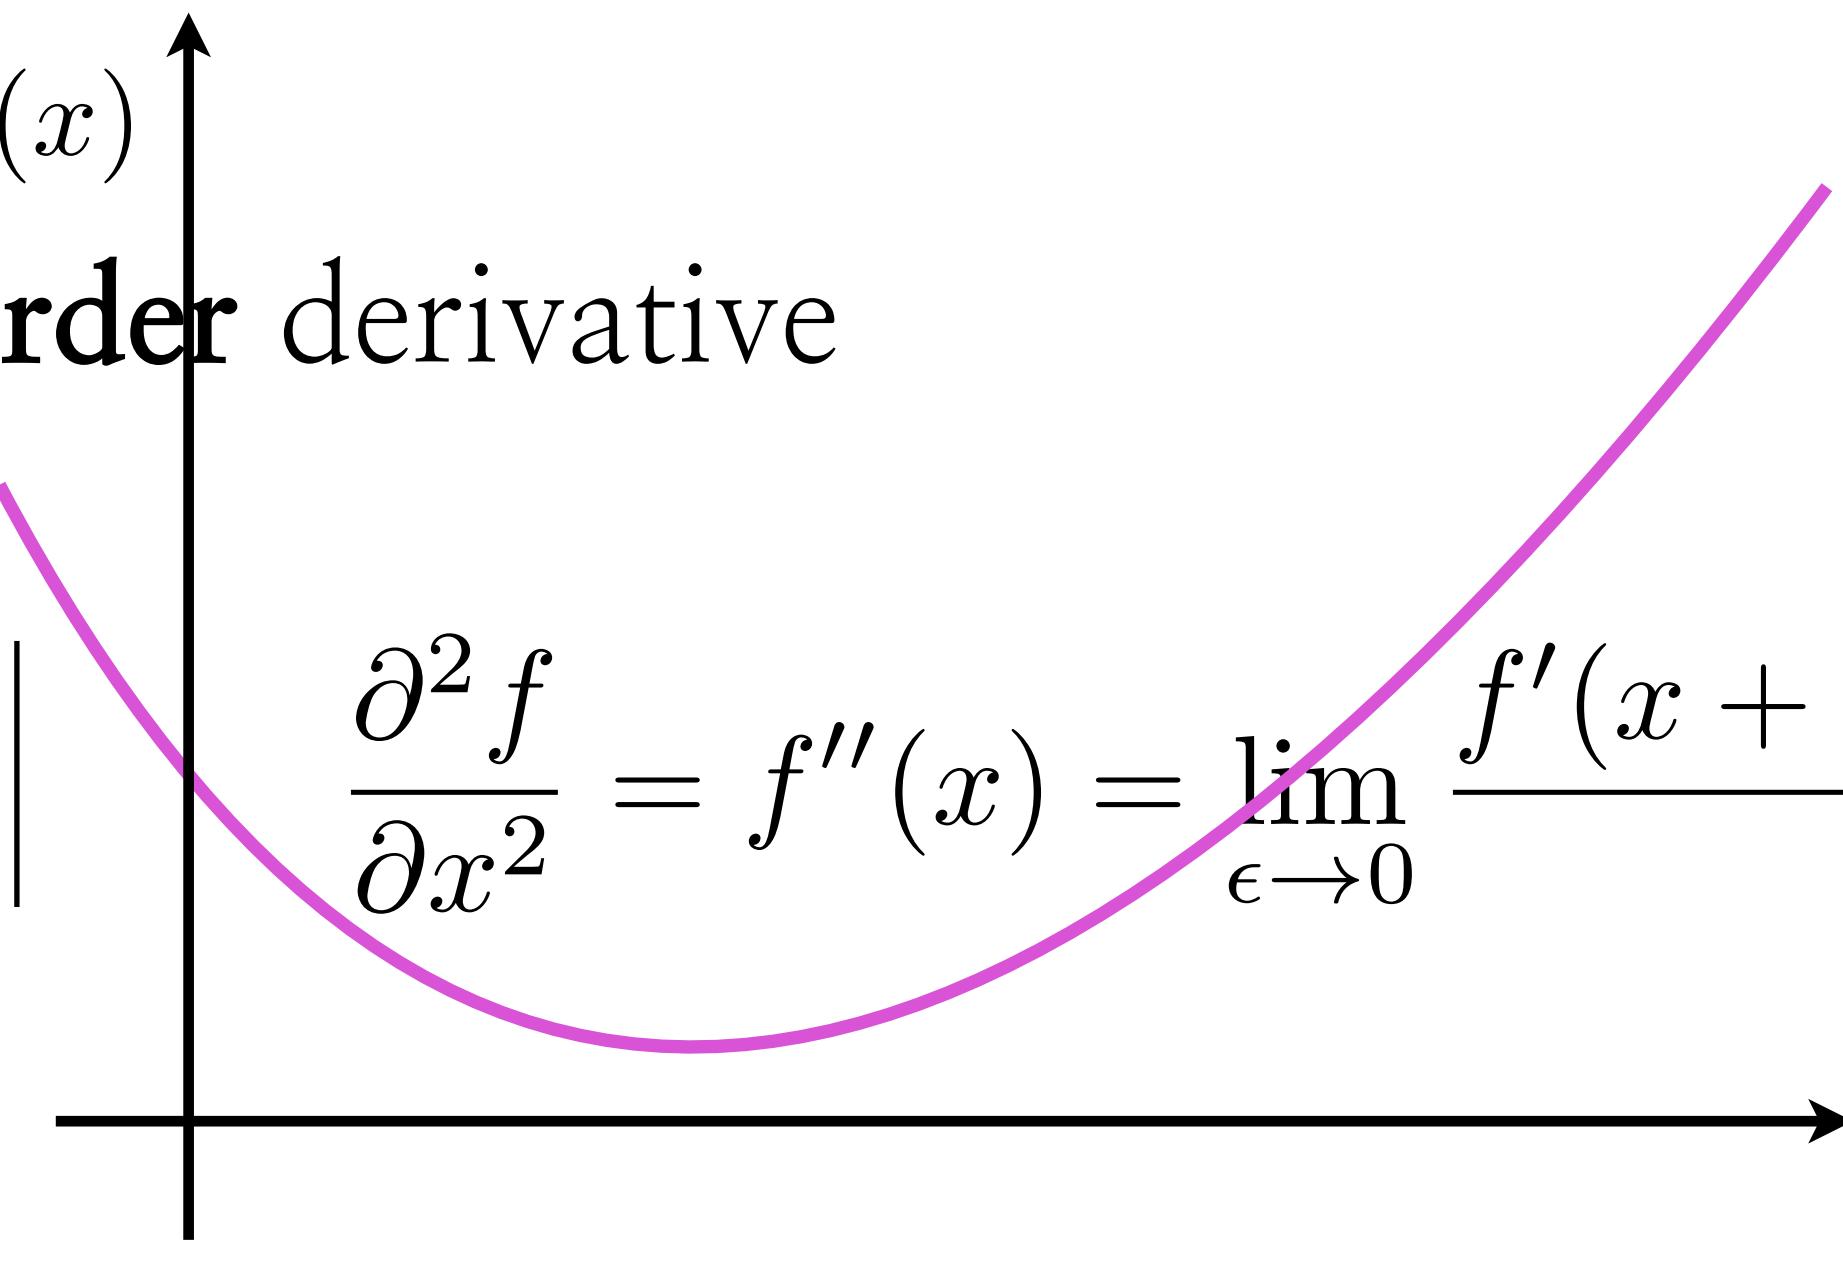
Derivatives and gradients

- Definition of a **derivative**

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad | \quad \frac{\partial f}{\partial x} = f'(x) = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x)}{\epsilon}$$

- Intuition: generate a sequence of points $\{f(x + \epsilon_i), \epsilon_i\}$, and compute the limit as $\epsilon_i \rightarrow 0$

- Definition of **second-order** derivative

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad | \quad \frac{\partial^2 f}{\partial x^2} = f''(x) = \lim_{\epsilon \rightarrow 0} \frac{f'(x + \epsilon) - f'(x)}{\epsilon}$$


Derivatives and gradients

- Definition of a **derivative**

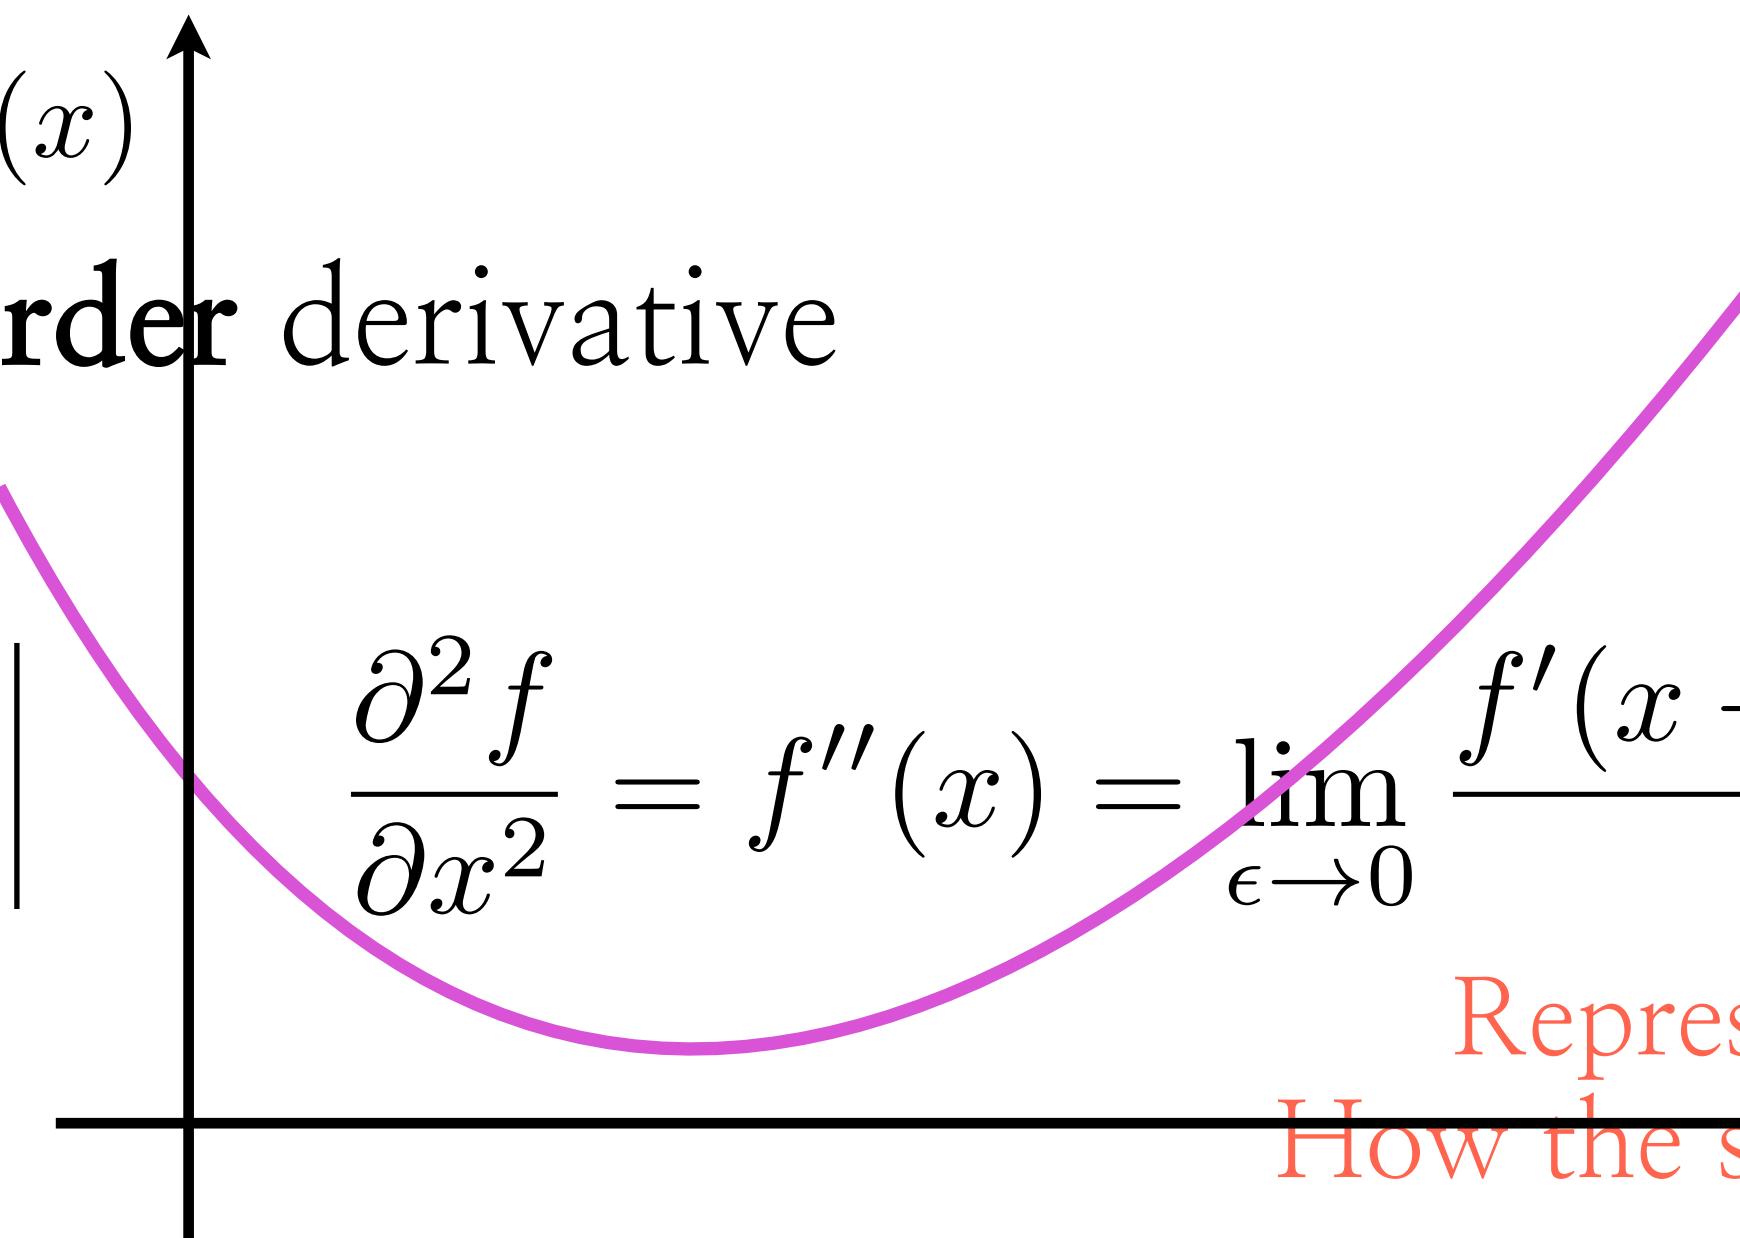
$$f : \mathbb{R} \rightarrow \mathbb{R} \quad | \quad \frac{\partial f}{\partial x} = f'(x) = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x)}{\epsilon}$$

- Intuition: generate a sequence of points $\{f(x + \epsilon_i), \epsilon_i\}$, and compute the limit as $\epsilon_i \rightarrow 0$

- Definition of **second-order** derivative

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad | \quad \frac{\partial^2 f}{\partial x^2} = f''(x) = \lim_{\epsilon \rightarrow 0} \frac{f'(x + \epsilon) - f'(x)}{\epsilon}$$

Represents the local curvature:
How the slope of the function changes



Derivatives and gradients

- Generalization to multiple components: **gradient**

$$f : \mathbb{R}^p \rightarrow \mathbb{R} \quad | \quad \nabla f(x) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_p} \end{bmatrix} \in \mathbb{R}^p$$

where

$$\frac{\partial f}{\partial x_i} = \lim_{\epsilon \rightarrow 0} \frac{f(x_1, \dots, x_{i-1}, x_i + \epsilon, x_{i+1}, \dots, x_p) - f(x_1, \dots, x_{i-1}, x_i, x_{i+1}, \dots, x_p)}{\epsilon} = \frac{f(x + \epsilon e_i) - f(x)}{\epsilon}$$

Derivatives and gradients

- **Jacobian** matrix (relates to neural networks)

$$f : \mathbb{R}^p \rightarrow \mathbb{R}^m \quad | \quad Df(x) = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \frac{\partial f_1}{\partial x_2} & \cdots & \frac{\partial f_1}{\partial x_p} \\ \vdots & \vdots & & \vdots \\ \frac{\partial f_m}{\partial x_1} & \frac{\partial f_m}{\partial x_2} & \cdots & \frac{\partial f_m}{\partial x_p} \end{bmatrix} \in \mathbb{R}^{m \times p}$$

- Generalizes the notion of gradient to multiple-output functions

Derivatives and gradients

- Hessian matrix

$$f : \mathbb{R}^p \rightarrow \mathbb{R} \quad | \quad \nabla^2 f(x) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_p} \\ \frac{\partial^2 f}{\partial x_1 \partial x_2} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_p} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_p \partial x_1} & \frac{\partial^2 f}{\partial x_p \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_p^2} \end{bmatrix} \in \mathbb{R}^{p \times p}$$

Derivatives and gradients

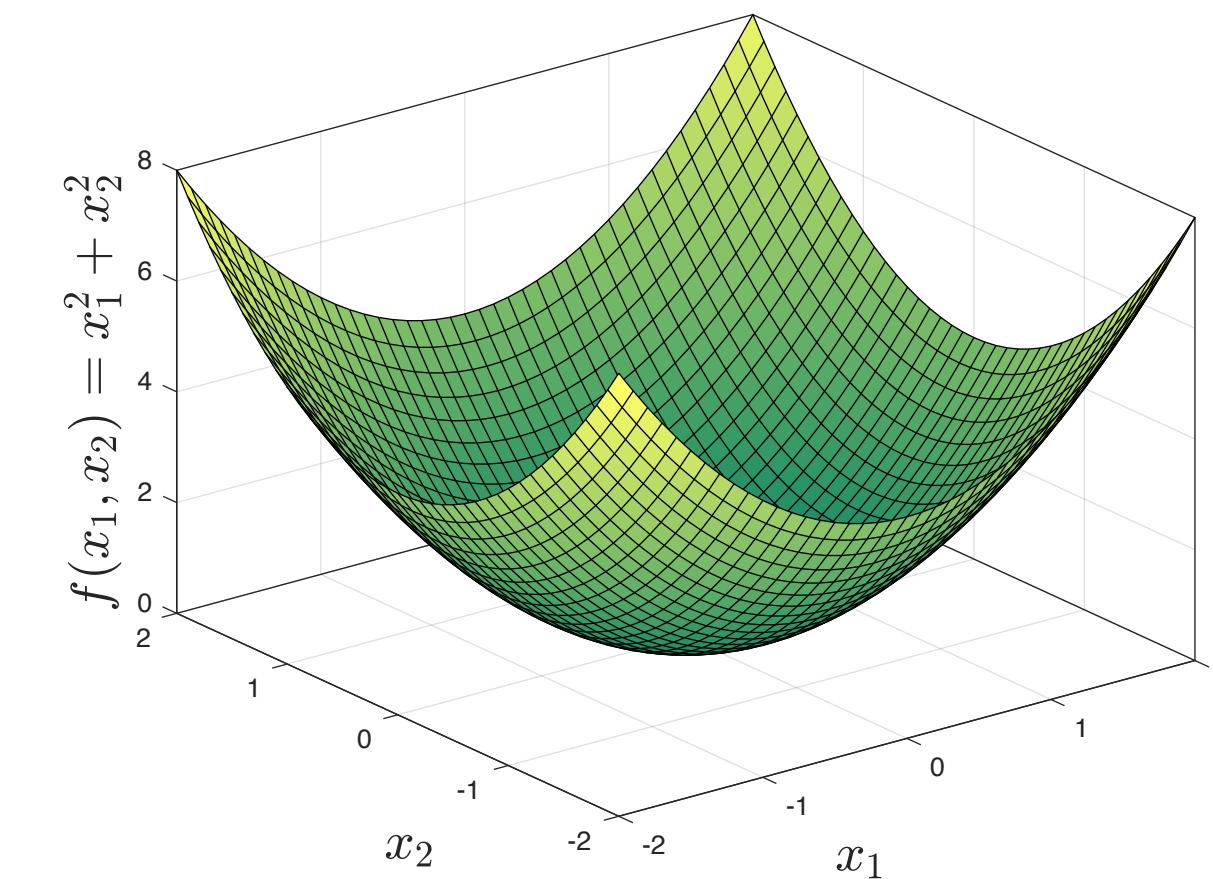
- Hessian matrix

$$f : \mathbb{R}^p \rightarrow \mathbb{R}$$

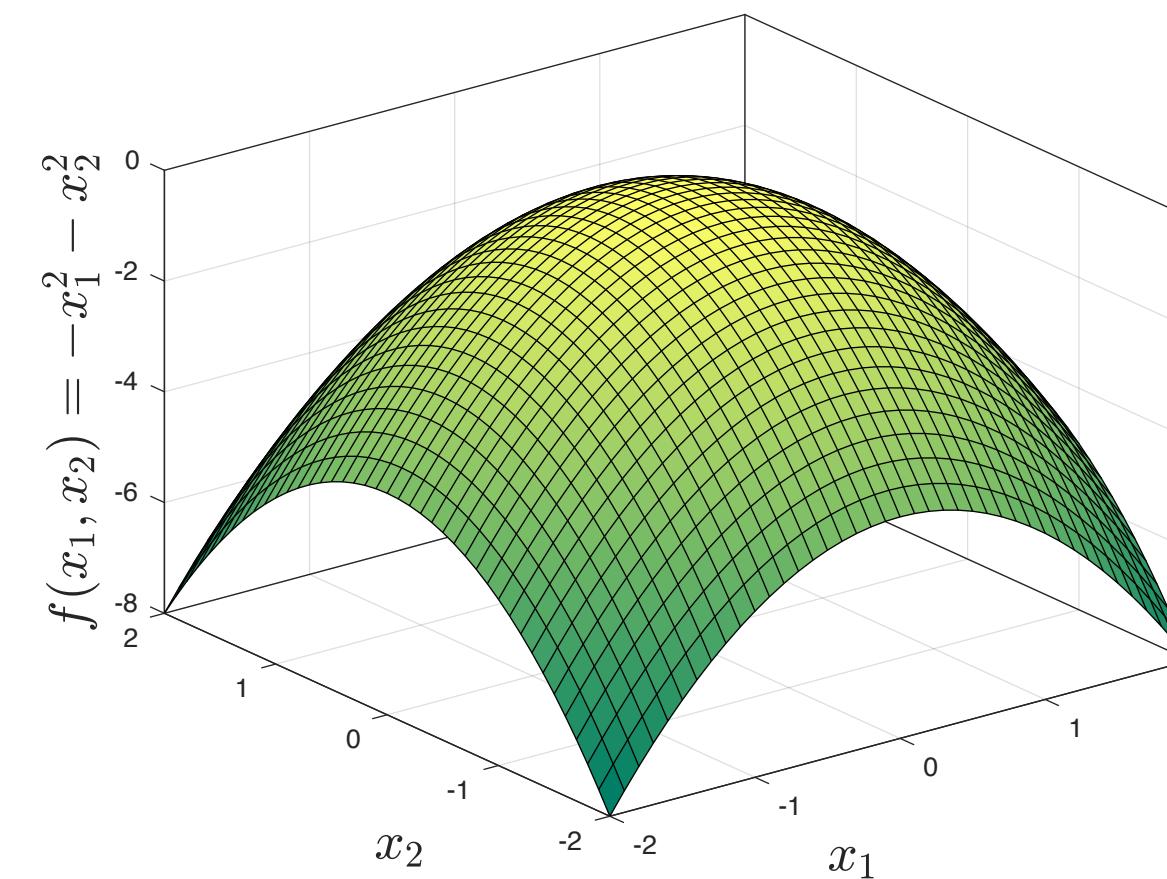
$$\mid$$

$$\nabla^2 f(x) =$$

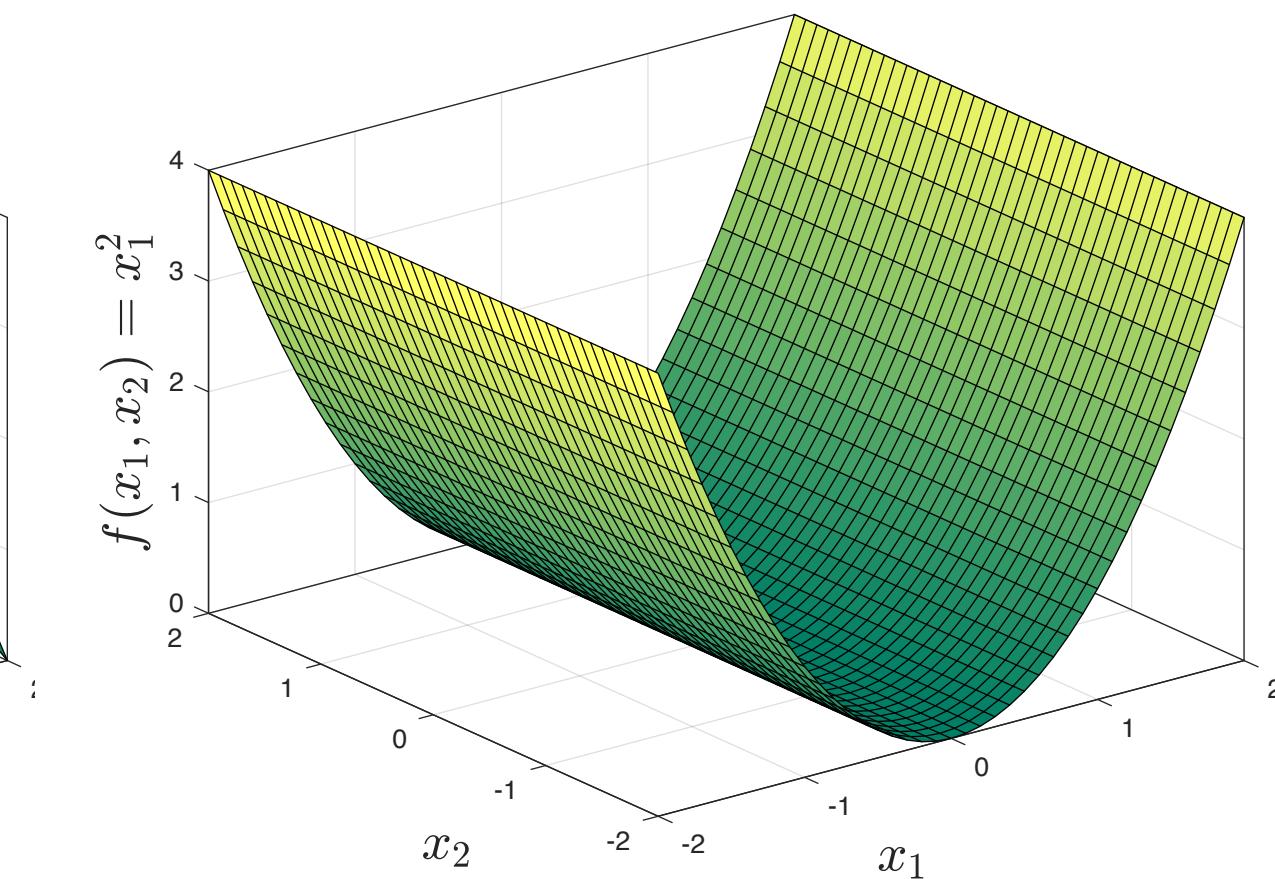
$$\begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_p} \\ \frac{\partial^2 f}{\partial x_1 \partial x_2} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_p} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_p \partial x_1} & \frac{\partial^2 f}{\partial x_p \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_p^2} \end{bmatrix} \in \mathbb{R}^{p \times p}$$



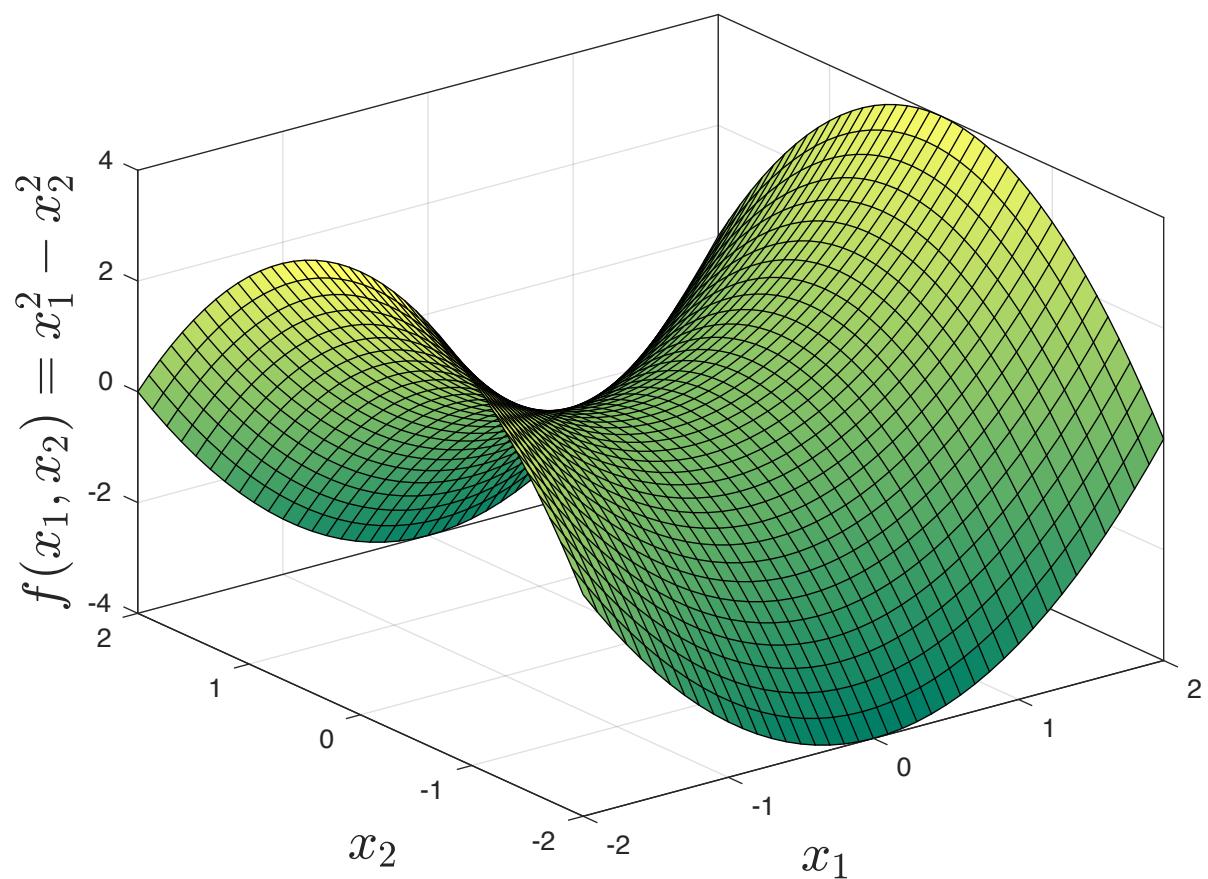
$\nabla^2 f(\cdot) \succ 0$



$\nabla^2 f(\cdot) \prec 0$



$\nabla^2 f(\cdot) \succcurlyeq 0$

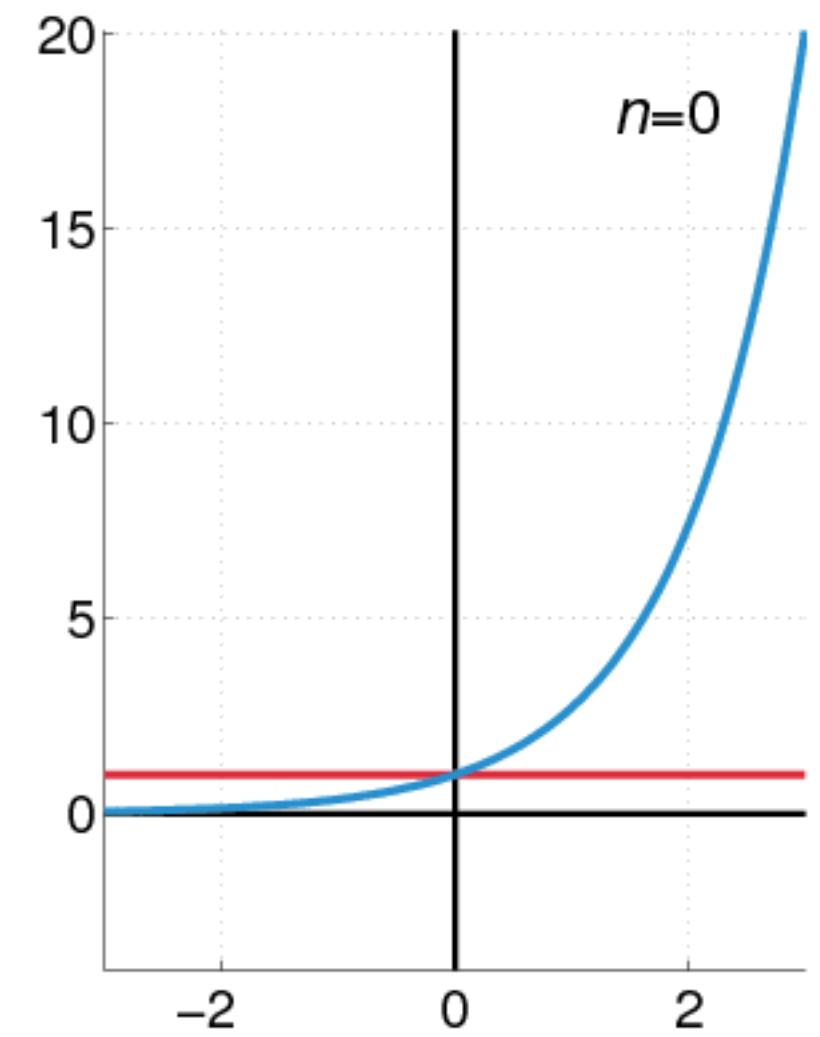


Indefinite

Taylor's expansion

- Taylor's expansion: used for (locally) approximating a function

$$f(x)\Big|_{x=\alpha} = f(\alpha) + f'(\alpha)(x - \alpha) + \frac{f''(\alpha)}{2!}(x - \alpha)^2 + \cdots + \frac{f^{(n)}(\alpha)}{n!}(x - \alpha)^n + R_n$$

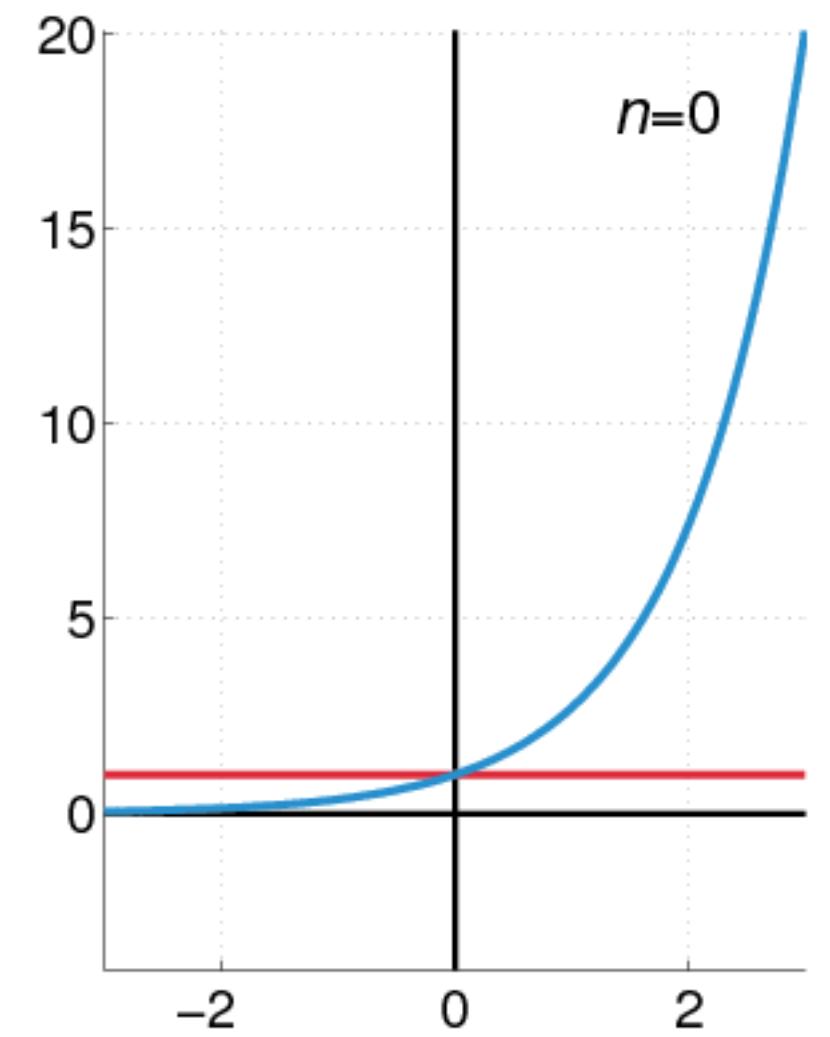


Example:
exponential function

Taylor's expansion

- Taylor's expansion: used for (locally) approximating a function

$$f(x)\Big|_{x=\alpha} = f(\alpha) + f'(\alpha)(x - \alpha) + \frac{f''(\alpha)}{2!}(x - \alpha)^2 + \cdots + \frac{f^{(n)}(\alpha)}{n!}(x - \alpha)^n + R_n$$

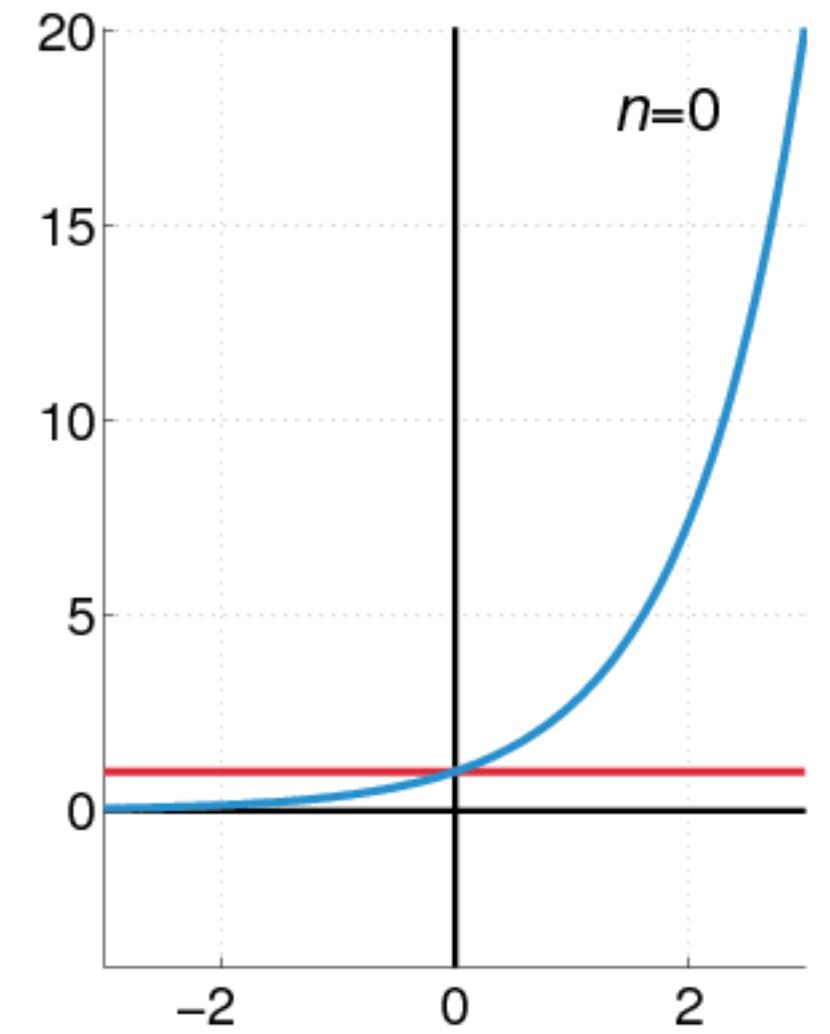


Example:
exponential function

Taylor's expansion

- Taylor's expansion: used for (locally) approximating a function

$$f(x)\Big|_{x=\alpha} = f(\alpha) + f'(\alpha)(x - \alpha) + \frac{f''(\alpha)}{2!}(x - \alpha)^2 + \cdots + \frac{f^{(n)}(\alpha)}{n!}(x - \alpha)^n + R_n$$



- Key properties/assumptions:
 - Function f is differentiable as many times we'd like
 - Provides (locally) a good approximation of the function

Example:
exponential function

Taylor's expansion

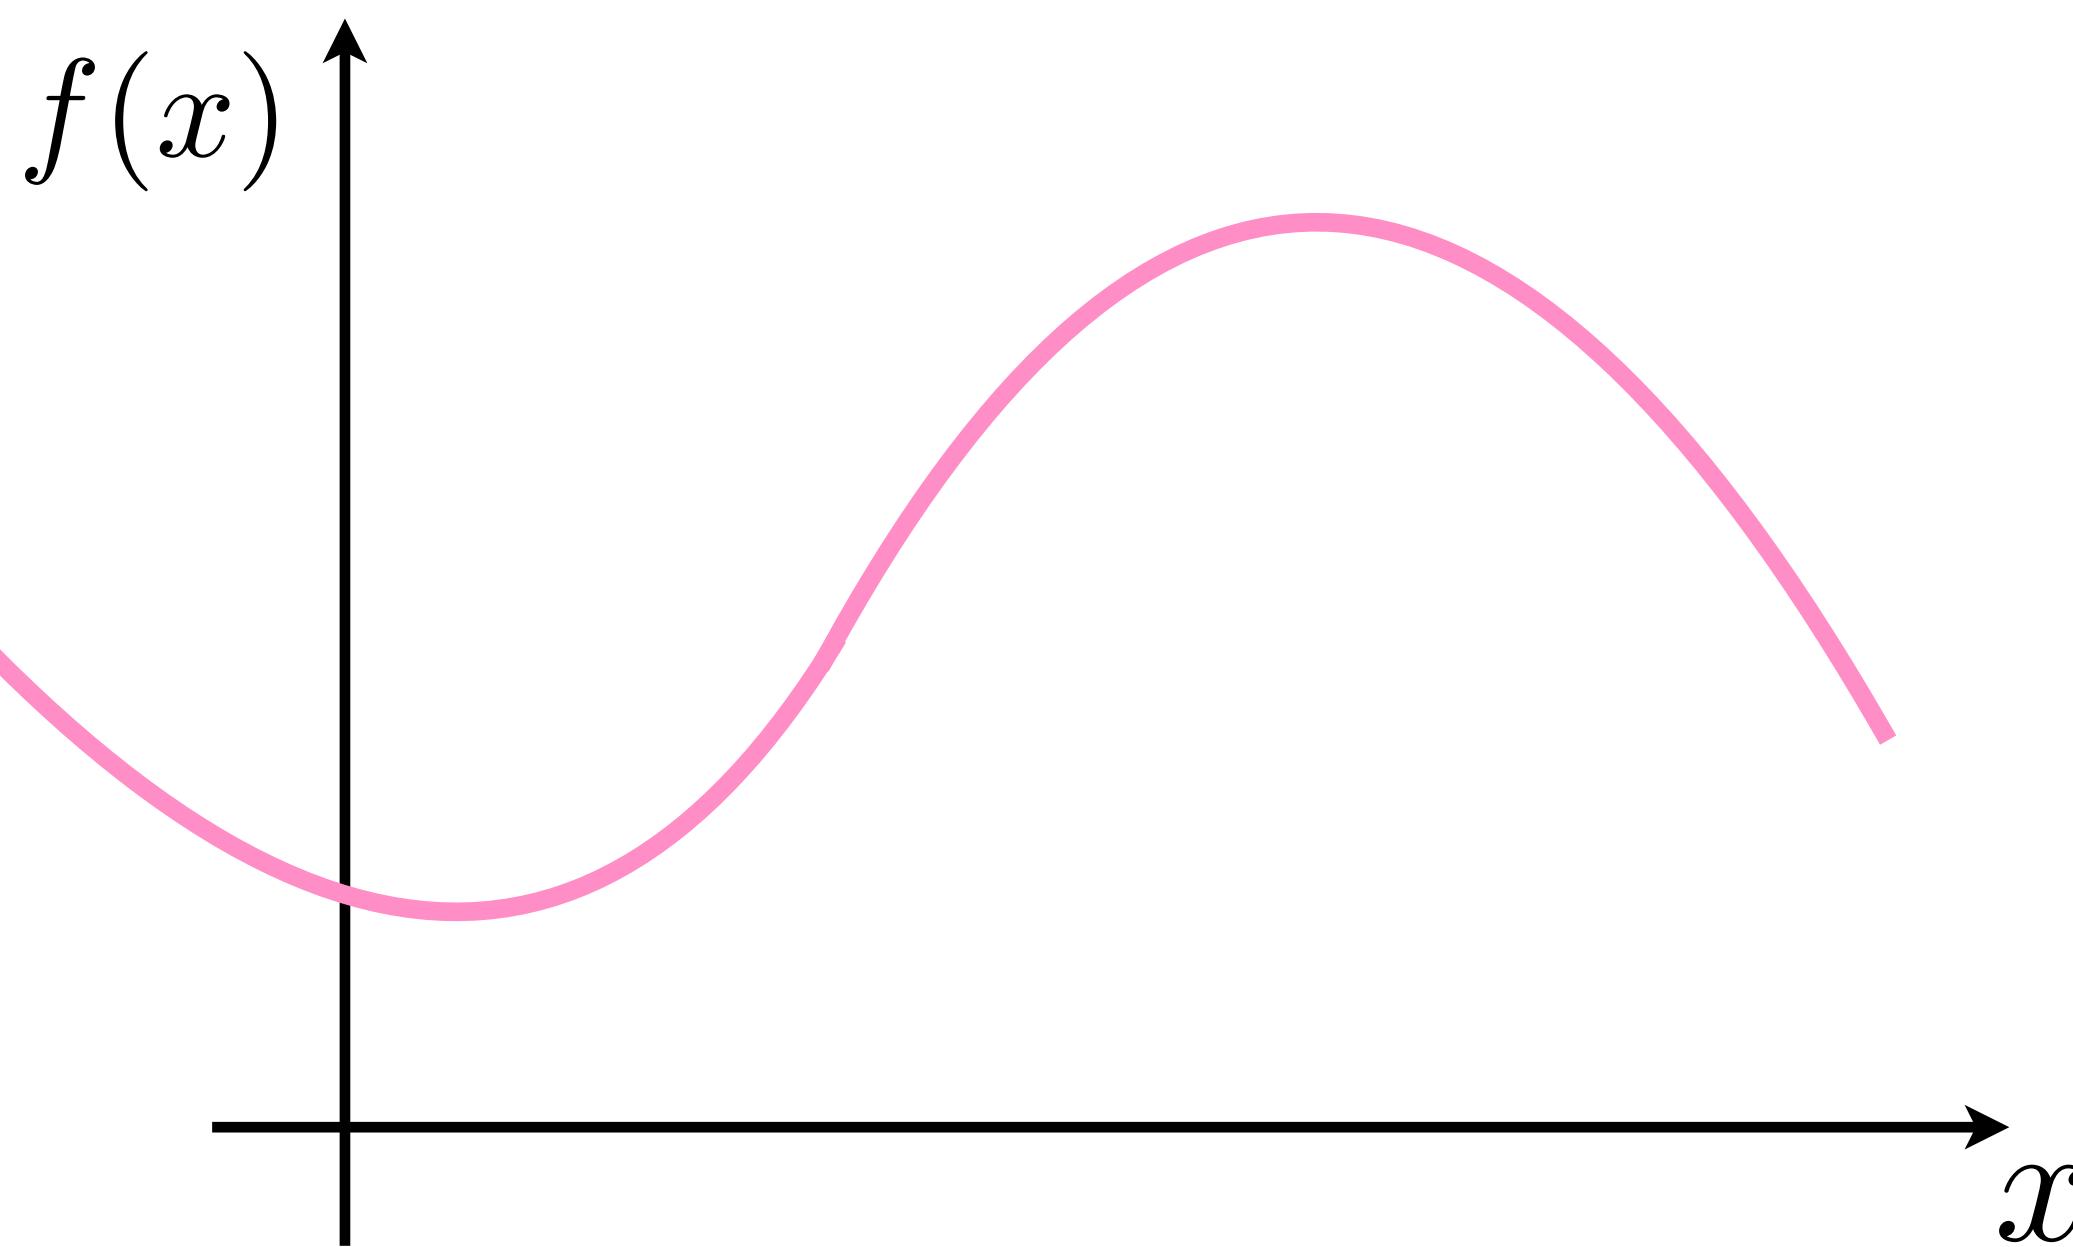
- First-order Taylor's approximation

$$f : \mathbb{R}^p \rightarrow \mathbb{R} \quad \Big| \quad f(x) \approx f(\alpha) + \langle \nabla f(\alpha), x - \alpha \rangle, \alpha \in \mathbb{R}^p$$

Taylor's expansion

- First-order Taylor's approximation

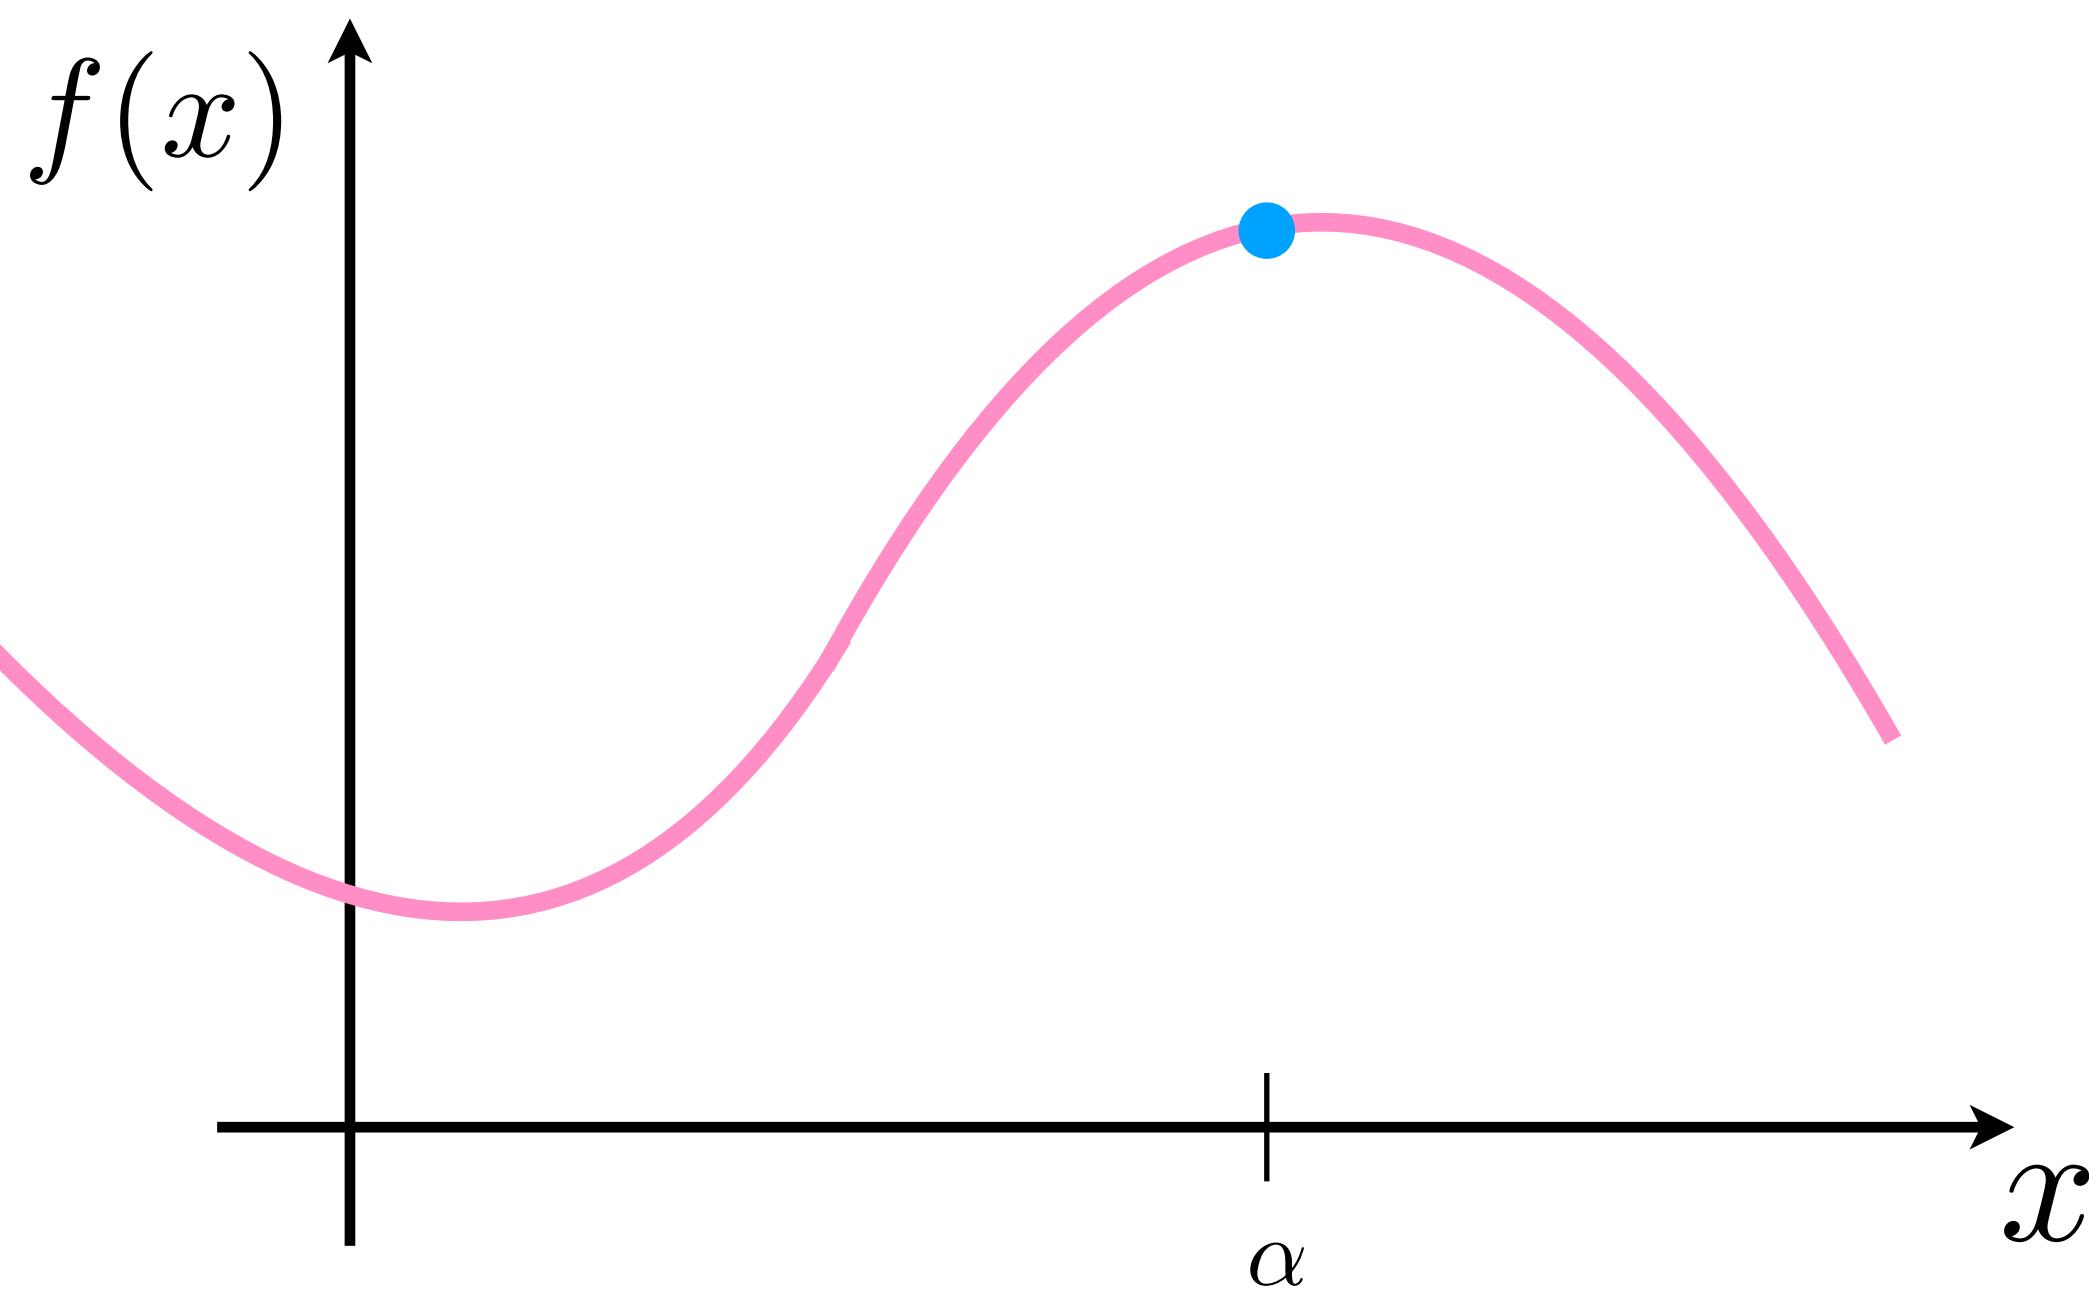
$$f : \mathbb{R}^p \rightarrow \mathbb{R} \quad | \quad f(x) \approx f(\alpha) + \langle \nabla f(\alpha), x - \alpha \rangle, \alpha \in \mathbb{R}^p$$



Taylor's expansion

- First-order Taylor's approximation

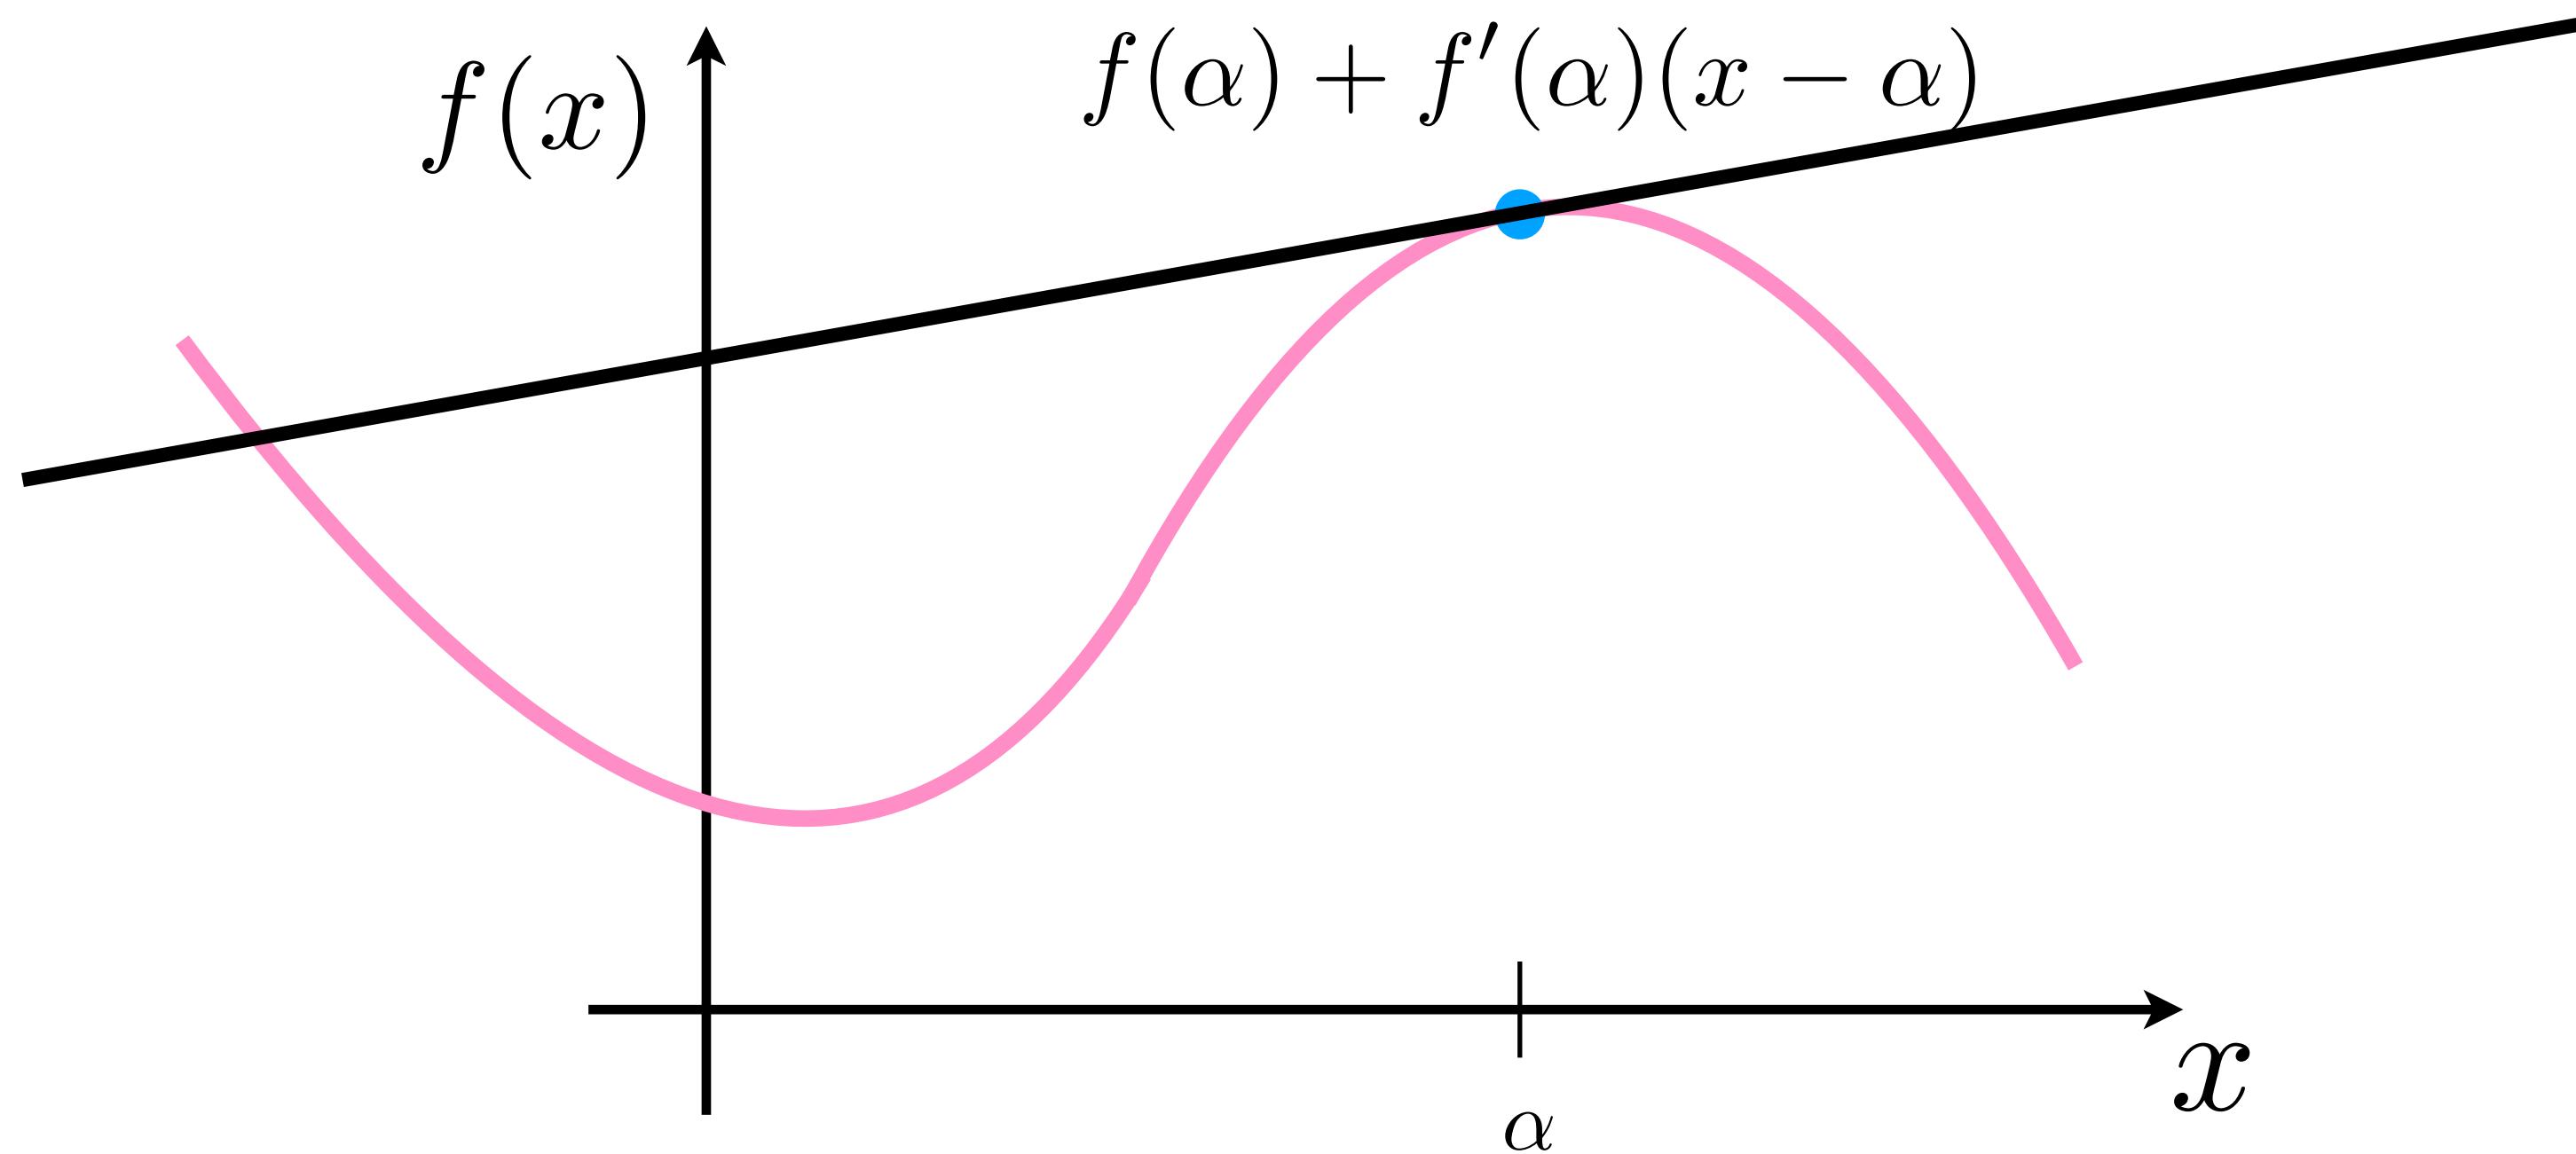
$$f : \mathbb{R}^p \rightarrow \mathbb{R} \quad | \quad f(x) \approx f(\alpha) + \langle \nabla f(\alpha), x - \alpha \rangle, \alpha \in \mathbb{R}^p$$



Taylor's expansion

- First-order Taylor's approximation

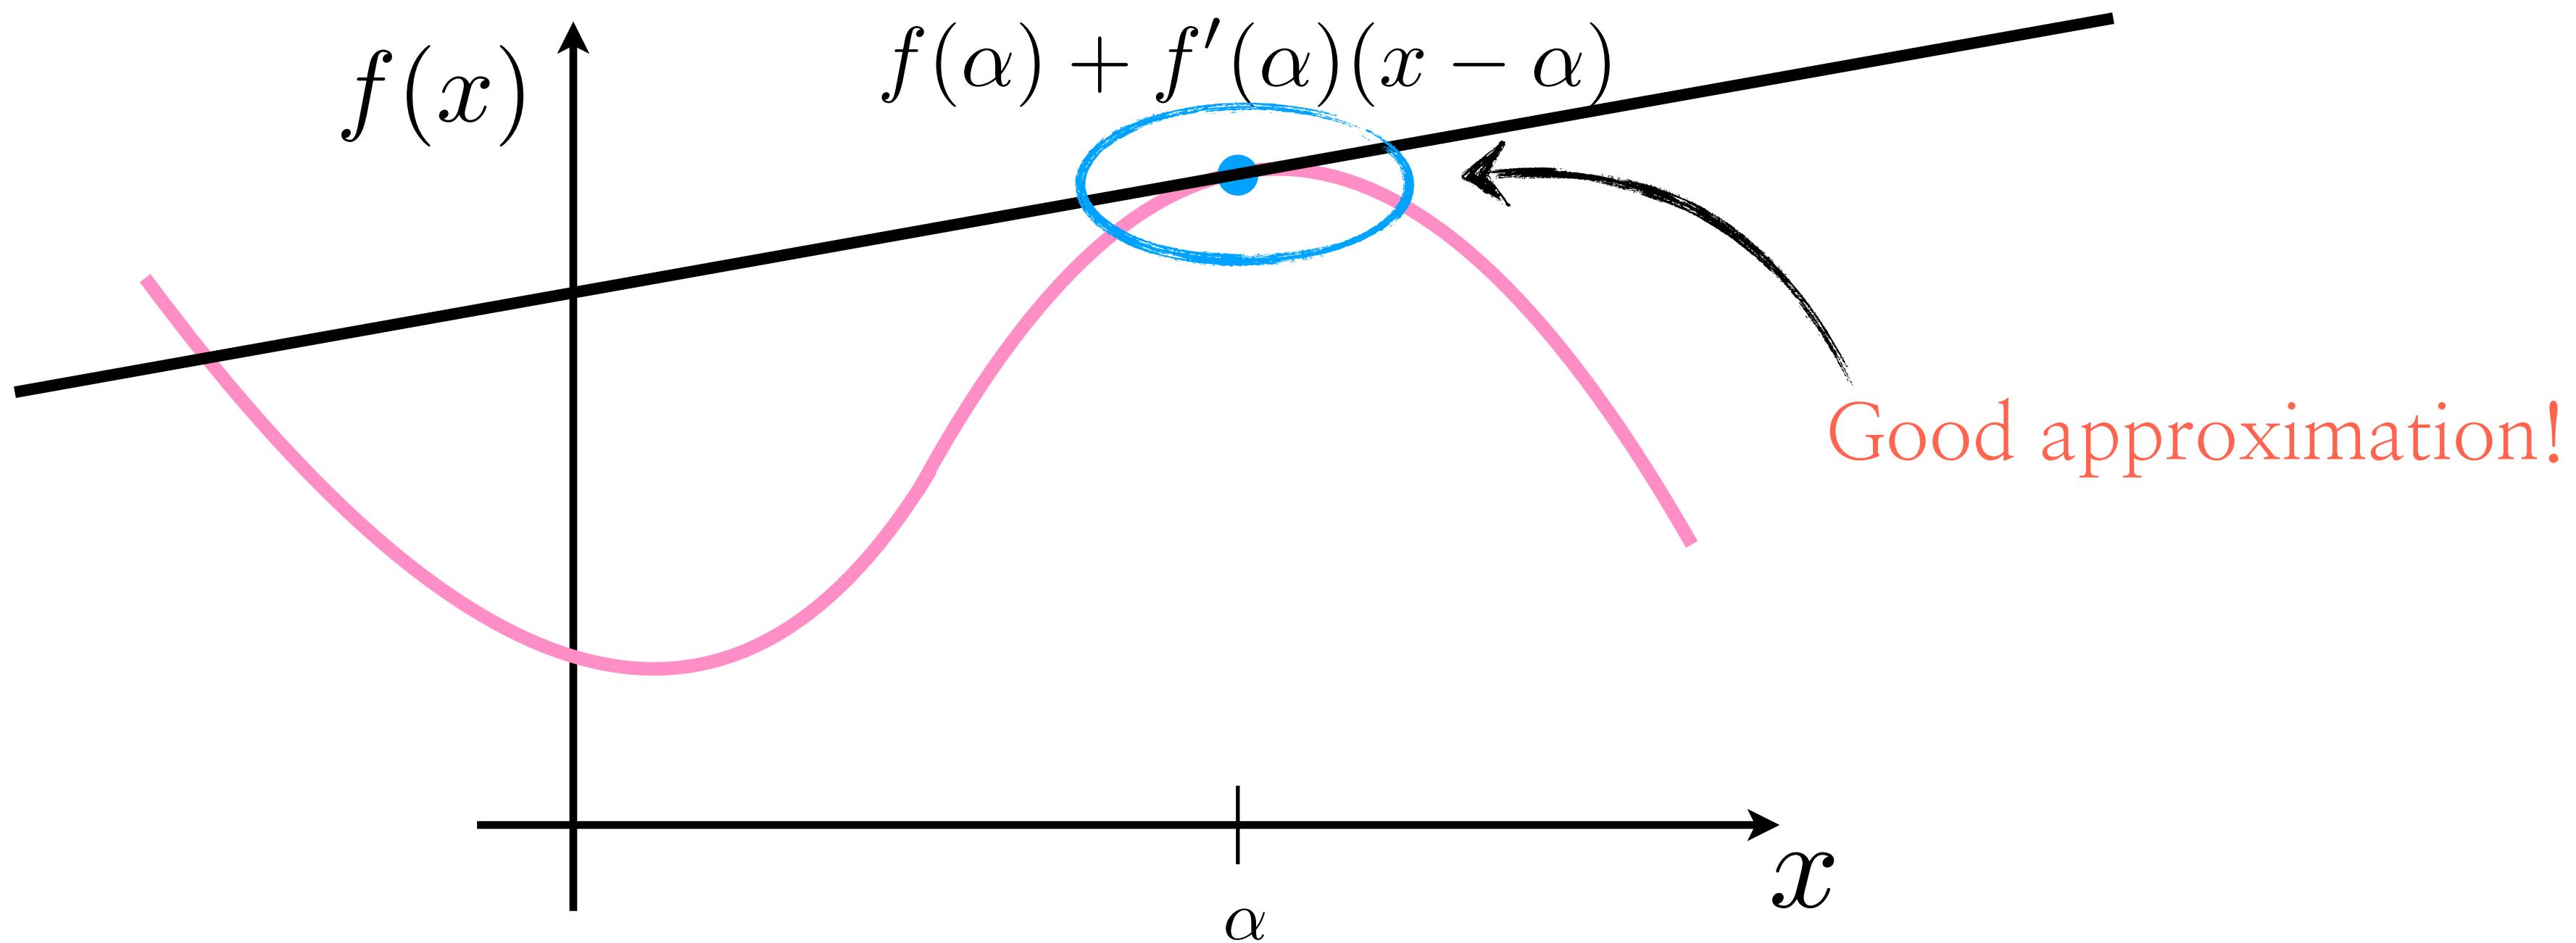
$$f : \mathbb{R}^p \rightarrow \mathbb{R} \quad | \quad f(x) \approx f(\alpha) + \langle \nabla f(\alpha), x - \alpha \rangle, \alpha \in \mathbb{R}^p$$



Taylor's expansion

- First-order Taylor's approximation

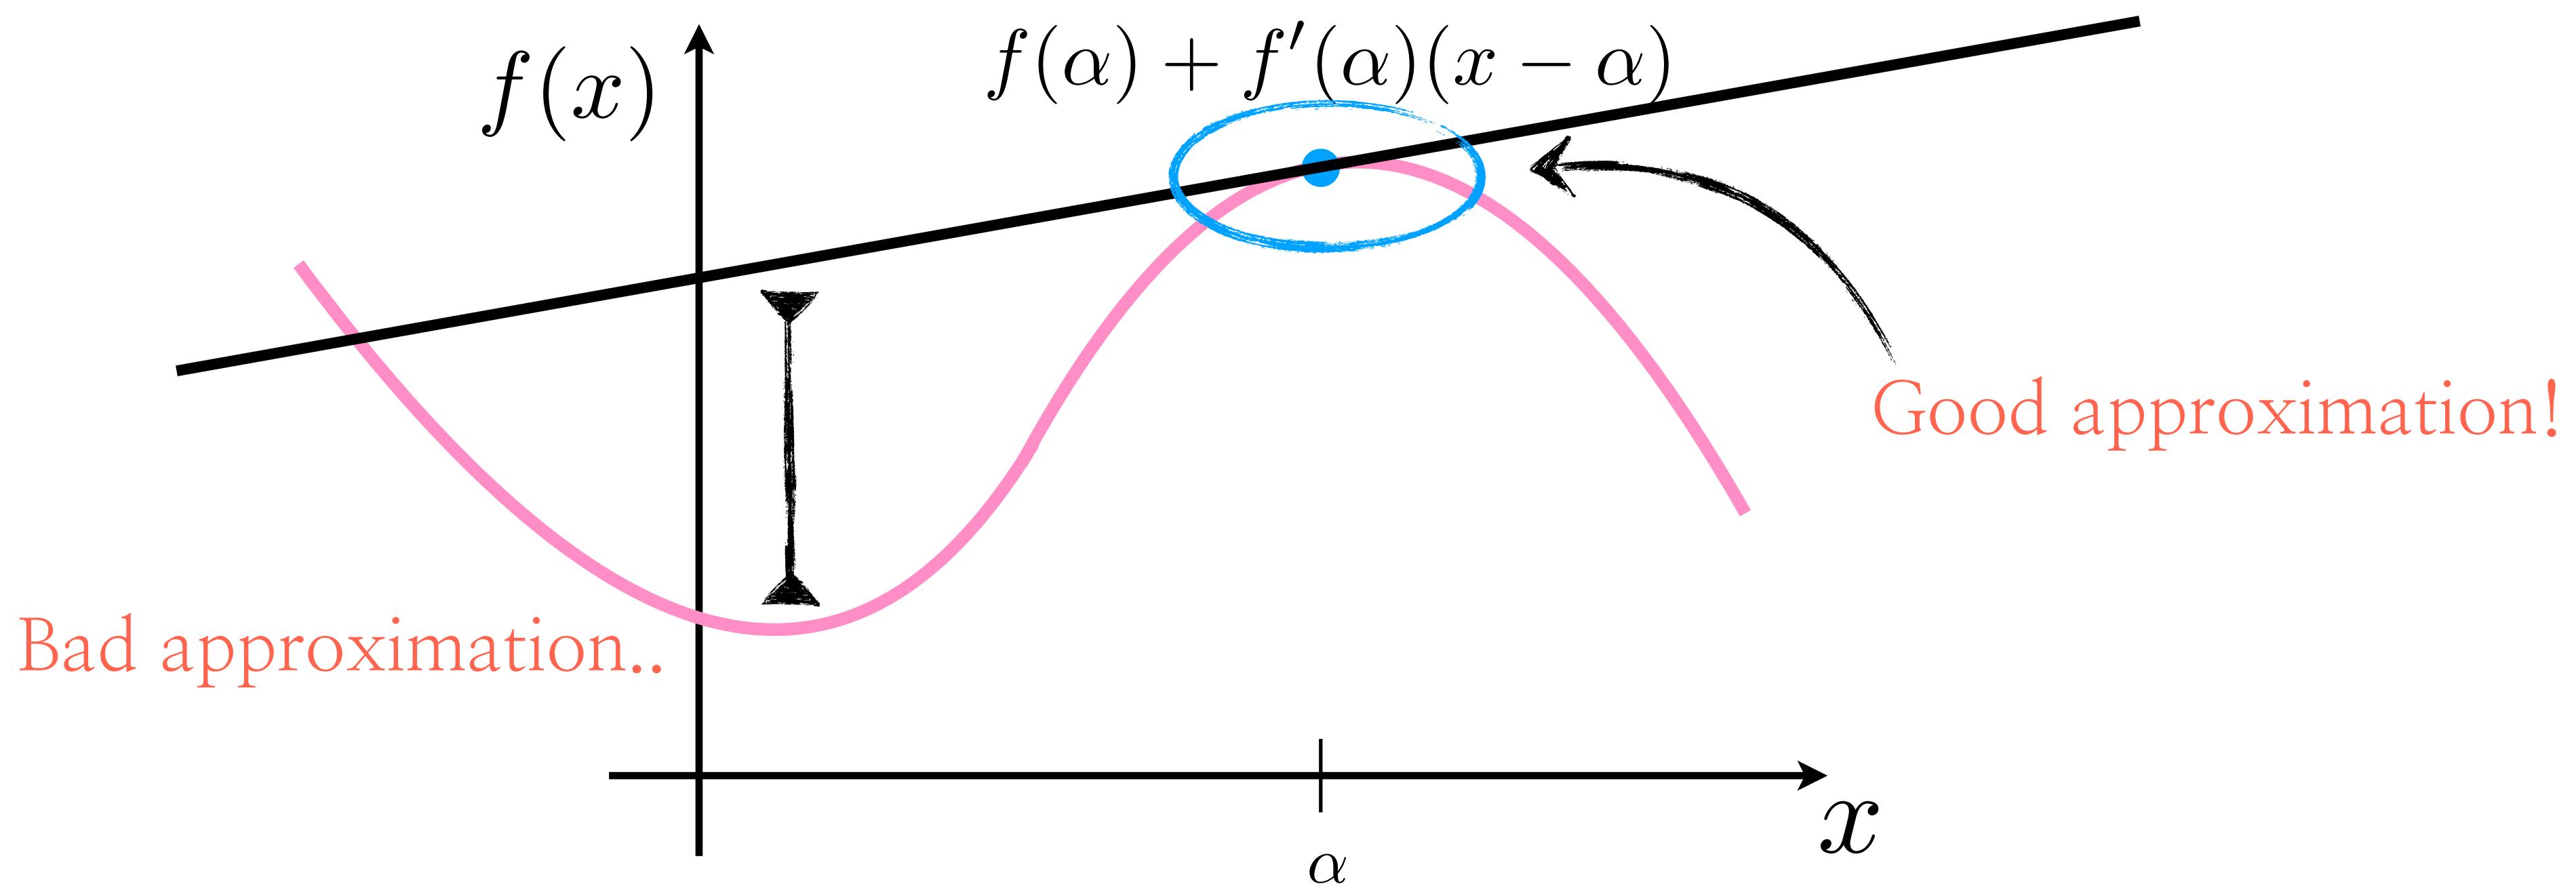
$$f : \mathbb{R}^p \rightarrow \mathbb{R} \quad | \quad f(x) \approx f(\alpha) + \langle \nabla f(\alpha), x - \alpha \rangle, \alpha \in \mathbb{R}^p$$



Taylor's expansion

- First-order Taylor's approximation

$$f : \mathbb{R}^p \rightarrow \mathbb{R} \quad | \quad f(x) \approx f(\alpha) + \langle \nabla f(\alpha), x - \alpha \rangle, \alpha \in \mathbb{R}^p$$



Taylor's expansion

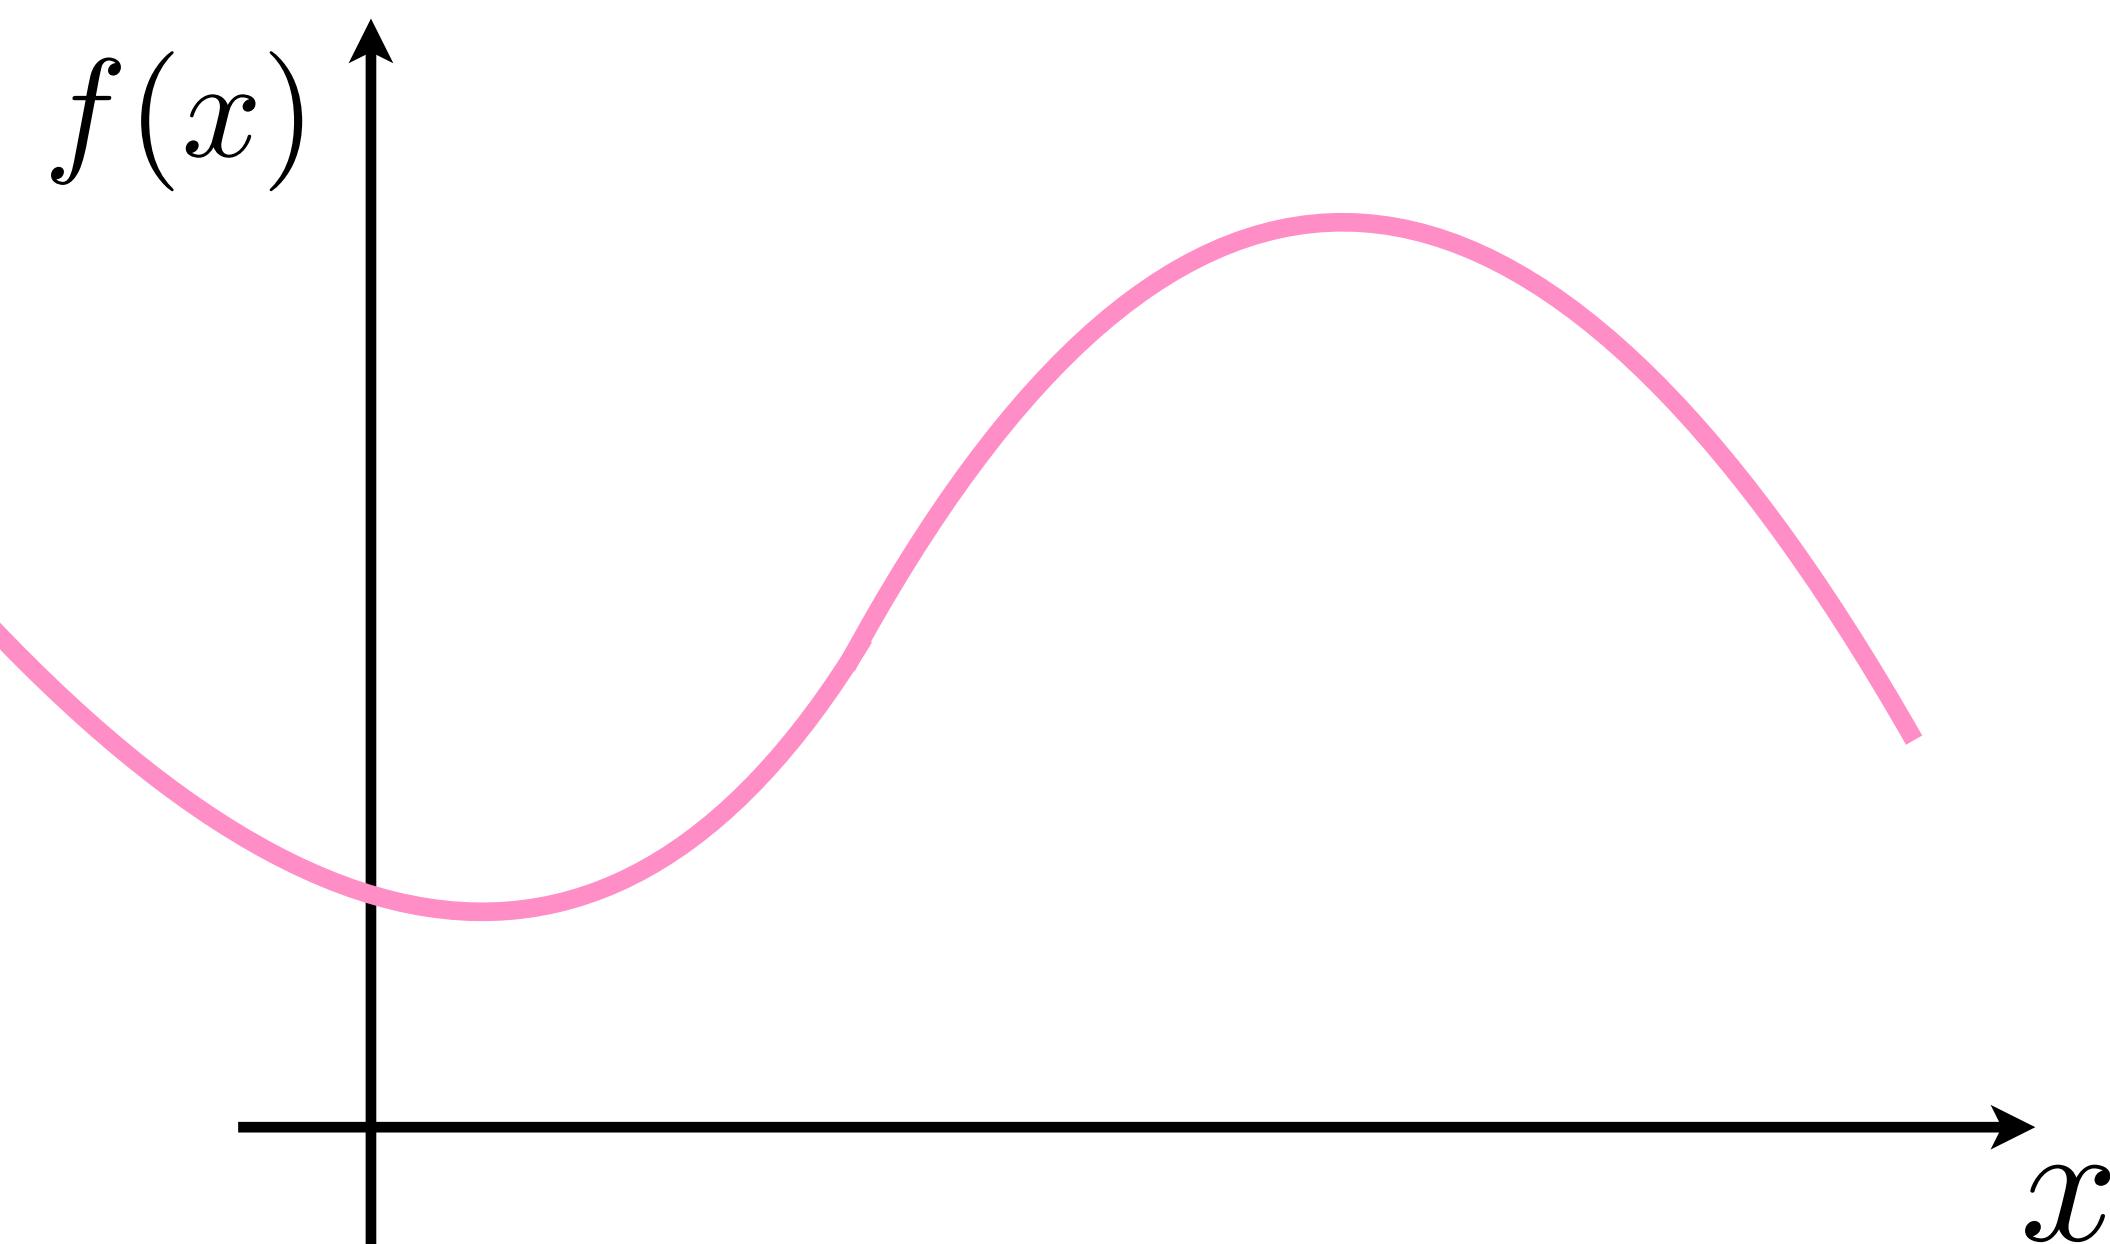
- Second-order Taylor's approximation

$$f : \mathbb{R}^p \rightarrow \mathbb{R} \quad \Big| \quad f(x) \approx f(\alpha) + \langle \nabla f(\alpha), x - \alpha \rangle + \frac{1}{2} \langle \nabla^2 f(\alpha)(x - \alpha), x - \alpha \rangle, \alpha \in \mathbb{R}^p$$

Taylor's expansion

- Second-order Taylor's approximation

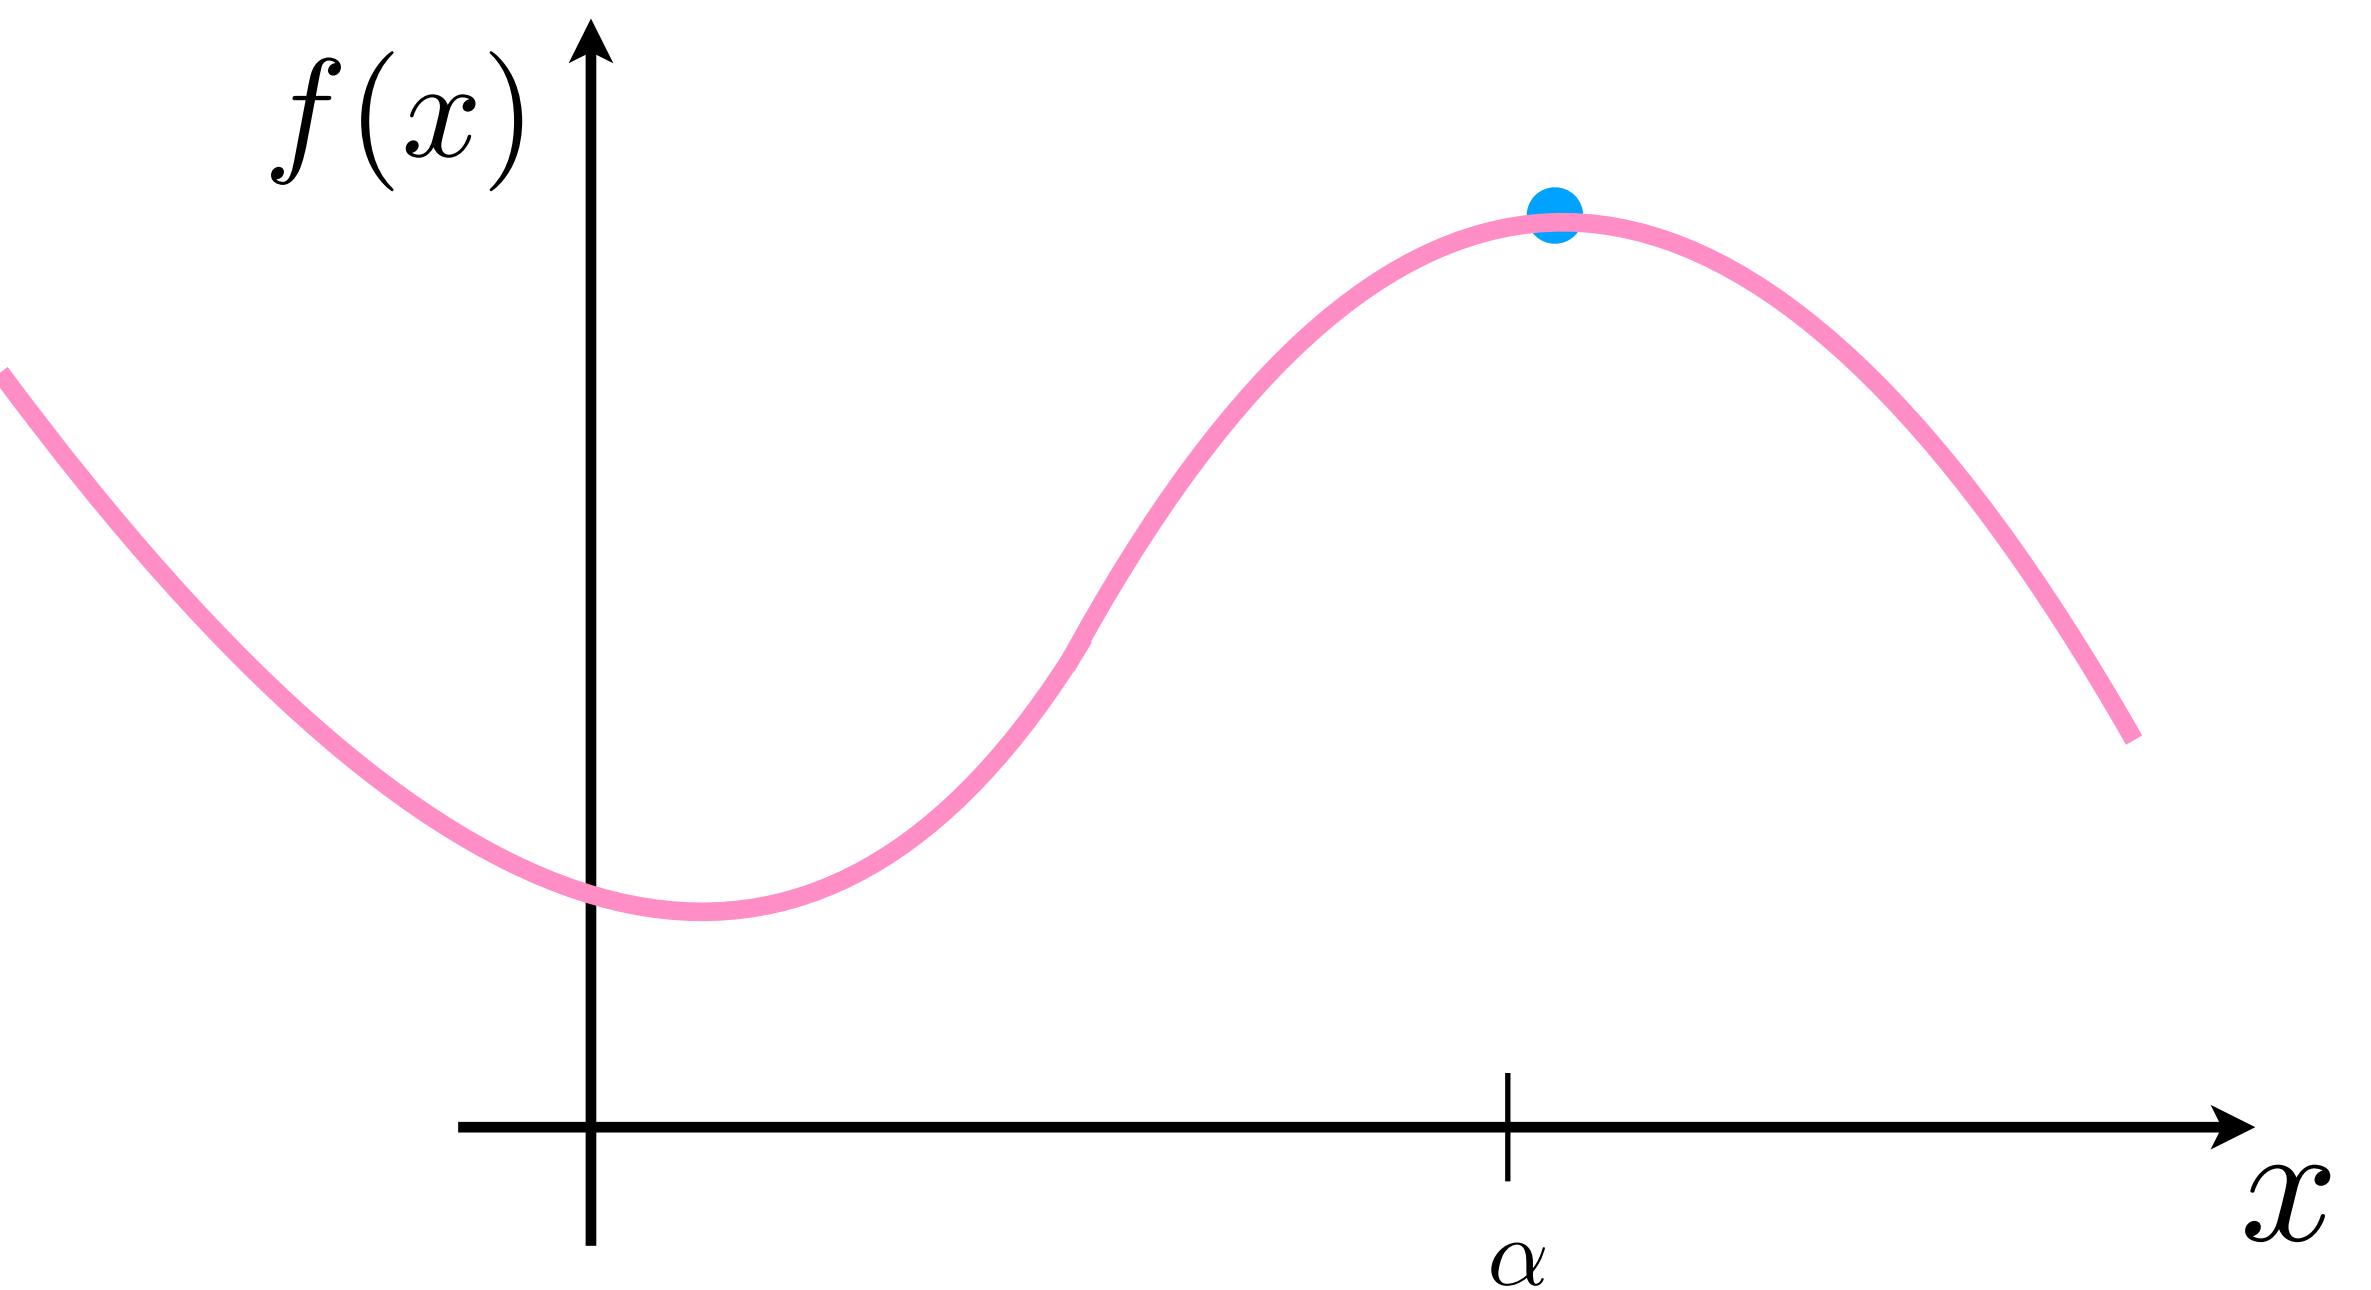
$$f : \mathbb{R}^p \rightarrow \mathbb{R} \quad | \quad f(x) \approx f(\alpha) + \langle \nabla f(\alpha), x - \alpha \rangle + \frac{1}{2} \langle \nabla^2 f(\alpha)(x - \alpha), x - \alpha \rangle, \alpha \in \mathbb{R}^p$$



Taylor's expansion

- Second-order Taylor's approximation

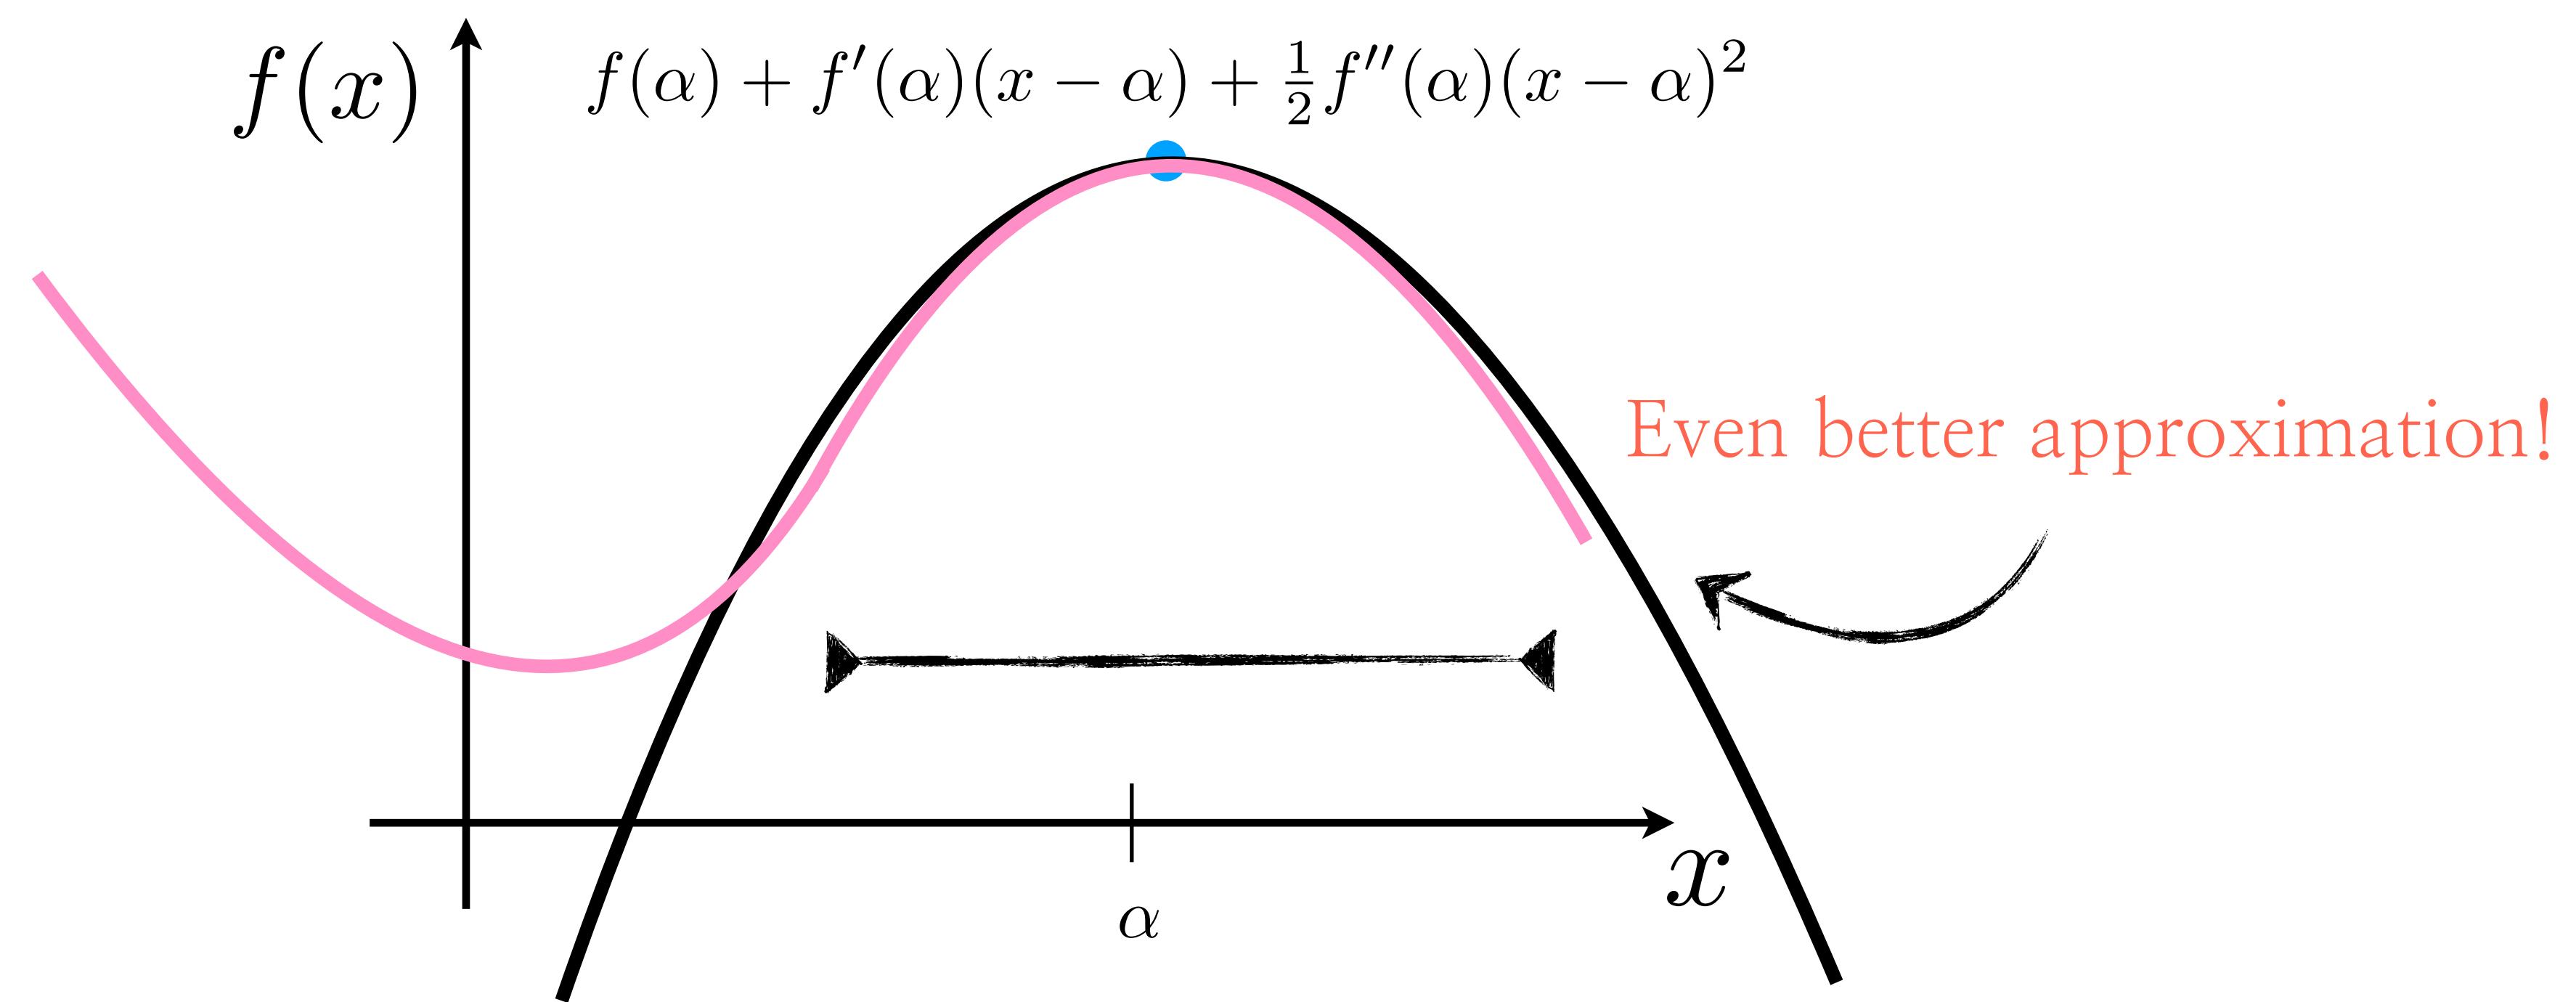
$$f : \mathbb{R}^p \rightarrow \mathbb{R} \quad | \quad f(x) \approx f(\alpha) + \langle \nabla f(\alpha), x - \alpha \rangle + \frac{1}{2} \langle \nabla^2 f(\alpha)(x - \alpha), x - \alpha \rangle, \alpha \in \mathbb{R}^p$$



Taylor's expansion

- Second-order Taylor's approximation

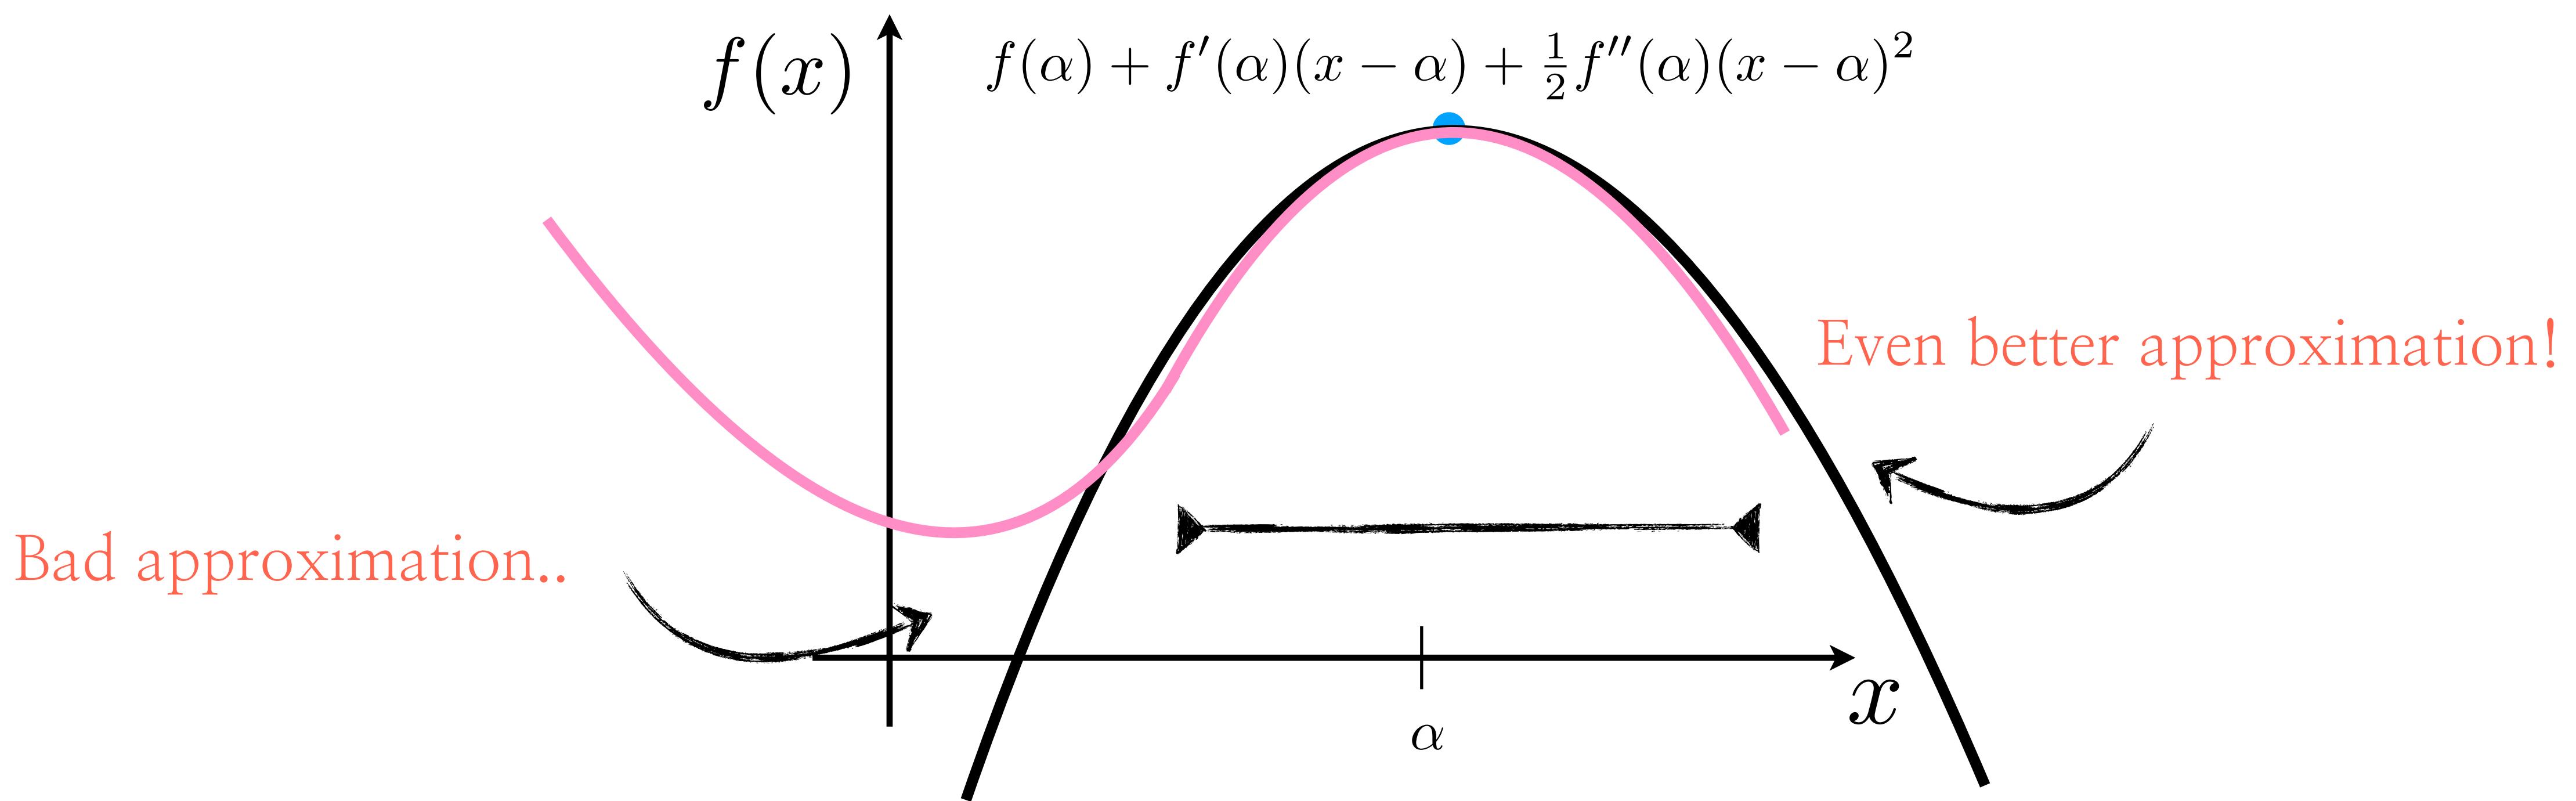
$$f : \mathbb{R}^p \rightarrow \mathbb{R} \quad | \quad f(x) \approx f(\alpha) + \langle \nabla f(\alpha), x - \alpha \rangle + \frac{1}{2} \langle \nabla^2 f(\alpha)(x - \alpha), x - \alpha \rangle, \alpha \in \mathbb{R}^p$$



Taylor's expansion

- Second-order Taylor's approximation

$$f : \mathbb{R}^p \rightarrow \mathbb{R} \quad | \quad f(x) \approx f(\alpha) + \langle \nabla f(\alpha), x - \alpha \rangle + \frac{1}{2} \langle \nabla^2 f(\alpha)(x - \alpha), x - \alpha \rangle, \alpha \in \mathbb{R}^p$$



Taylor's expansion

- Spoiler alert: “Why are all these useful?”
 - Often, we optimize a function through its local approximations
 - E.g., second order approximations are.. quadratic functions!

Taylor's expansion

- Spoiler alert: “Why are all these useful?”
 - Often, we optimize a function through its local approximations
 - E.g., second order approximations are.. quadratic functions!

$$\min_x f(x)$$

Taylor's expansion

- Spoiler alert: “Why are all these useful?”
 - Often, we optimize a function through its local approximations
 - E.g., second order approximations are.. quadratic functions!

Weirdest function ever
(but differentiable)

$$\min_x f(x)$$


Taylor's expansion

- Spoiler alert: “Why are all these useful?”
 - Often, we optimize a function through its local approximations
 - E.g., second order approximations are.. quadratic functions!

Weirdest function ever
(but differentiable)

$$\min_x f(x) \longrightarrow \min_x \left\{ f(x_0) + \nabla f(x_0)^\top (x - x_0) + \frac{1}{2}(x - x_0)^\top \nabla f(x_0)(x - x_0) \right\}$$

Taylor's expansion

- Spoiler alert: “Why are all these useful?”
 - Often, we optimize a function through its local approximations
 - E.g., second order approximations are.. quadratic functions!

Weirdest function ever
(but differentiable)

$$\min_x f(x) \longrightarrow \min_x \left\{ p^\top x + \frac{1}{2} x^\top H x \right\}$$

Taylor's expansion

- Spoiler alert: “Why are all these useful?”
 - Often, we optimize a function through its local approximations
 - E.g., second order approximations are.. quadratic functions!

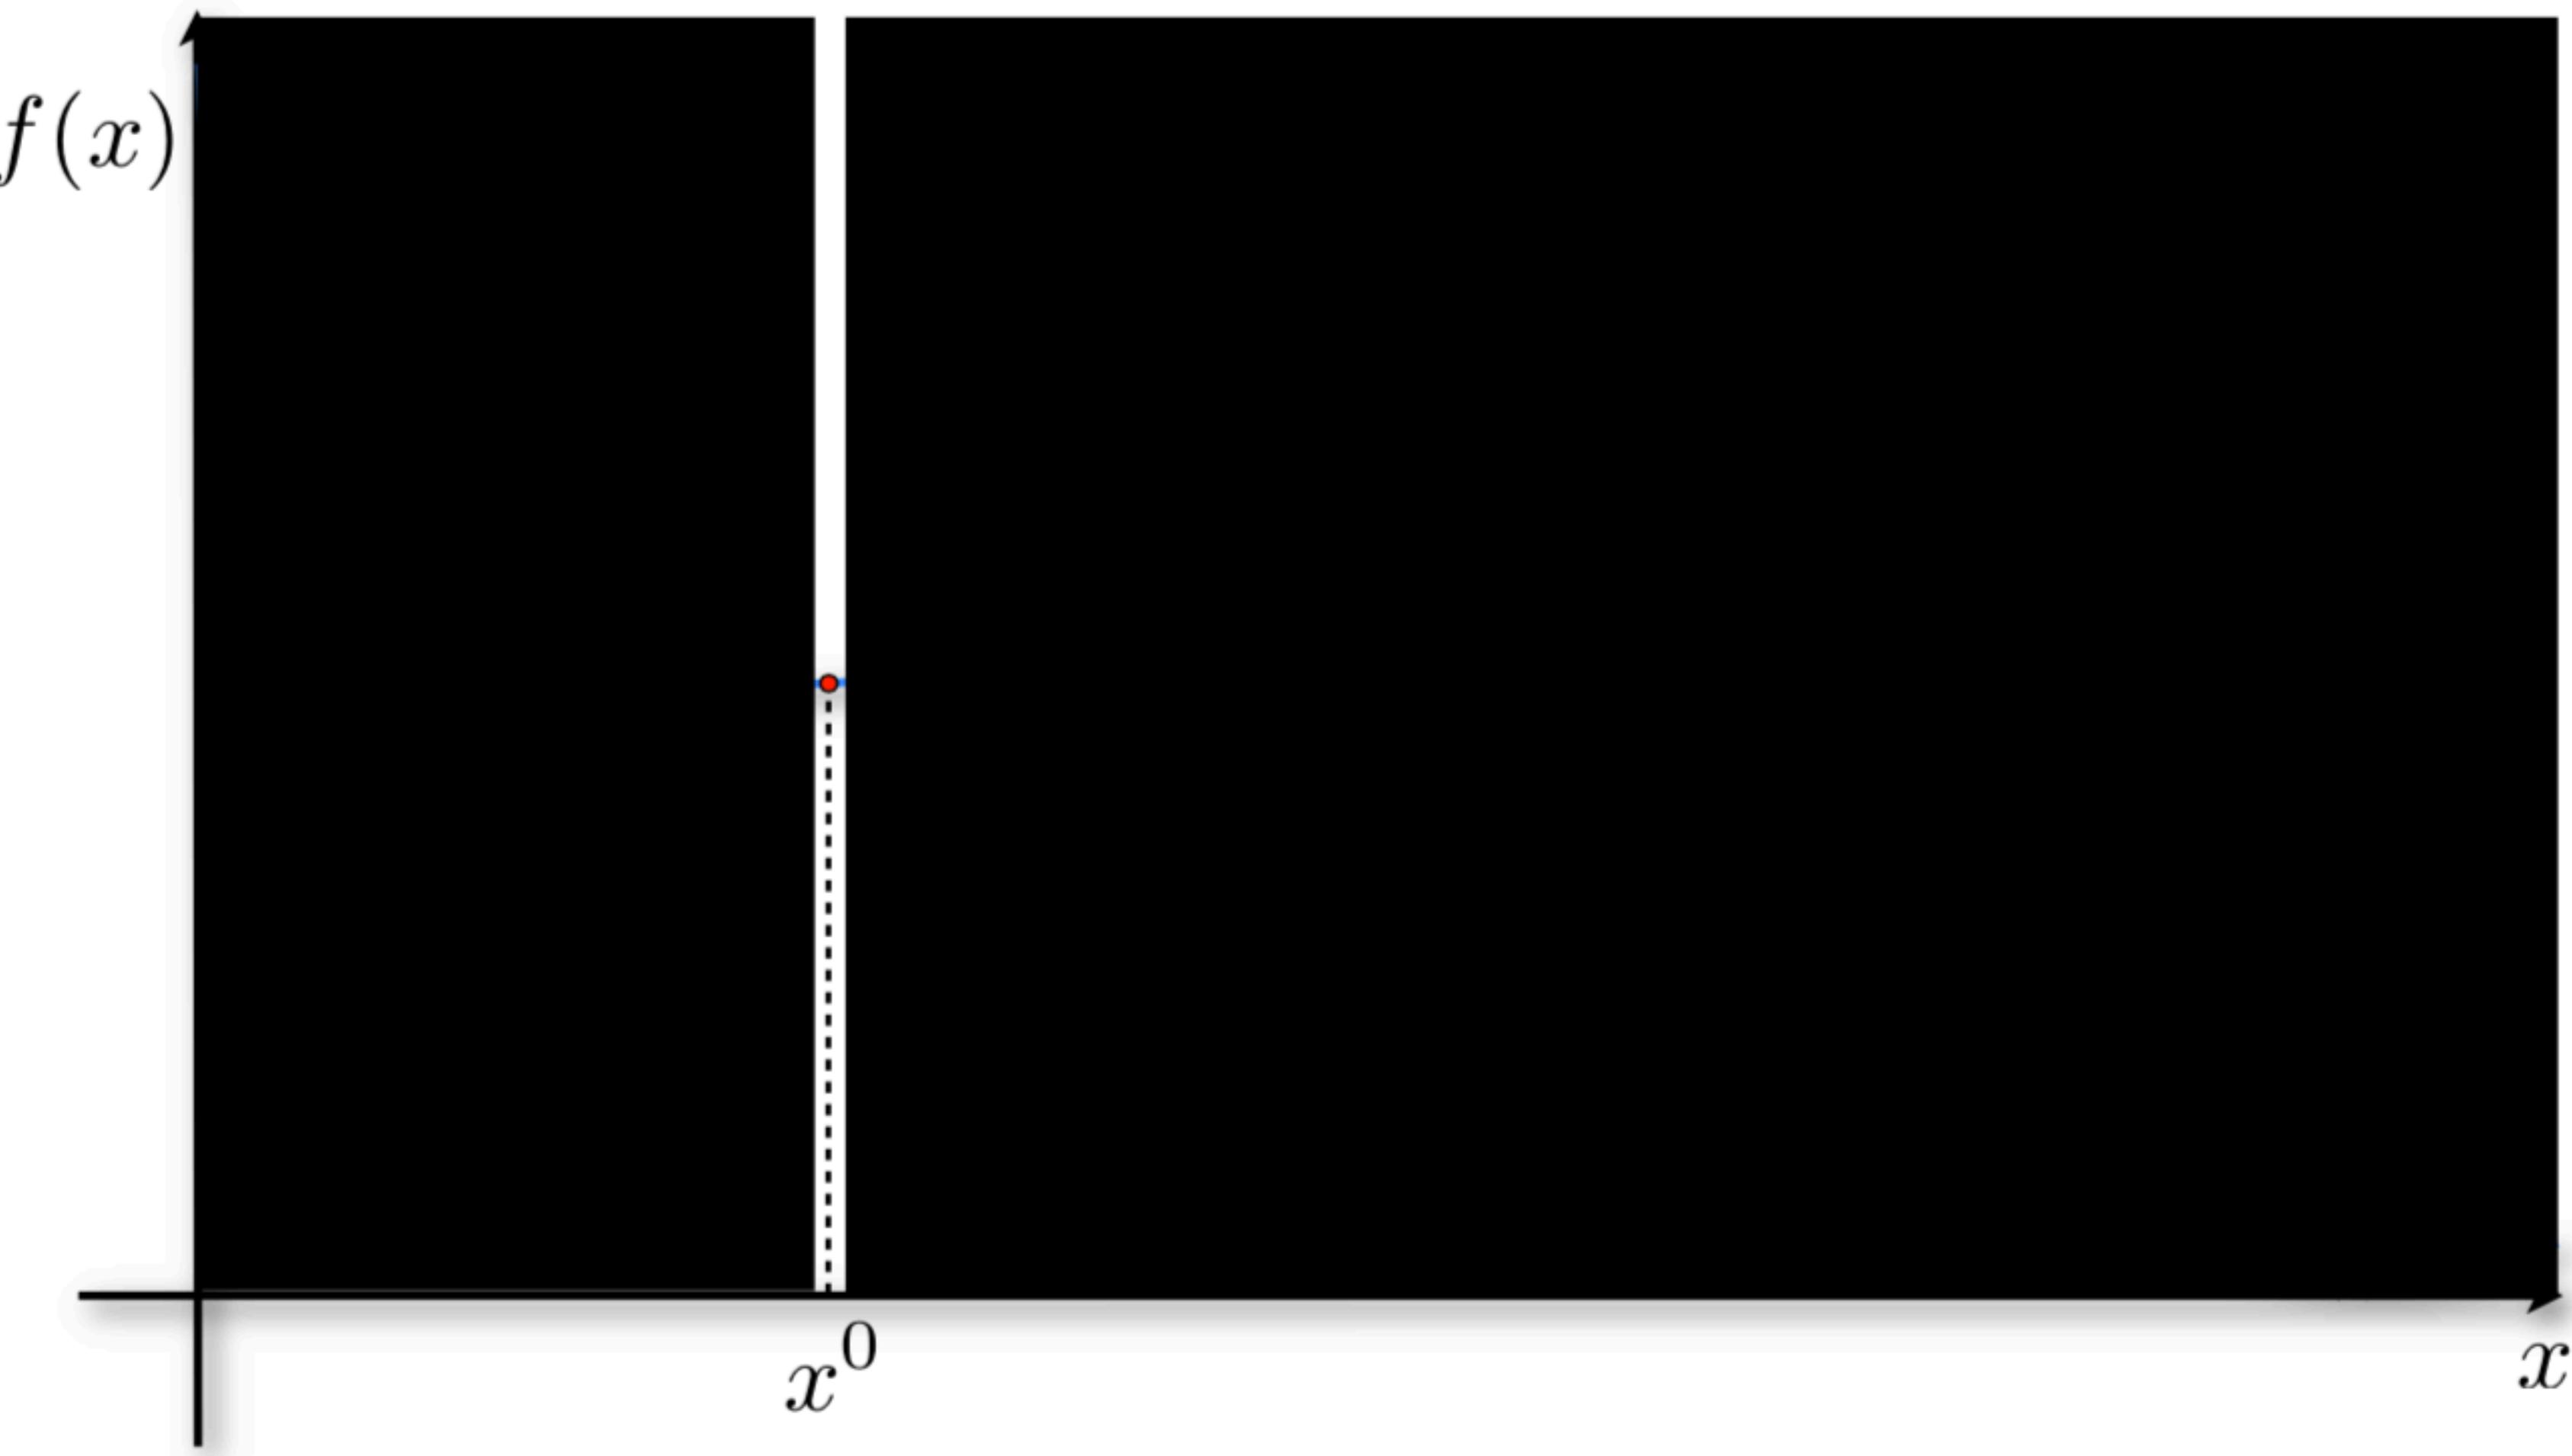
Weirdest function ever
(but differentiable)

$$\min_x f(x) \xrightarrow{\quad} \min_x \left\{ p^\top x + \frac{1}{2} x^\top H x \right\} \xrightarrow{\quad} \dots$$

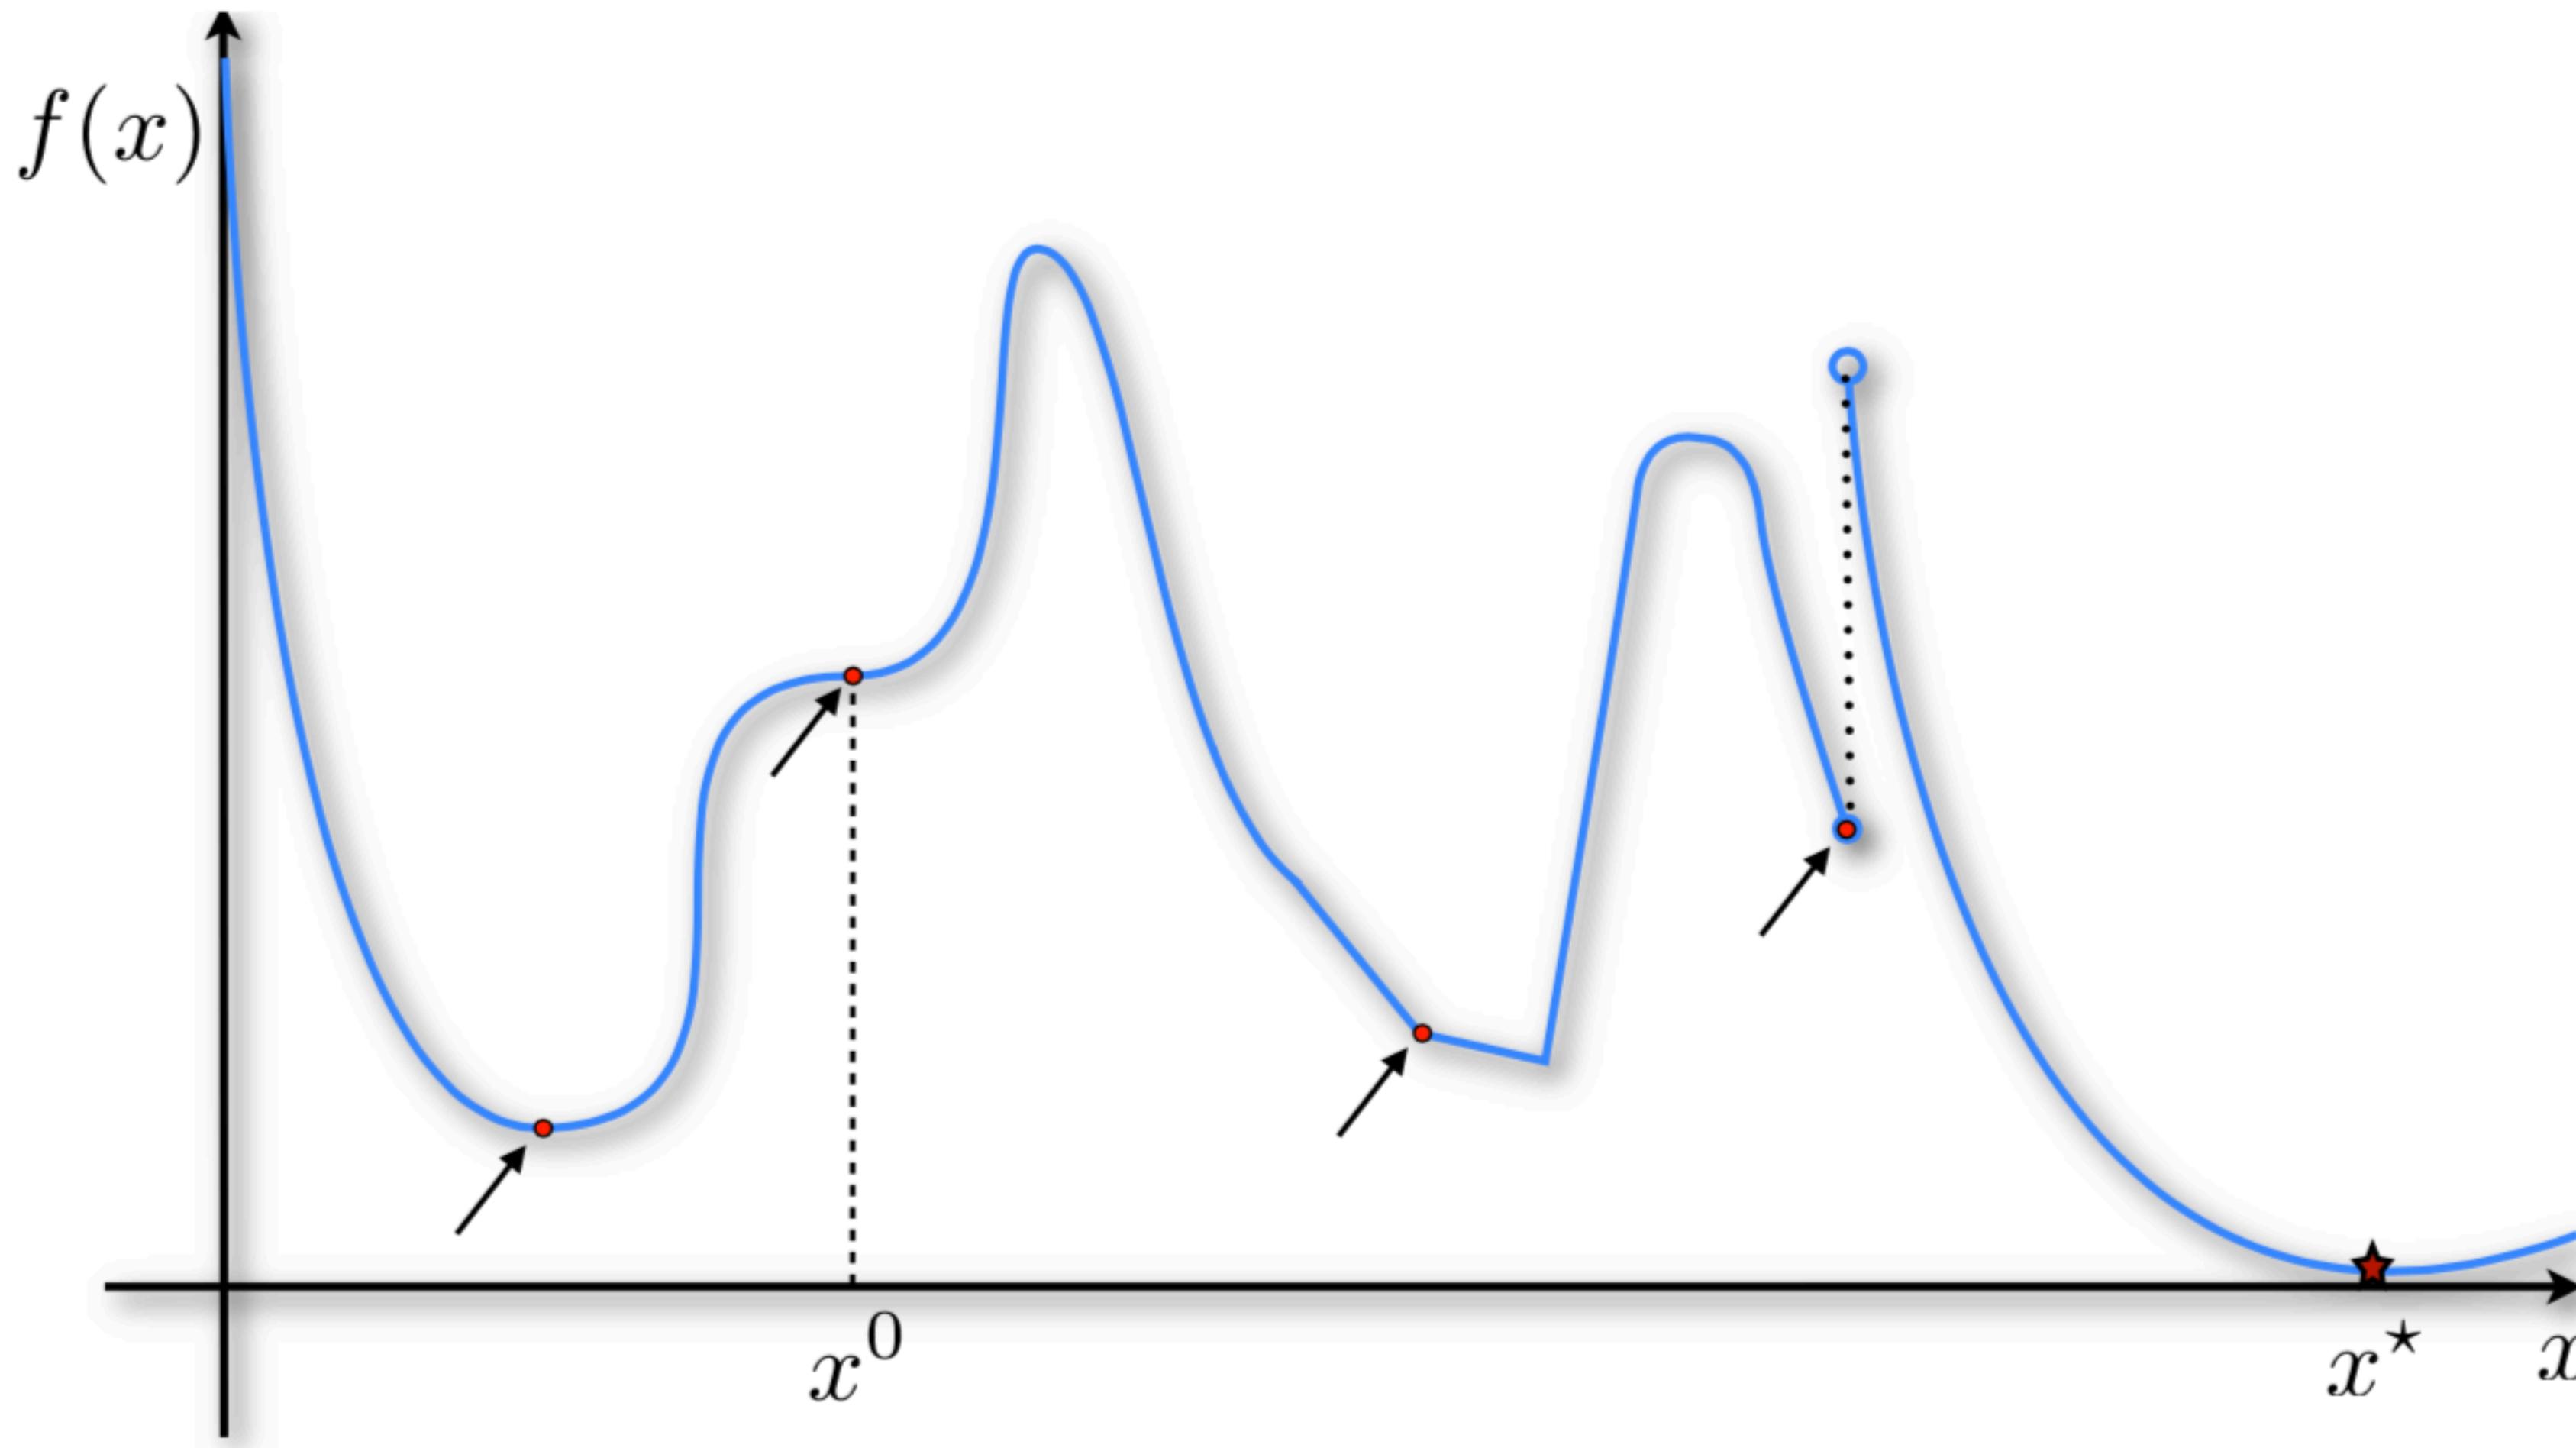
Agnostic optimization

Demo

Agnostic optimization



Agnostic optimization

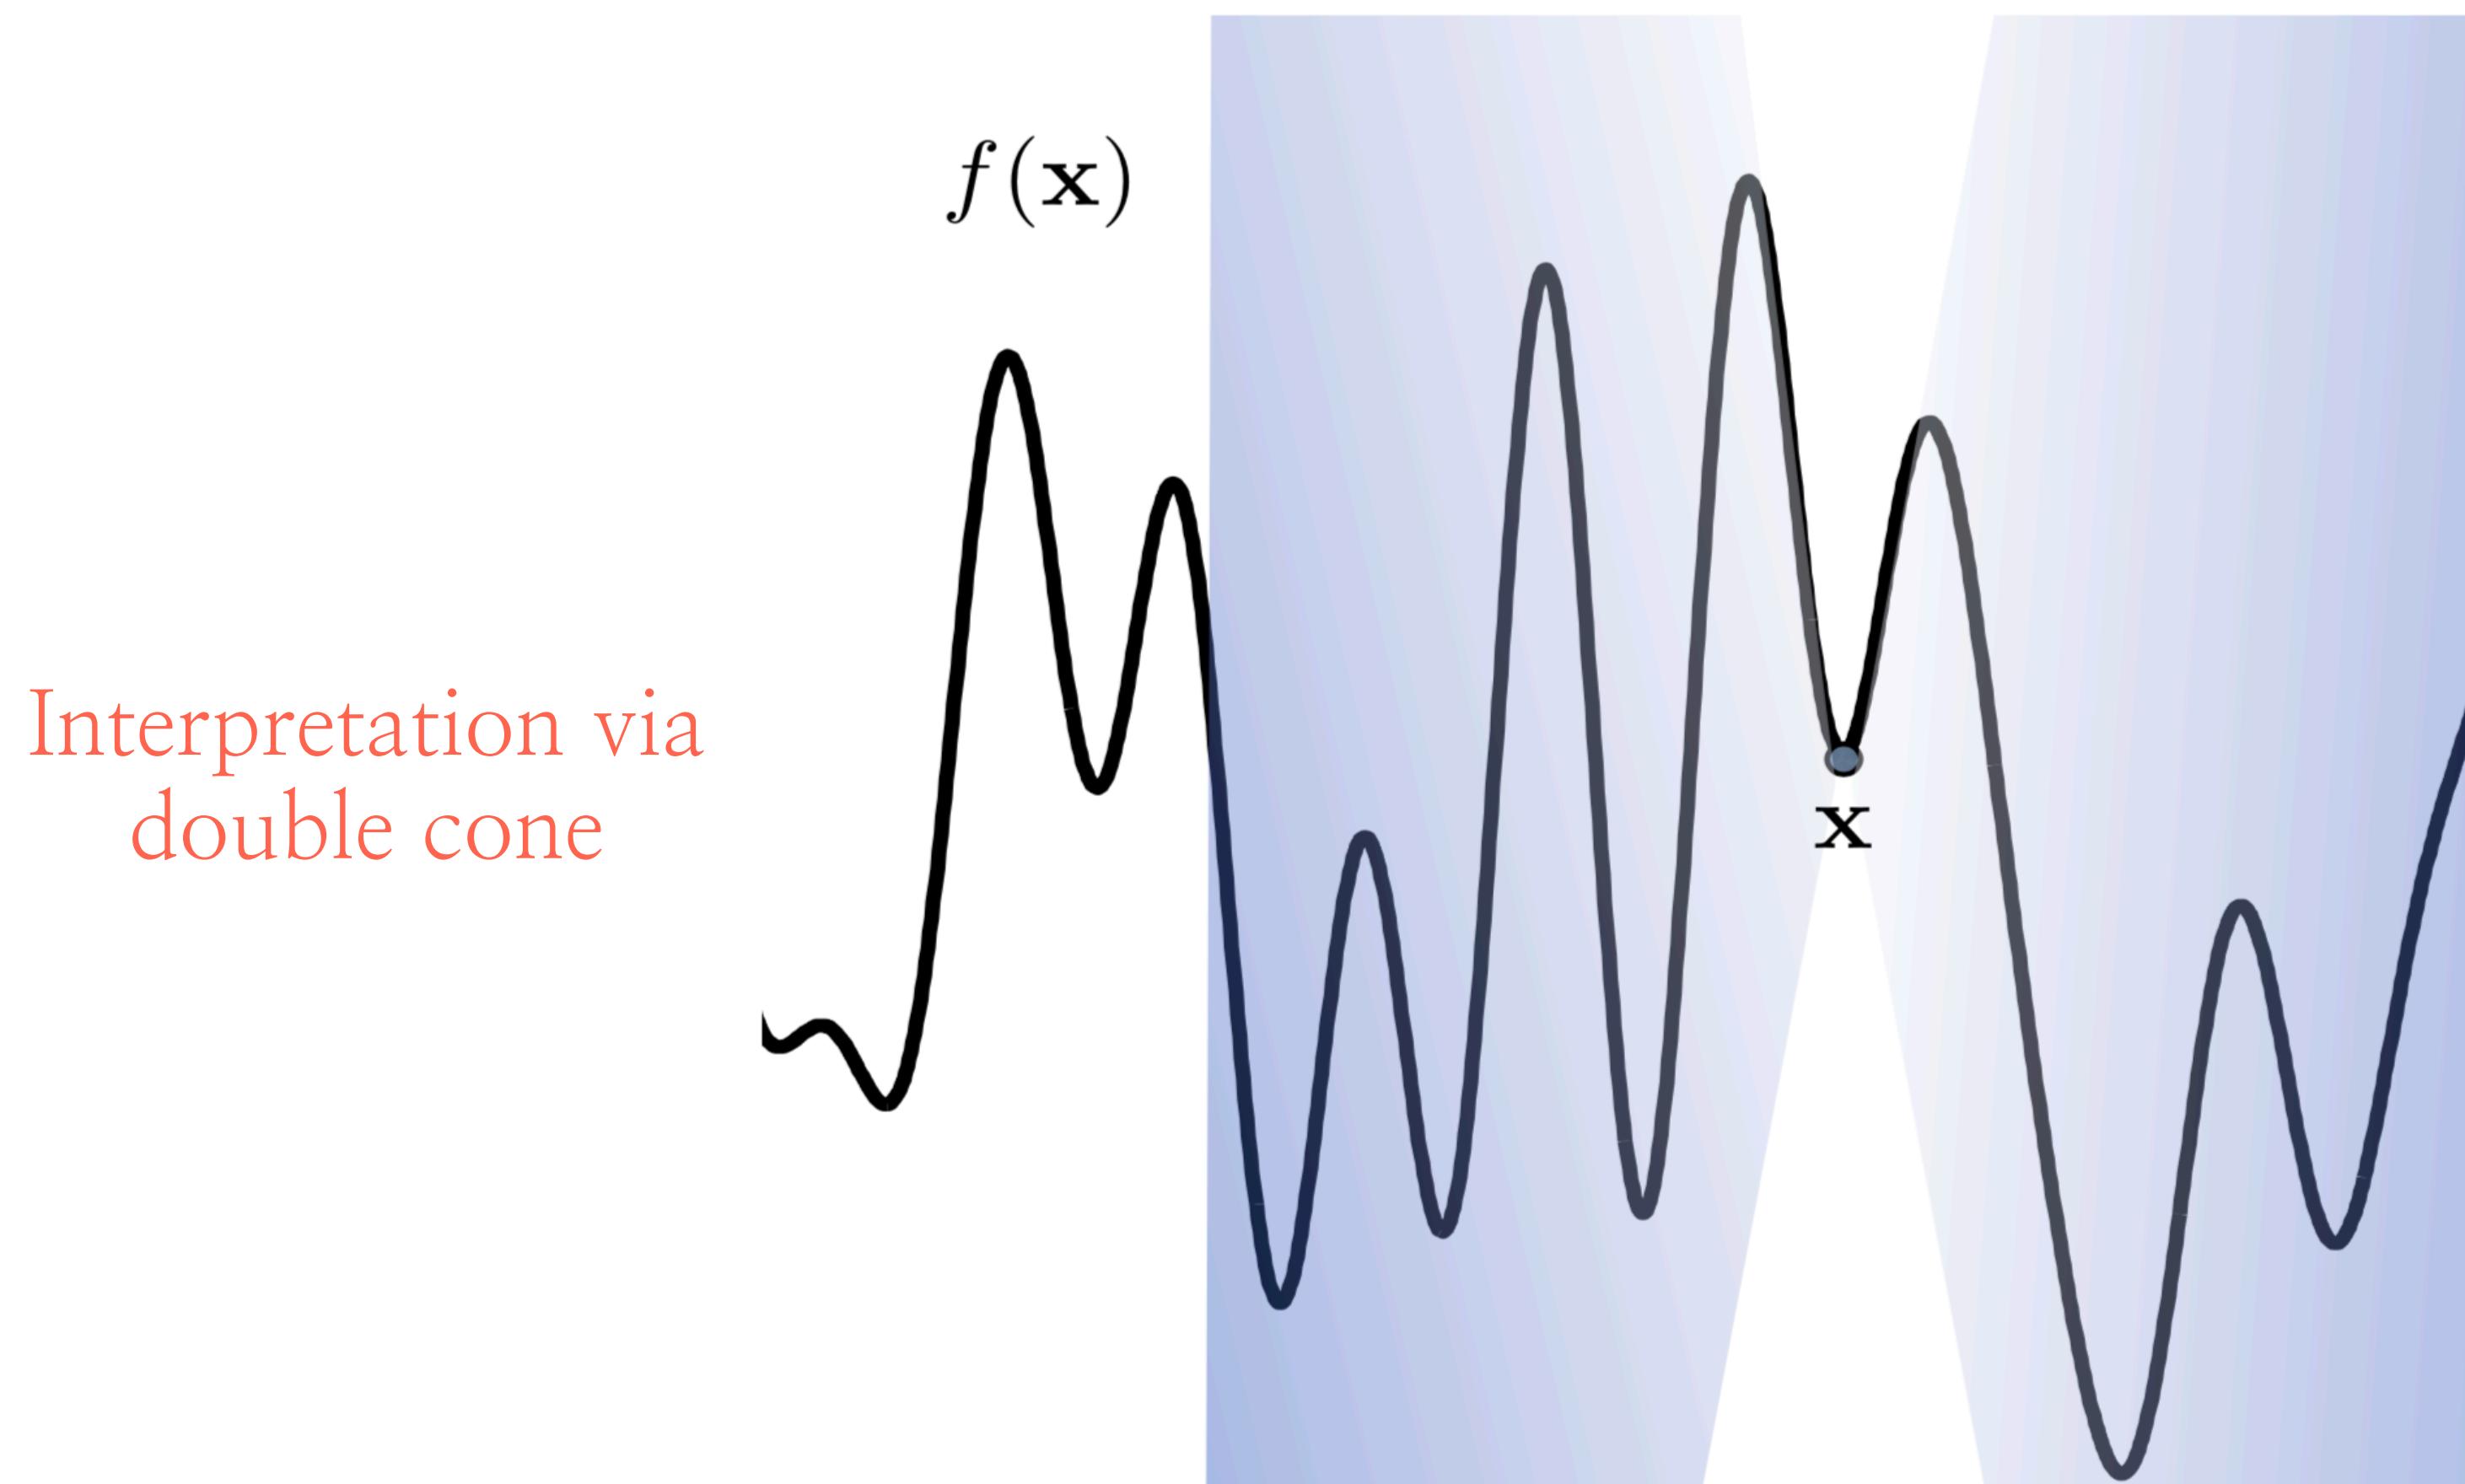


Lipschitz conditions

– Lipschitz continuity: $|f(x) - f(y)| \leq M\|x - y\|_2, \quad \forall x, y$

Lipschitz conditions

- Lipschitz continuity: $|f(x) - f(y)| \leq M\|x - y\|_2, \quad \forall x, y$



- Function examples:
1. Absolute value
 2. Trigonometric functions
 3. Quadratics (..)

Lipschitz conditions

- Lipschitz gradient continuity: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$

Lipschitz conditions

- Lipschitz gradient continuity: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$
- Intuition + comparison with Lipschitz continuity:

“Lipschitz continuity implies that f should not be too steep”

“Lipschitz gradient continuity implies that changes in the slope of f should not happen suddenly”

Lipschitz conditions

- Lipschitz gradient continuity: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$
- Intuition + comparison with Lipschitz continuity:

“Lipschitz continuity implies that f should not be too steep”

“Lipschitz gradient continuity implies that changes in the slope of f should not happen suddenly”

- Example: Quadratics are not globally Lipschitz continuous
(the function becomes arbitrarily steep as we approach infinity)

but: $\|\nabla f(x) - \nabla f(y)\|_2 \leq \|A^\top A\|_2 \cdot \|x - y\|_2$

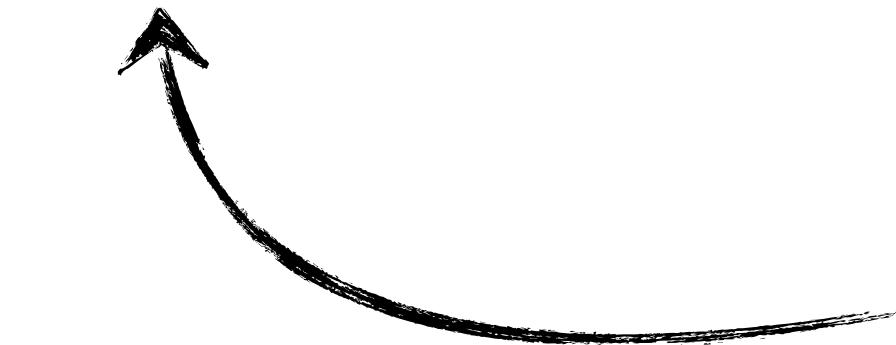
for: $f(x) = \frac{1}{2}\|Ax - b\|_2^2$

Lipschitz conditions

- Lipschitz gradient continuity: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$
- Intuition + comparison with Lipschitz continuity:
 - “Lipschitz continuity implies that f should not be too steep”
 - “Lipschitz gradient continuity implies that changes in the slope of f should not happen suddenly”
- Example: Quadratics are not globally Lipschitz continuous
 - (the function becomes arbitrarily steep as we approach infinity)

but: $\|\nabla f(x) - \nabla f(y)\|_2 \leq \|A^\top A\|_2 \cdot \|x - y\|_2$

for: $f(x) = \frac{1}{2}\|Ax - b\|_2^2$



Largest singular value

Lipschitz conditions

- Equivalent characterizations: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$

$$f(y) \leq f(x) + \langle \nabla f(x), y - x \rangle + \frac{L}{2} \|x - y\|_2^2$$

$$f(y) \geq f(x) + \langle \nabla f(x), y - x \rangle + \frac{1}{2L} \|\nabla f(x) - \nabla f(y)\|_2^2$$

$$\langle \nabla f(x) - \nabla f(y), x - y \rangle \leq L\|x - y\|_2^2$$

⋮ ⋮

Lipschitz conditions

- Equivalent characterizations: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$

$$f(y) \leq f(x) + \langle \nabla f(x), y - x \rangle + \frac{L}{2}\|x - y\|_2^2$$

$$f(y) \geq f(x) + \langle \nabla f(x), y - x \rangle + \frac{1}{2L}\|\nabla f(x) - \nabla f(y)\|_2^2$$

$$\langle \nabla f(x) - \nabla f(y), x - y \rangle \leq L\|x - y\|_2^2$$

$$\begin{matrix} \vdots & \vdots \end{matrix}$$

- Another important one:

$$\nabla^2 f(x) \preceq L I$$

Lipschitz conditions

- Equivalent characterizations: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$

$$f(y) \leq f(x) + \langle \nabla f(x), y - x \rangle + \frac{L}{2} \|x - y\|_2^2$$

$$f(y) \geq f(x) + \langle \nabla f(x), y - x \rangle + \frac{1}{2L} \|\nabla f(x) - \nabla f(y)\|_2^2$$

$$\langle \nabla f(x) - \nabla f(y), x - y \rangle \leq L\|x - y\|_2^2$$

$$\begin{matrix} \vdots & \vdots \end{matrix}$$

- Another important one:

$$\nabla^2 f(x) \preceq L I$$

Interpretation?

Lipschitz conditions

- Equivalent characterizations: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$

$$f(y) \leq f(x) + \langle \nabla f(x), y - x \rangle + \frac{L}{2} \|x - y\|_2^2$$

$$f(y) \geq f(x) + \langle \nabla f(x), y - x \rangle + \frac{1}{2L} \|\nabla f(x) - \nabla f(y)\|_2^2$$

$$\langle \nabla f(x) - \nabla f(y), x - y \rangle \leq L\|x - y\|_2^2$$

$$\begin{matrix} \vdots & \vdots \end{matrix}$$

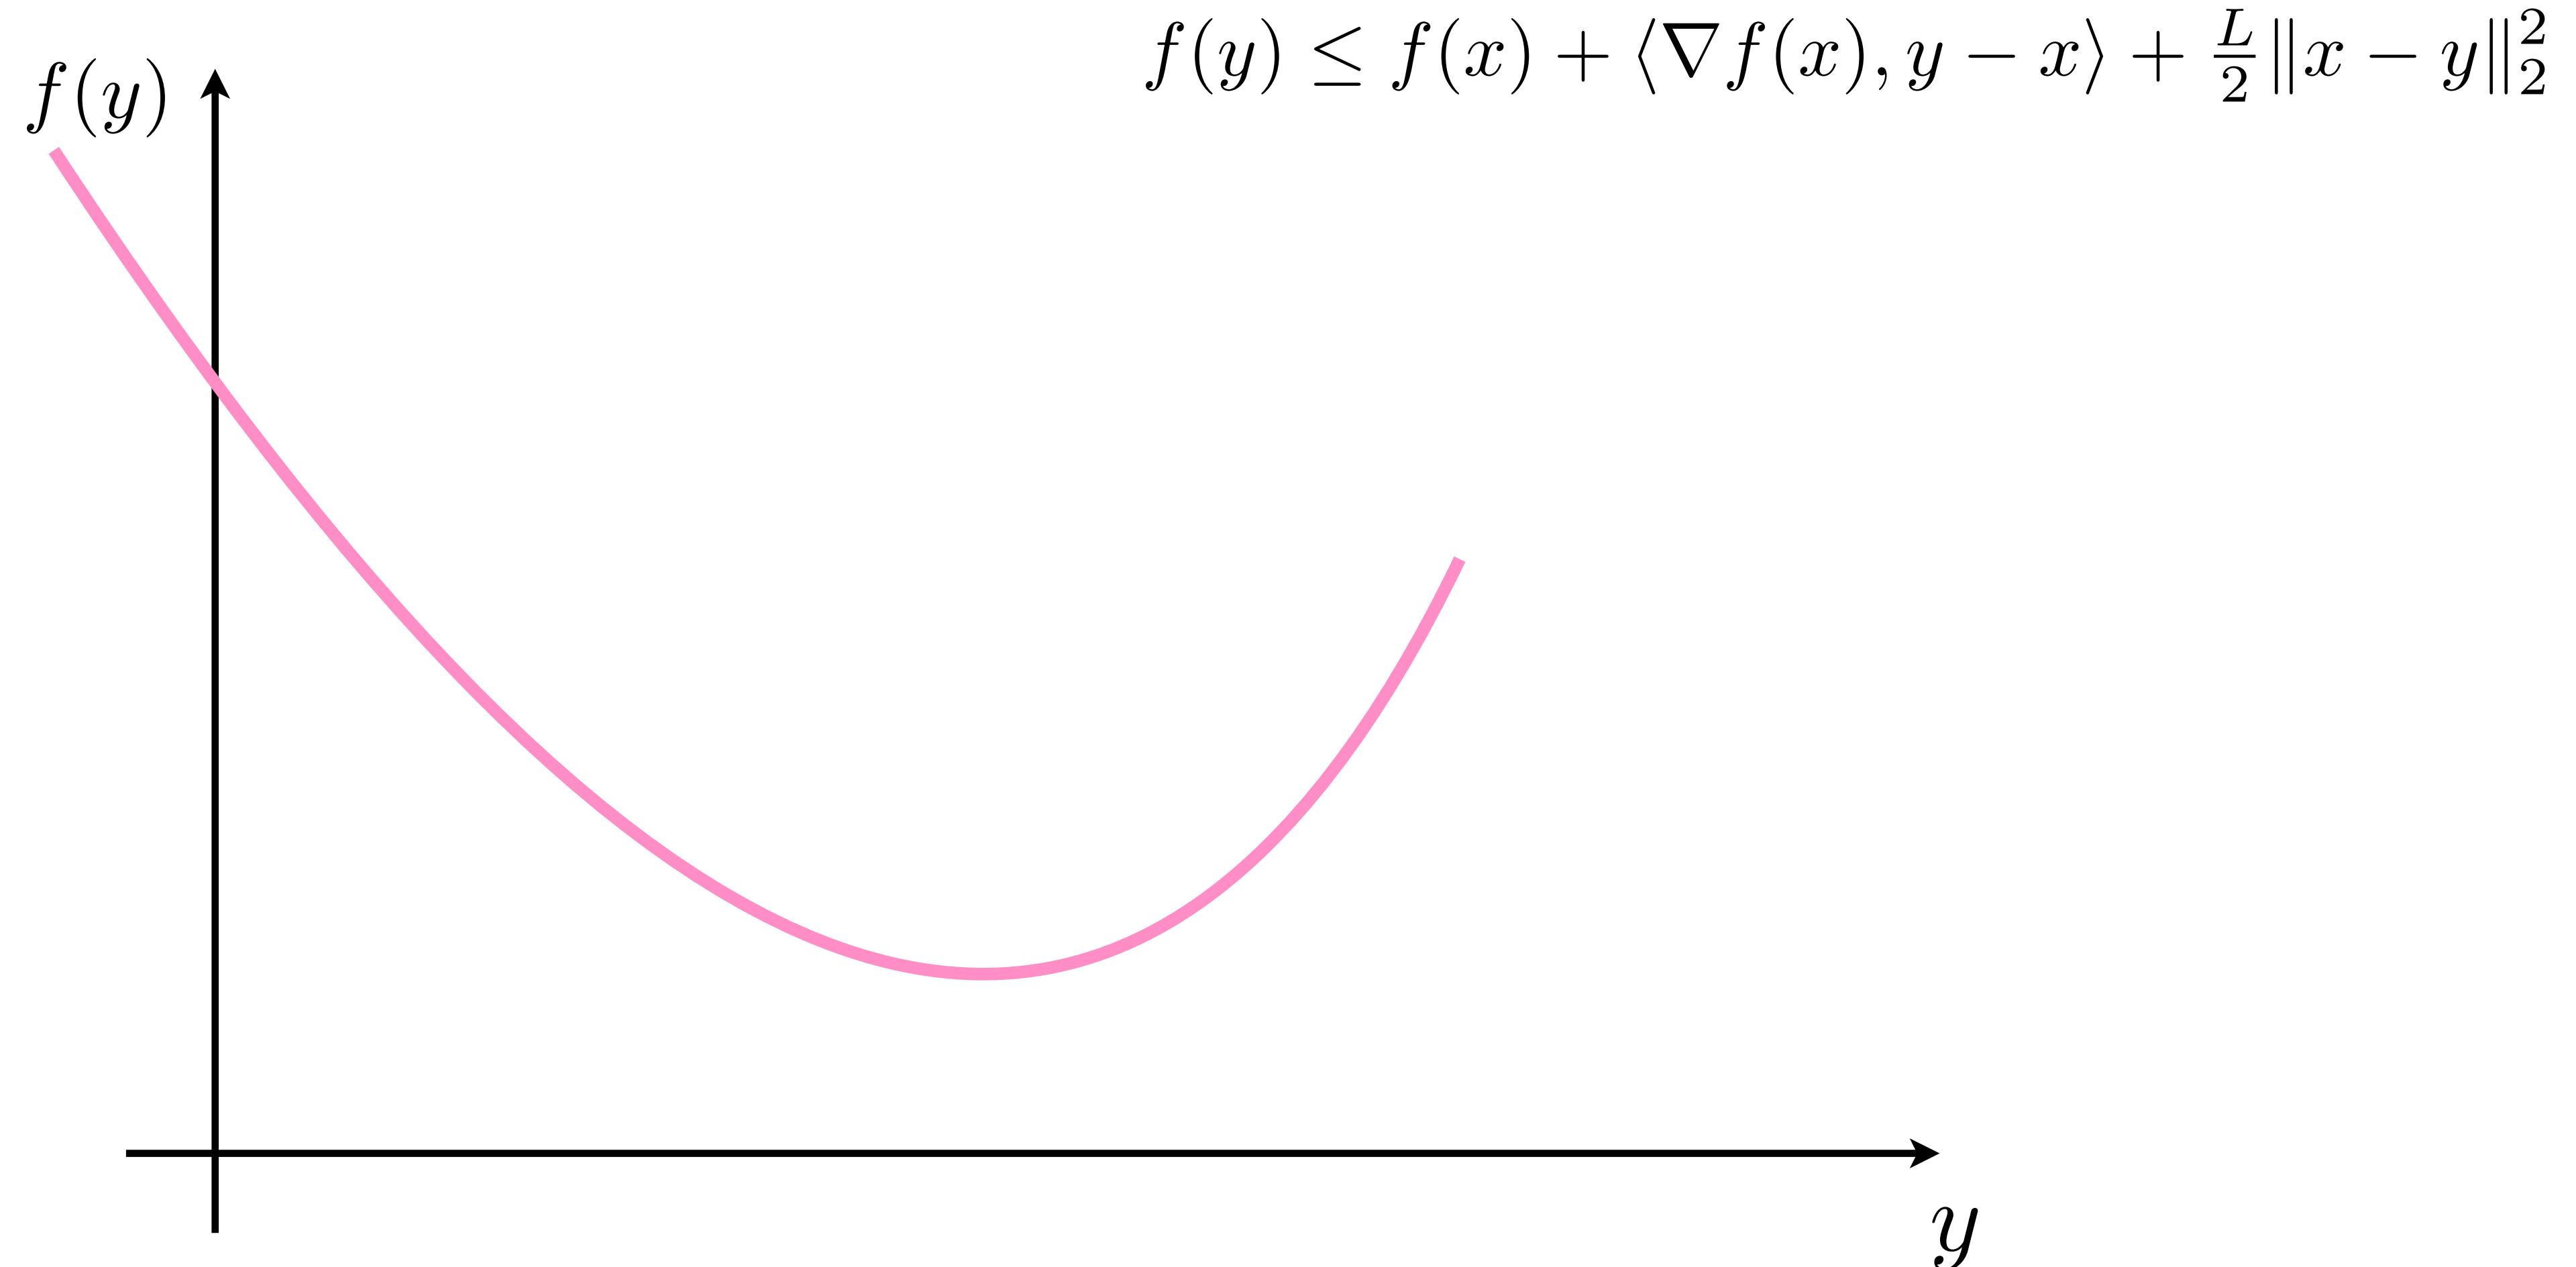
- Another important one:

$$\nabla^2 f(x) \preceq L I$$

Interpretation?

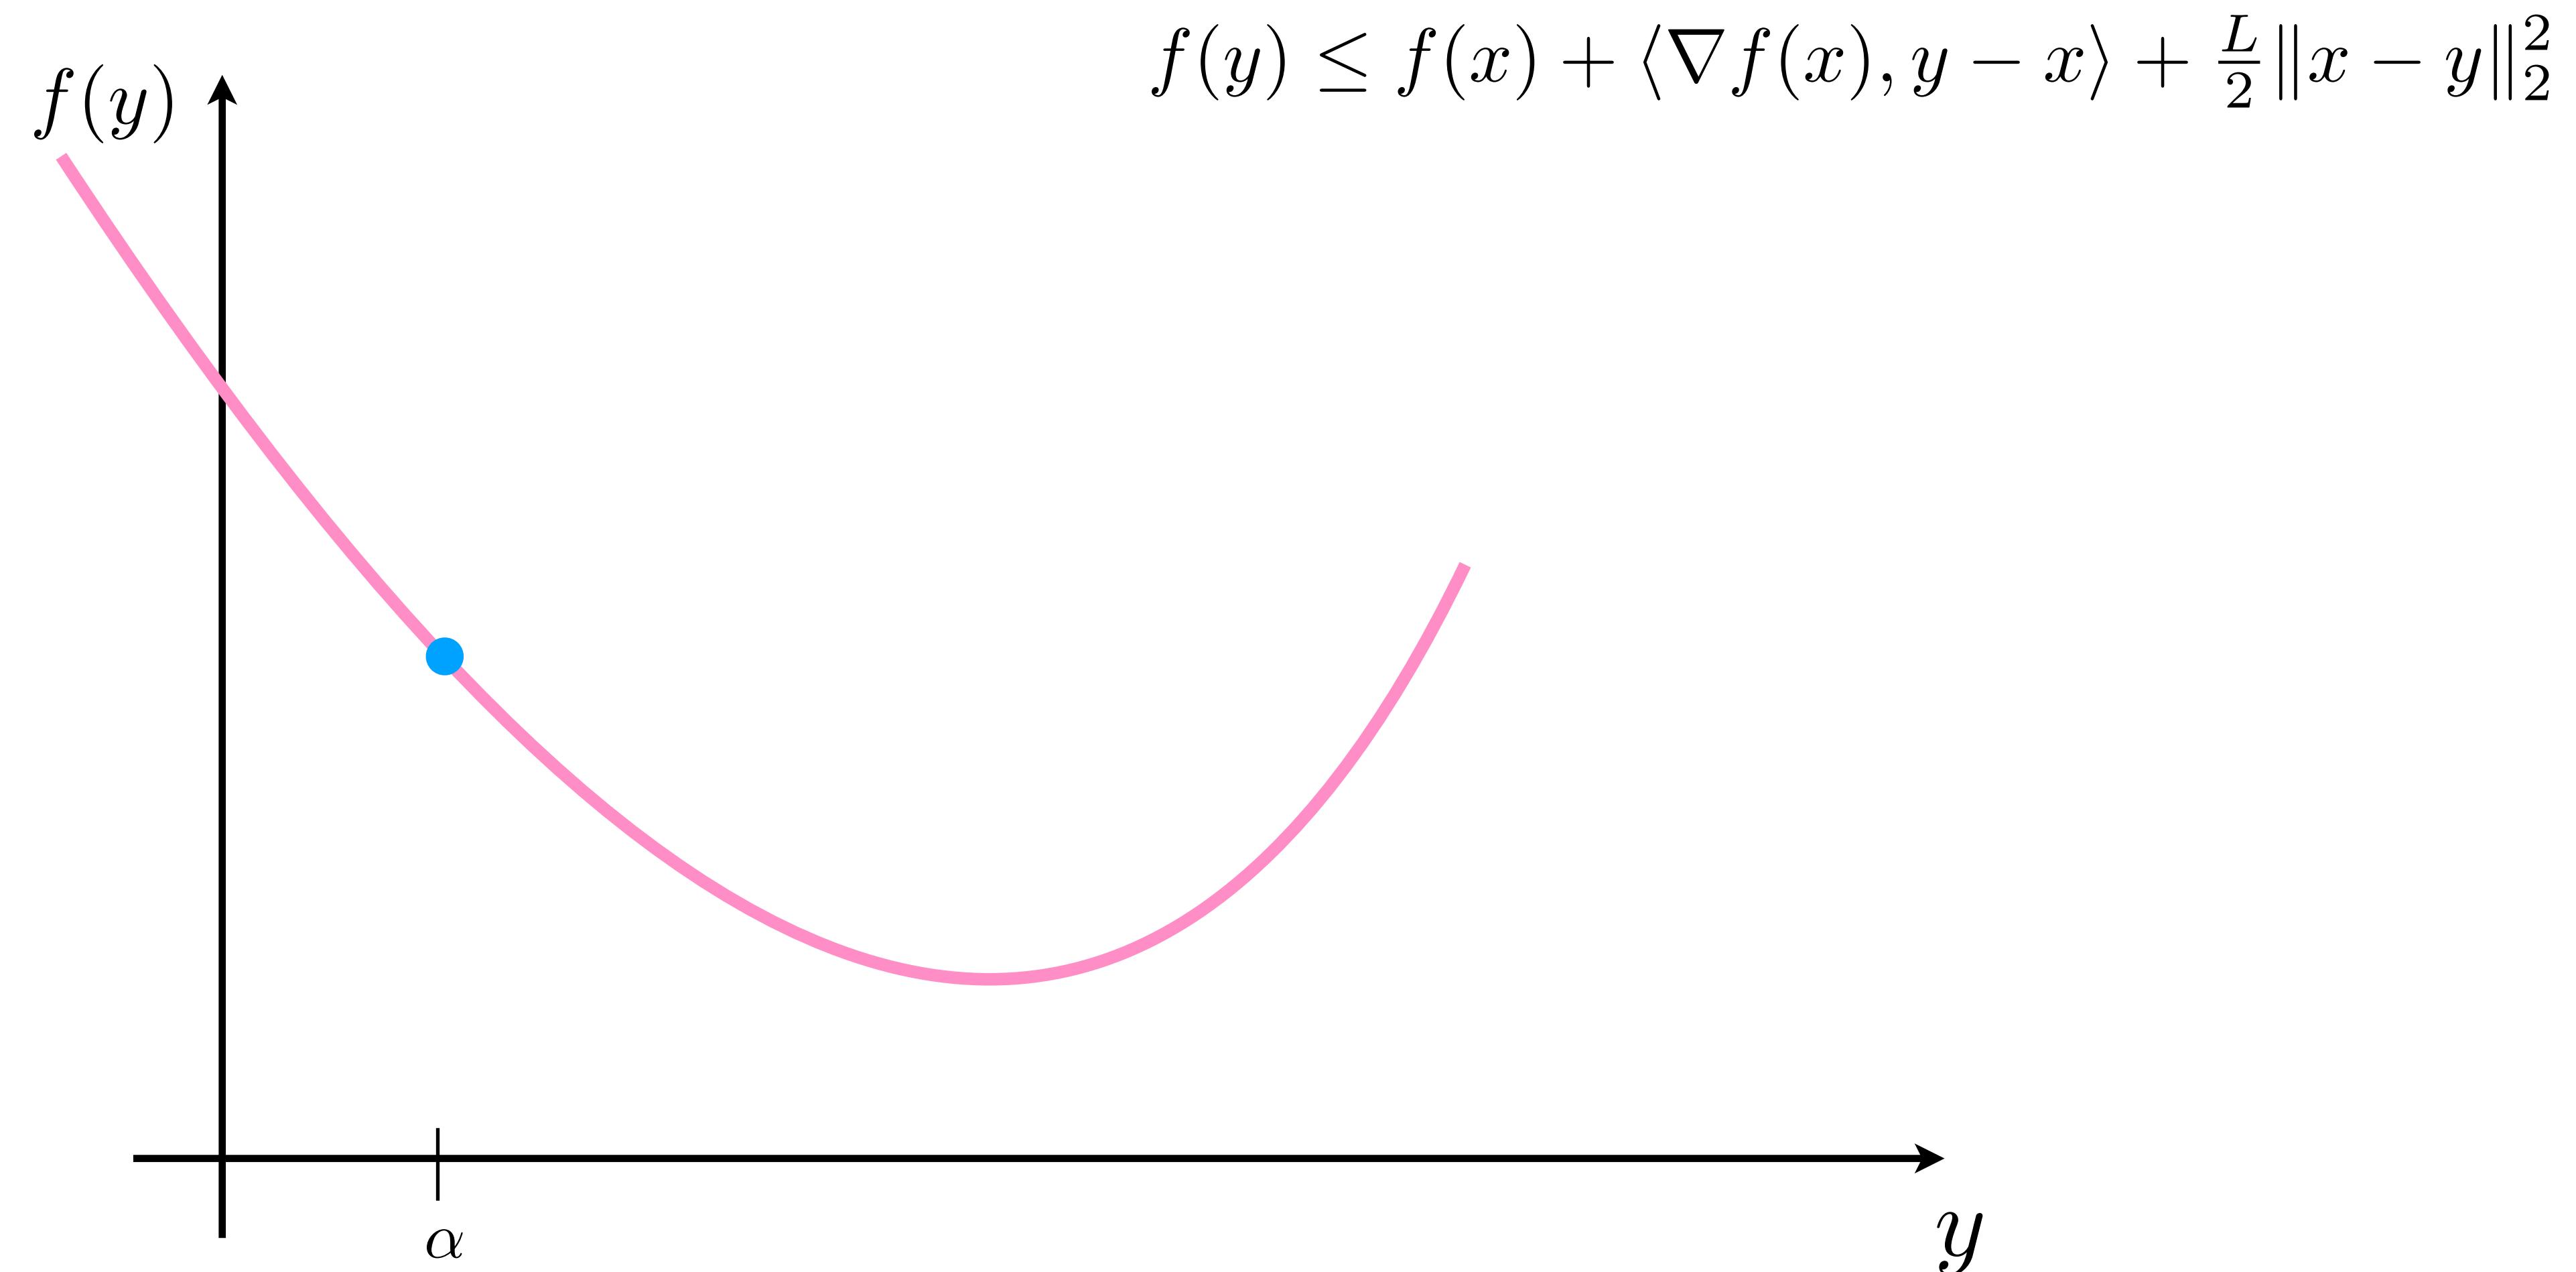
Lipschitz conditions

- Equivalent characterizations: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$



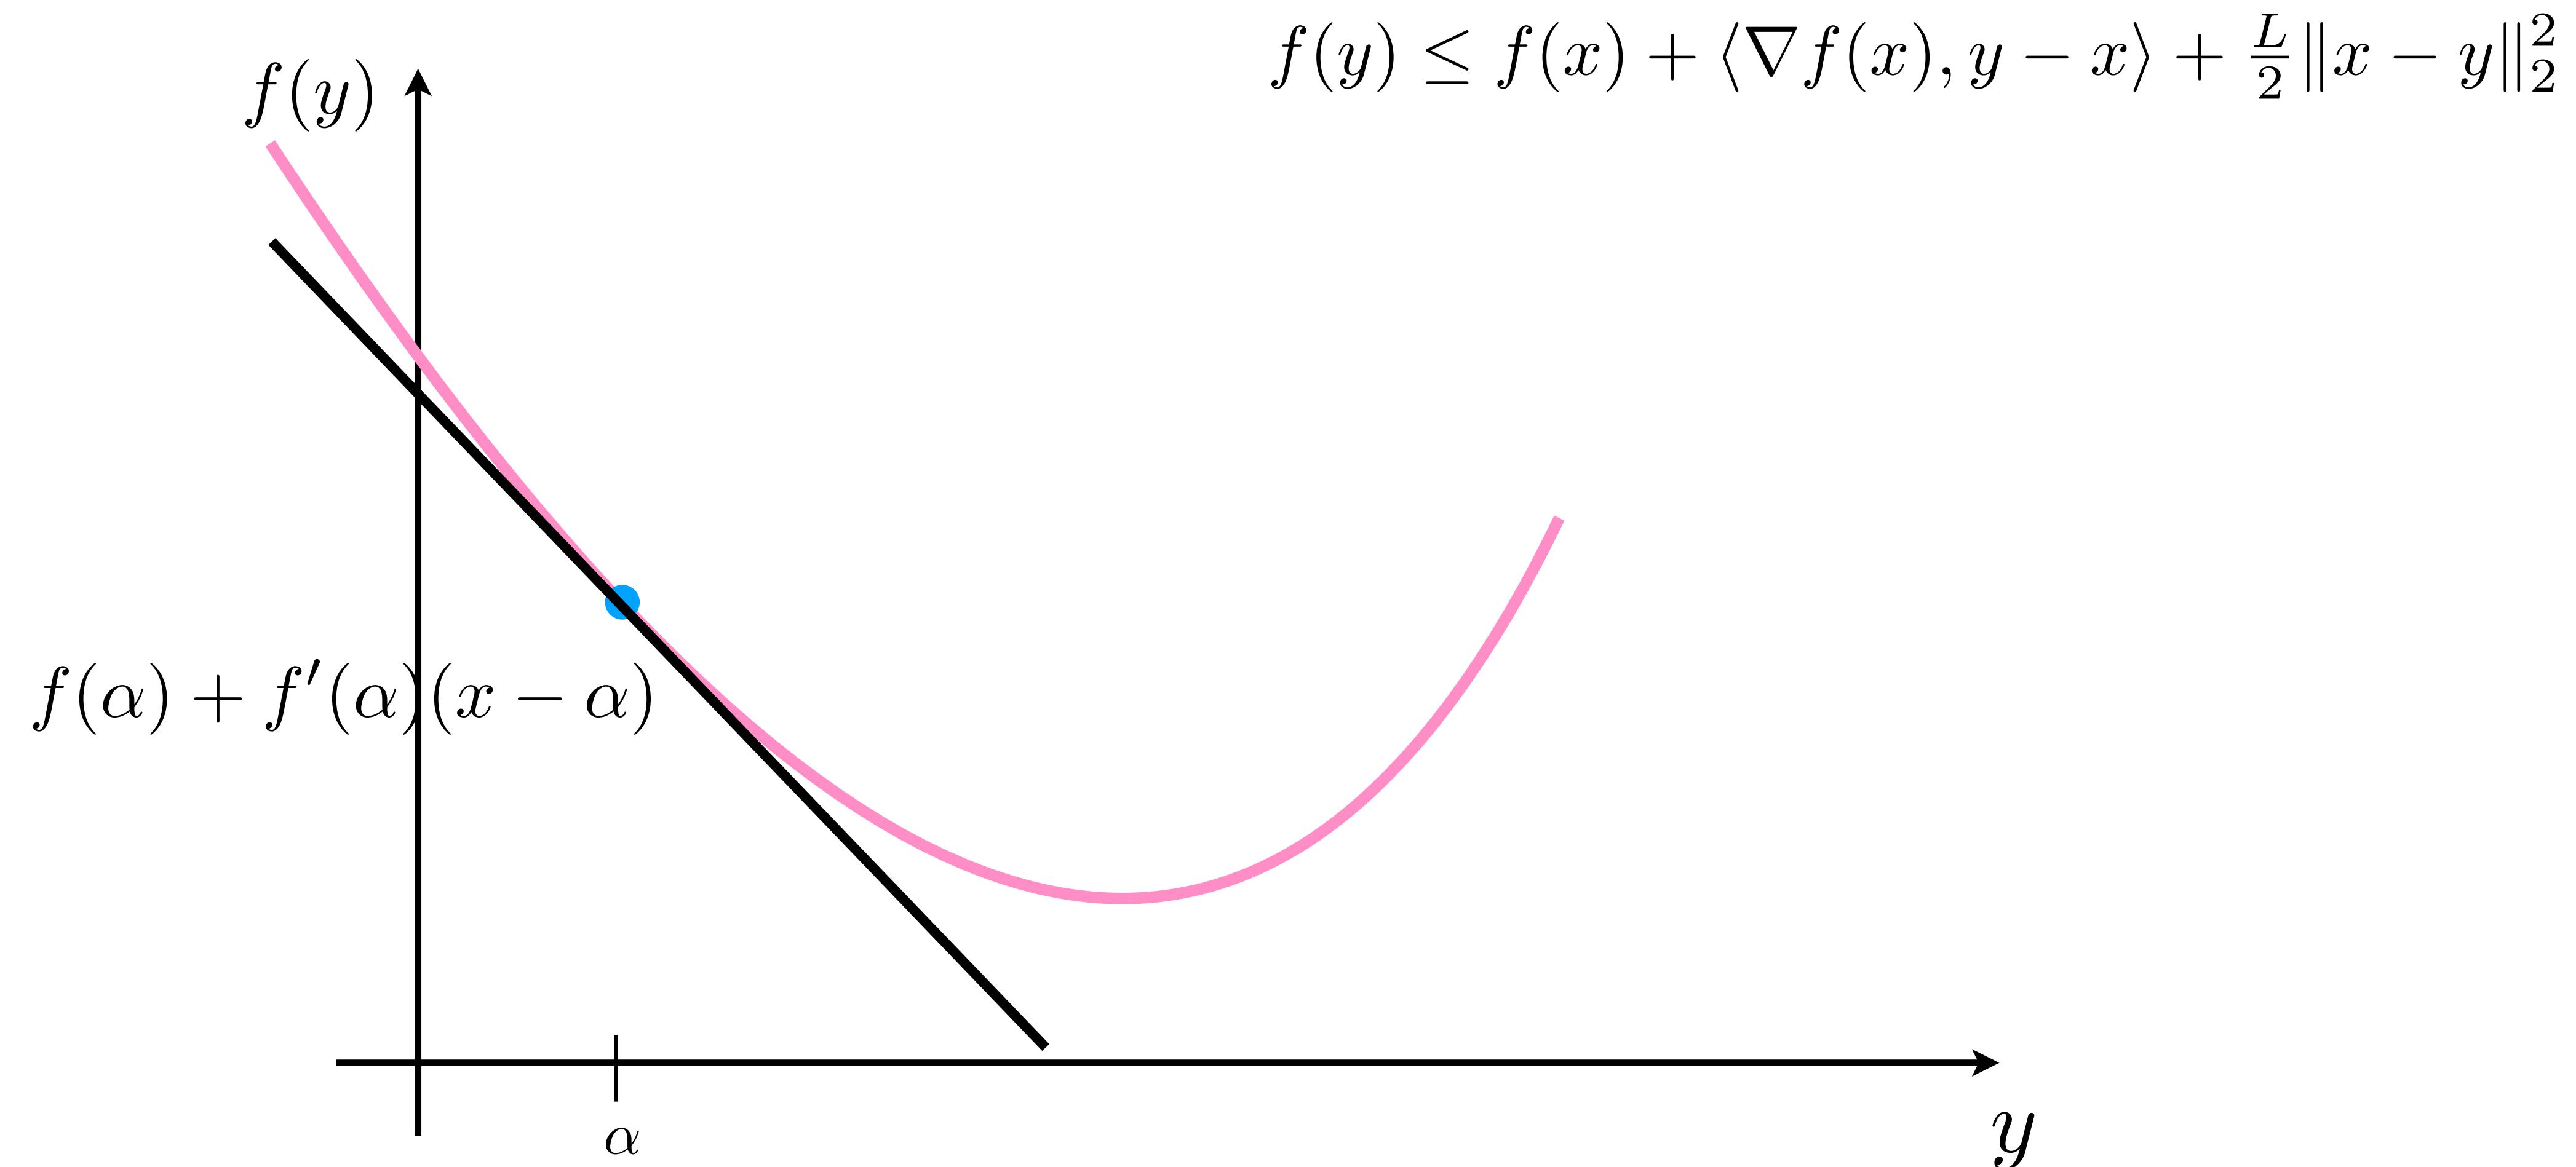
Lipschitz conditions

- Equivalent characterizations: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$



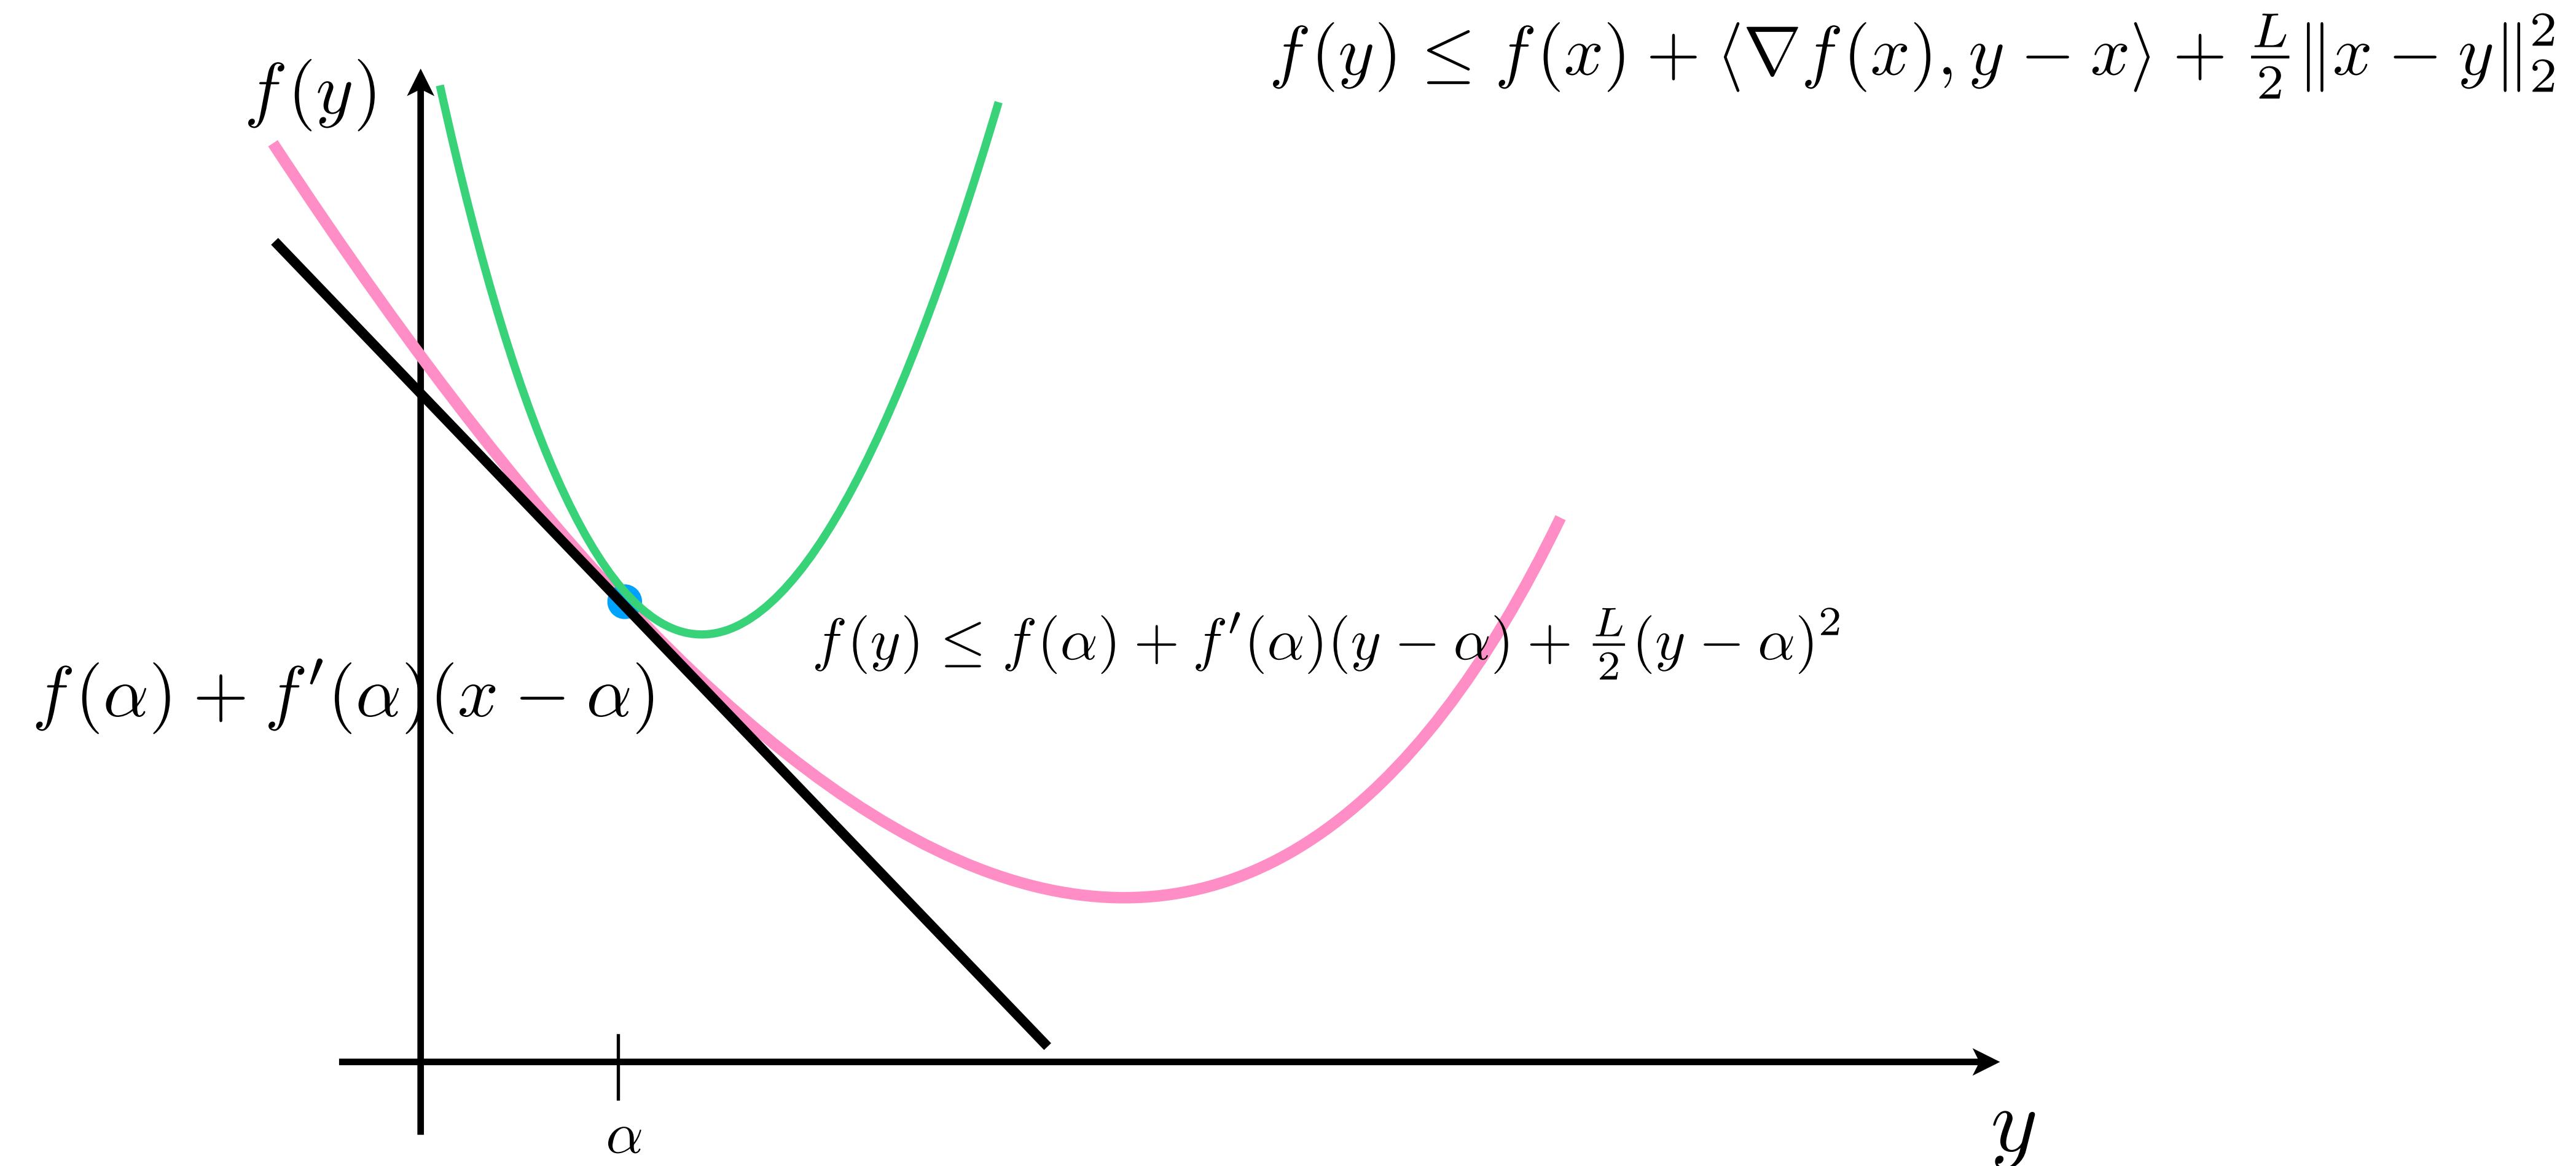
Lipschitz conditions

- Equivalent characterizations: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$



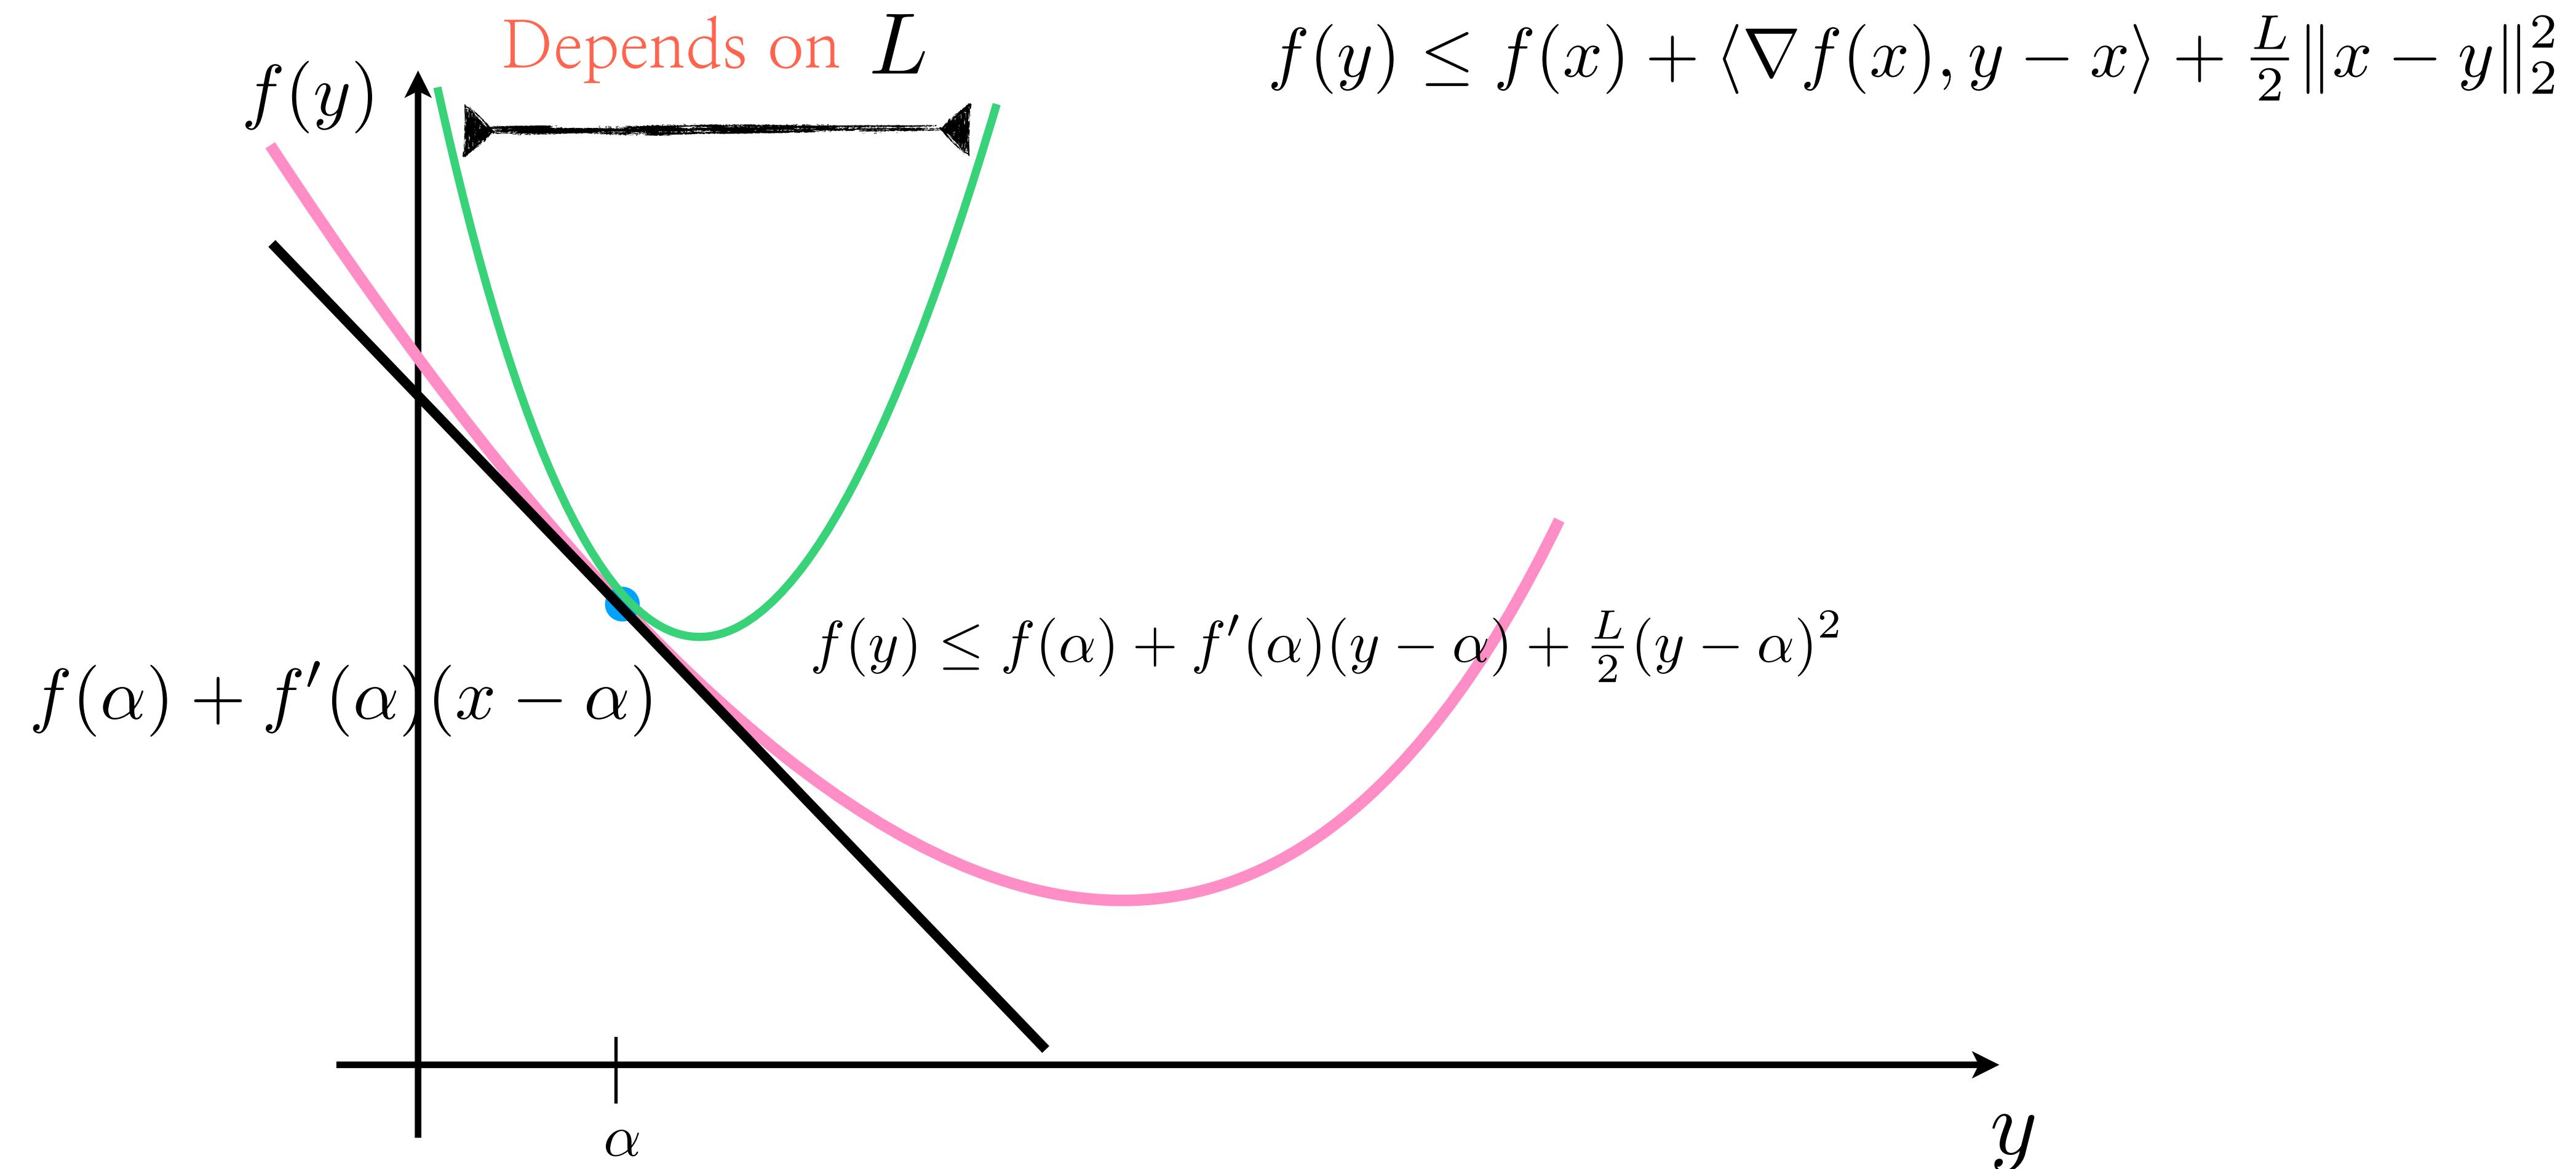
Lipschitz conditions

- Equivalent characterizations: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$



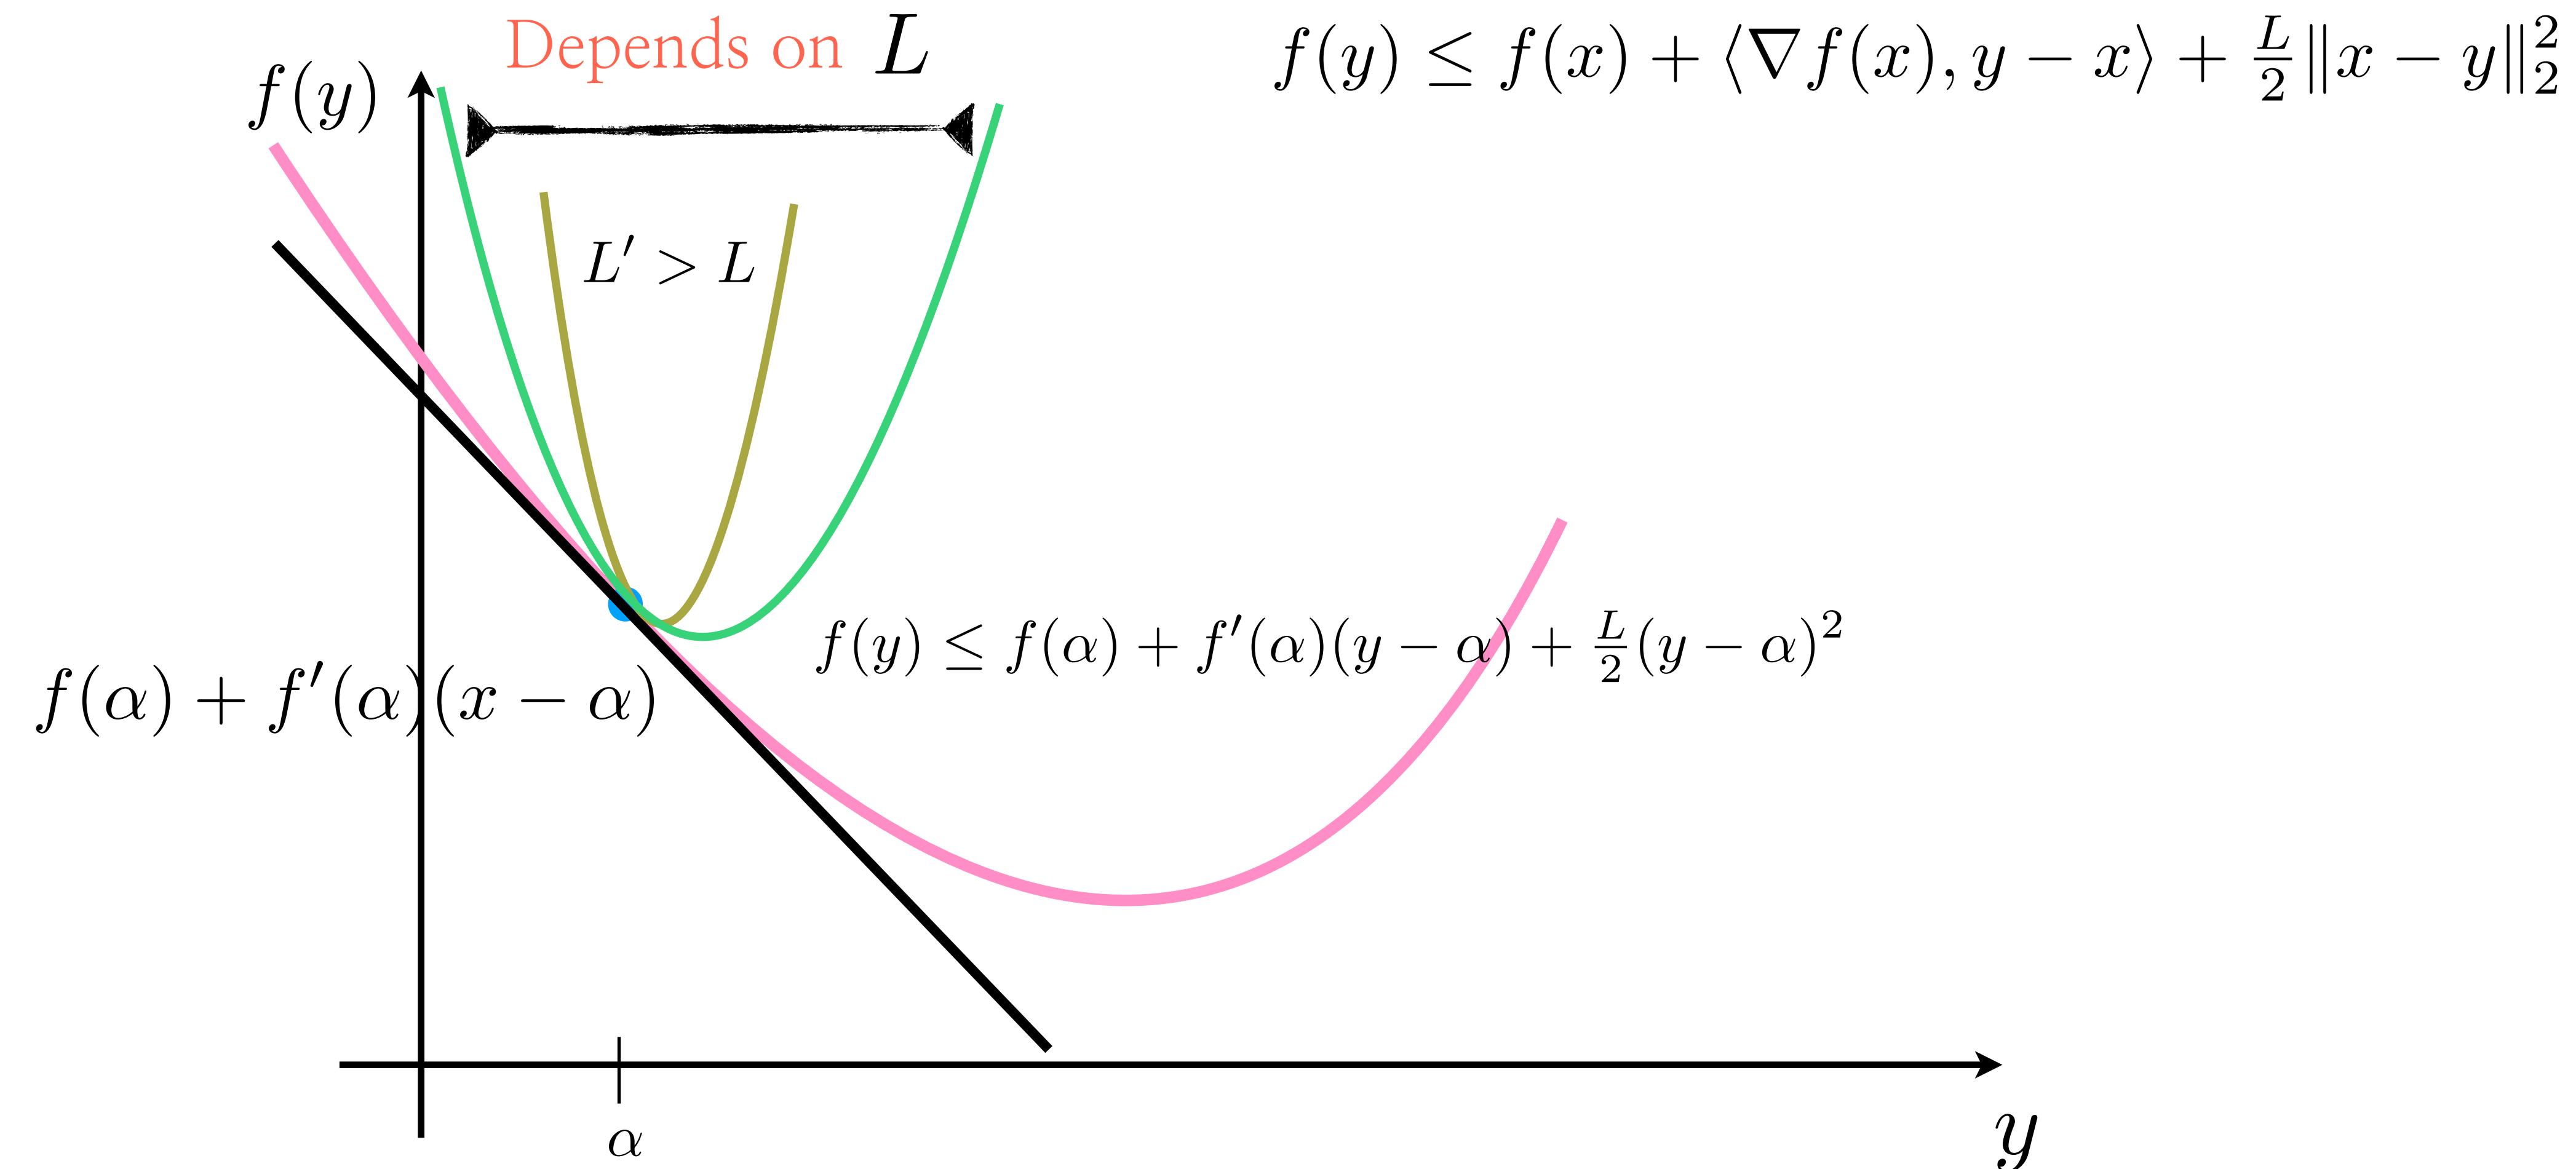
Lipschitz conditions

- Equivalent characterizations: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$



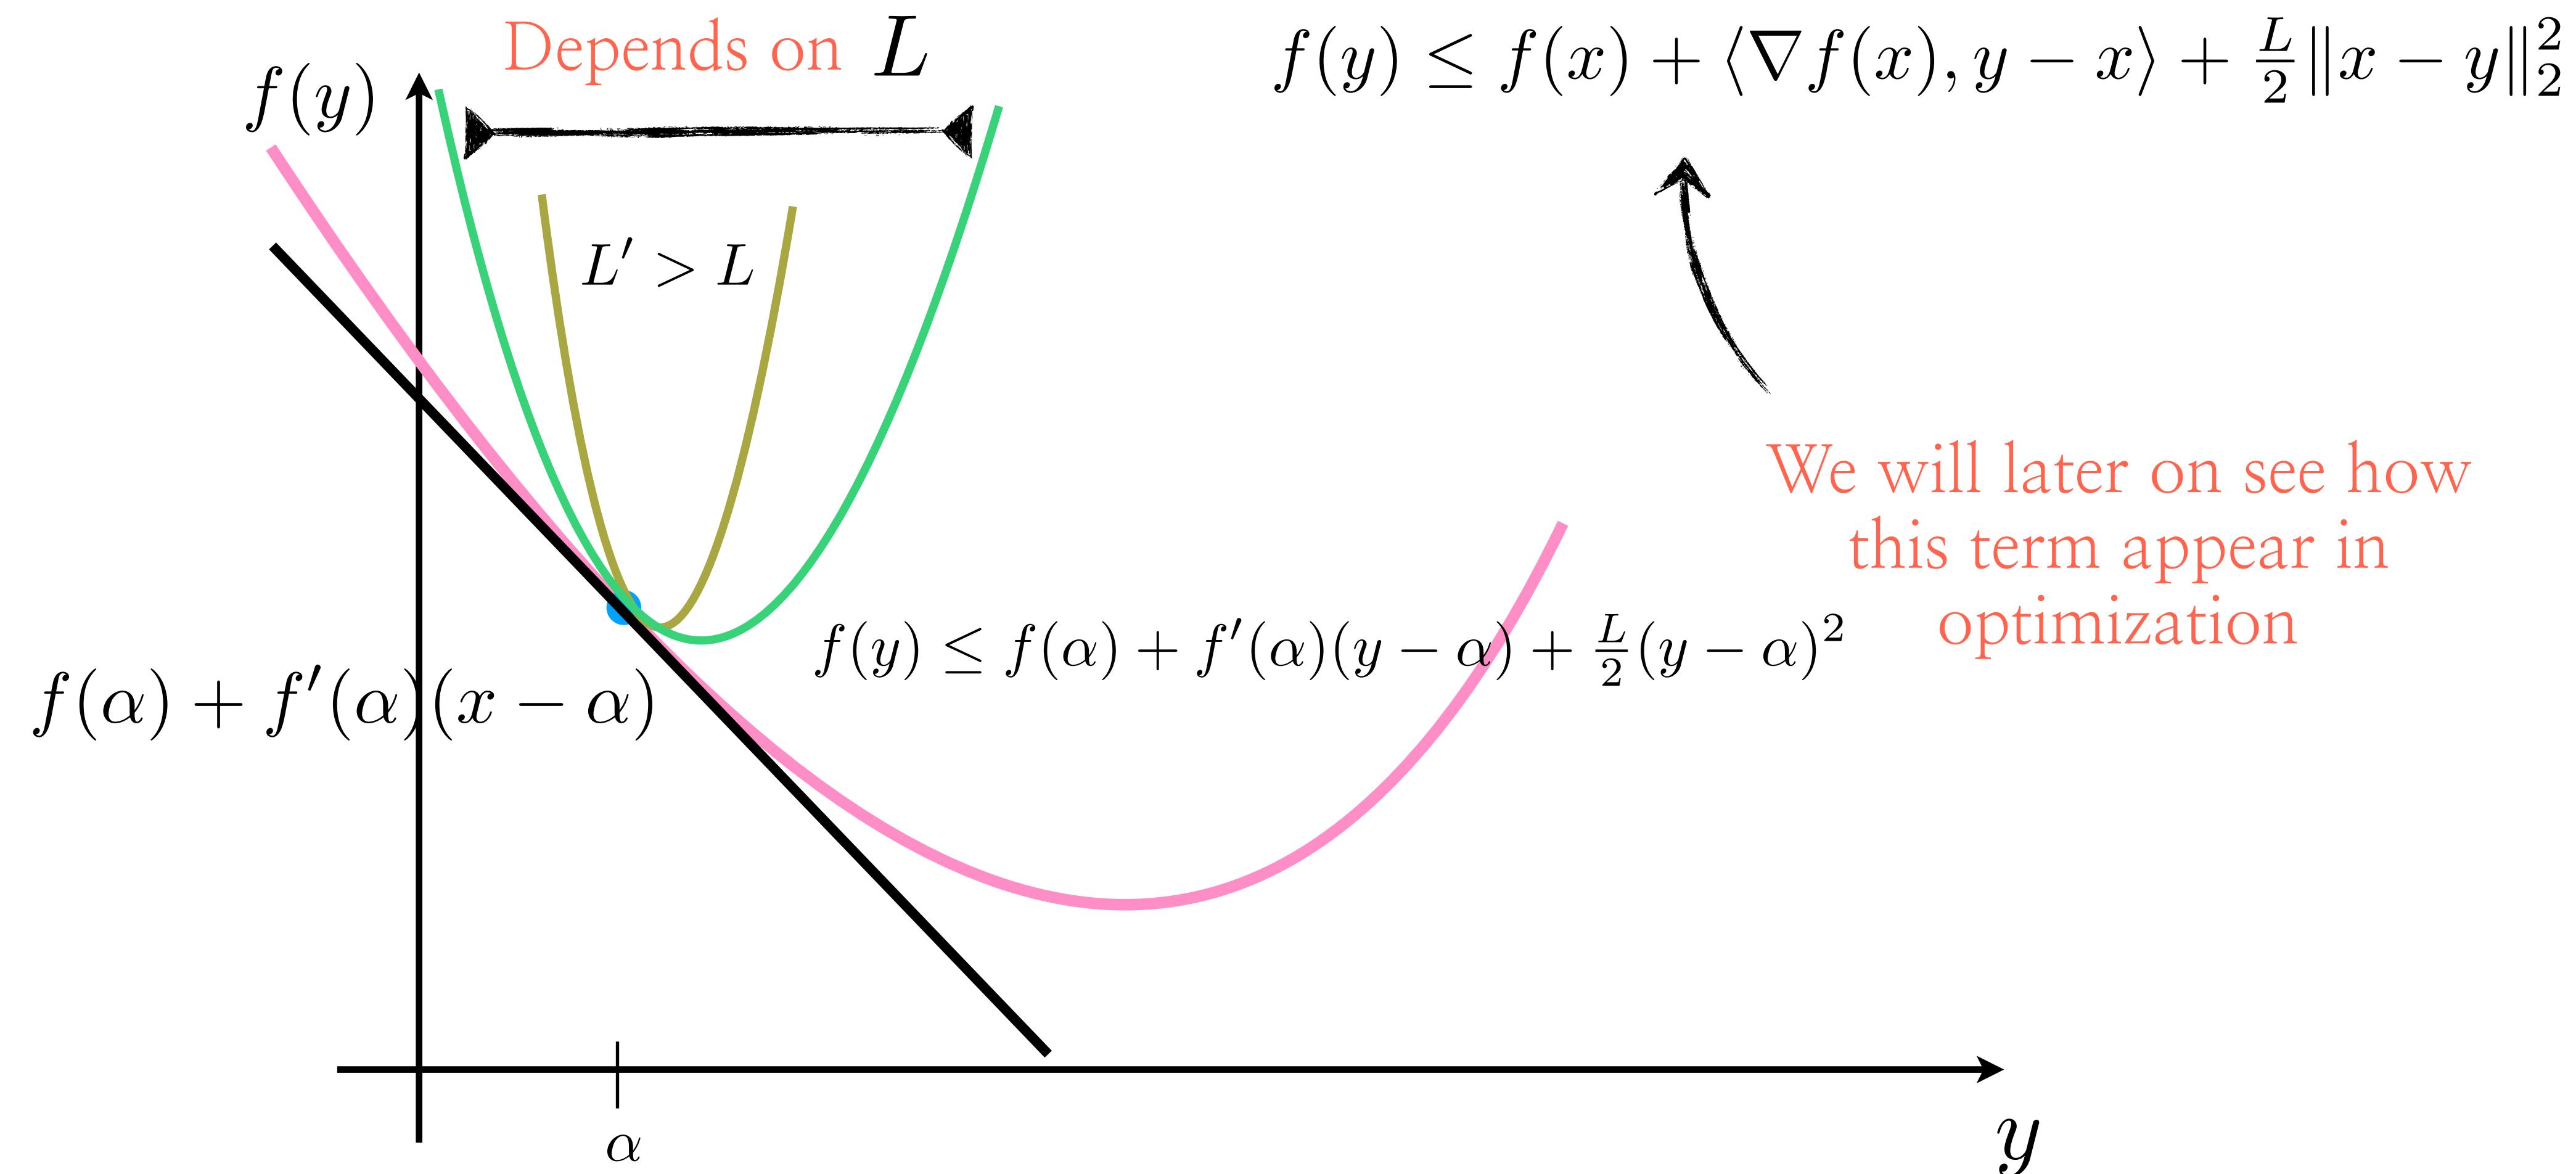
Lipschitz conditions

- Equivalent characterizations: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$



Lipschitz conditions

- Equivalent characterizations: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2, \quad \forall x, y$



Lipschitz conditions

- How does this relate to Taylor's expansion?

Lipschitz conditions

- How does this relate to Taylor's expansion?

$$f(y) = f(x) + \langle \nabla f(x), y - x \rangle + \frac{1}{2} \langle \nabla^2 f(z)(y - x), y - x \rangle$$

Lipschitz conditions

- How does this relate to Taylor's expansion?

$$f(y) = f(x) + \langle \nabla f(x), y - x \rangle + \frac{1}{2} \langle \nabla^2 f(z)(y - x), y - x \rangle$$

- From $\nabla^2 f(x) \preceq LI$, we have:

$$\nabla^2 f(x) \preceq LI \Rightarrow \|\nabla^2 f(x)\|_2 \leq \|LI\|_2 \Rightarrow \|\nabla^2 f(x)\|_2 \leq L$$

Lipschitz conditions

- How does this relate to Taylor's expansion?

$$f(y) = f(x) + \langle \nabla f(x), y - x \rangle + \frac{1}{2} \langle \nabla^2 f(z)(y - x), y - x \rangle$$

- From $\nabla^2 f(x) \preceq LI$, we have:

$$\nabla^2 f(x) \preceq LI \Rightarrow \|\nabla^2 f(x)\|_2 \leq \|LI\|_2 \Rightarrow \|\nabla^2 f(x)\|_2 \leq L$$

- Then:

$$\frac{1}{2} \langle \nabla^2 f(x)(y - x), y - x \rangle \leq \frac{1}{2} \|\nabla^2 f(x)(y - x)\|_2 \cdot \|y - x\|_2 \leq \|\nabla^2 f(x)\|_2 \|y - x\|_2^2$$

Lipschitz conditions

- How does this relate to Taylor's expansion?

$$f(y) = f(x) + \langle \nabla f(x), y - x \rangle + \frac{1}{2} \langle \nabla^2 f(z)(y - x), y - x \rangle$$

- From $\nabla^2 f(x) \preceq LI$, we have:

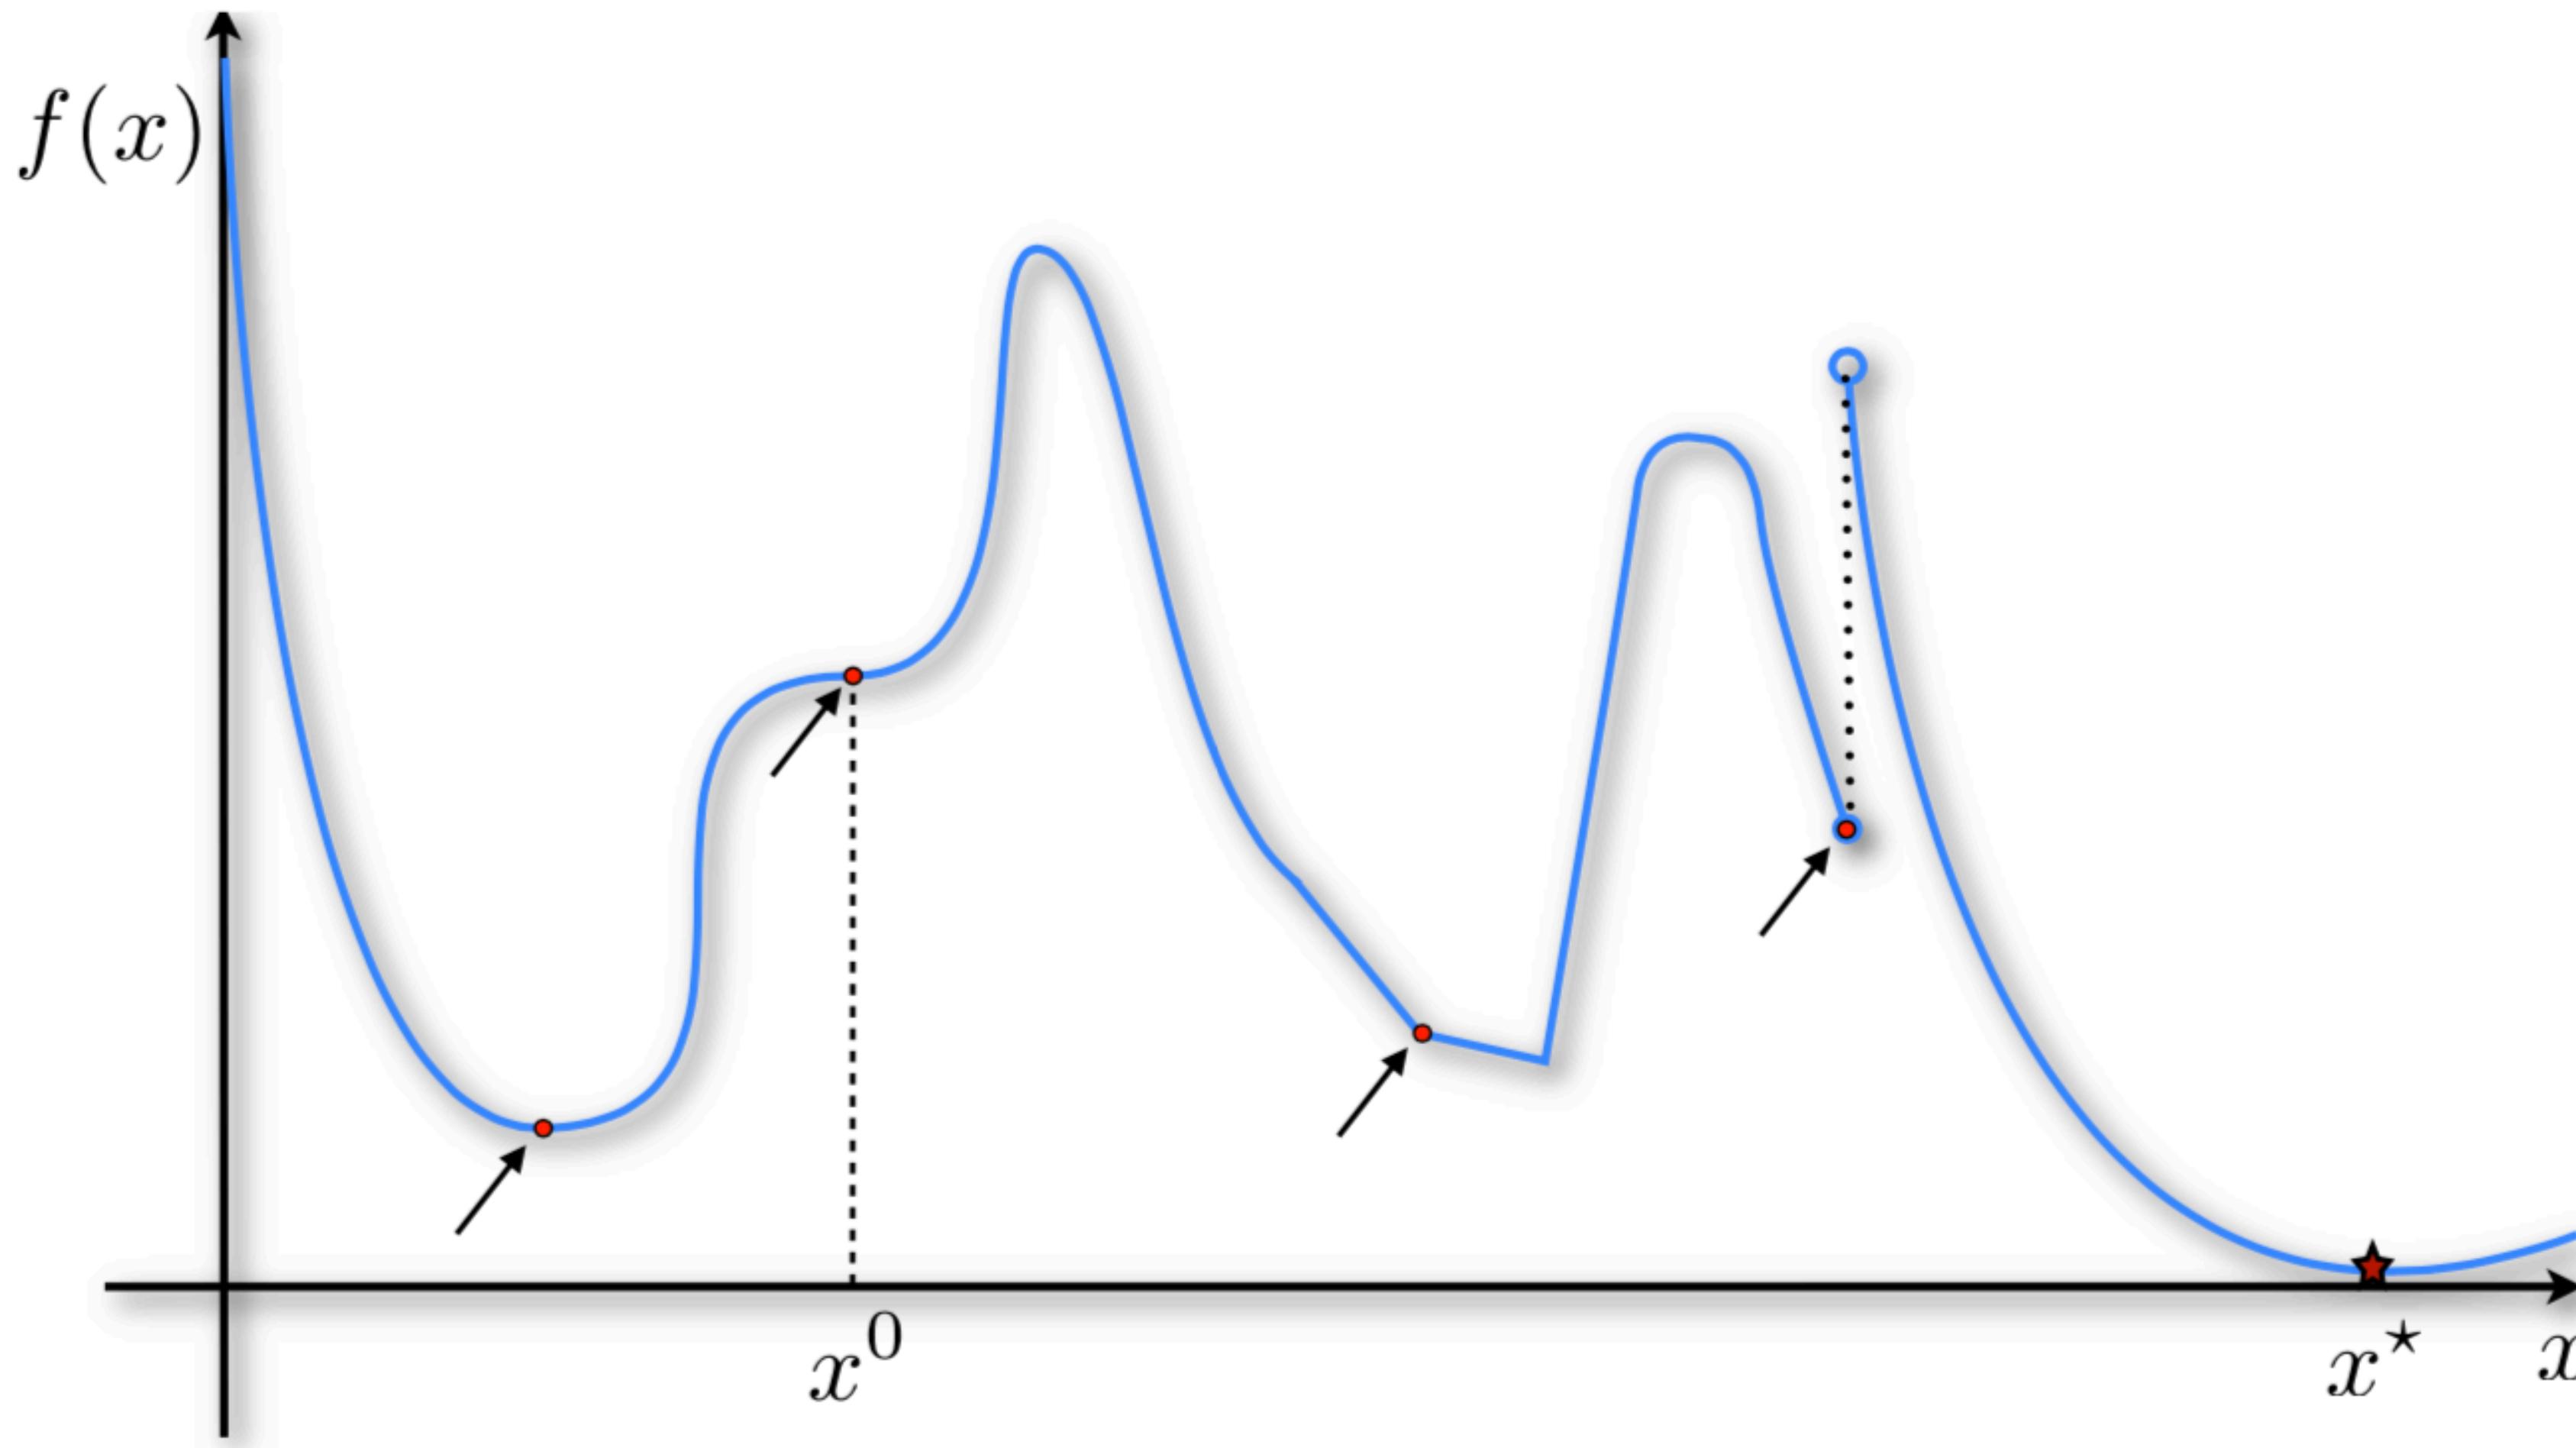
$$\nabla^2 f(x) \preceq LI \Rightarrow \|\nabla^2 f(x)\|_2 \leq \|LI\|_2 \Rightarrow \|\nabla^2 f(x)\|_2 \leq L$$

- Then:

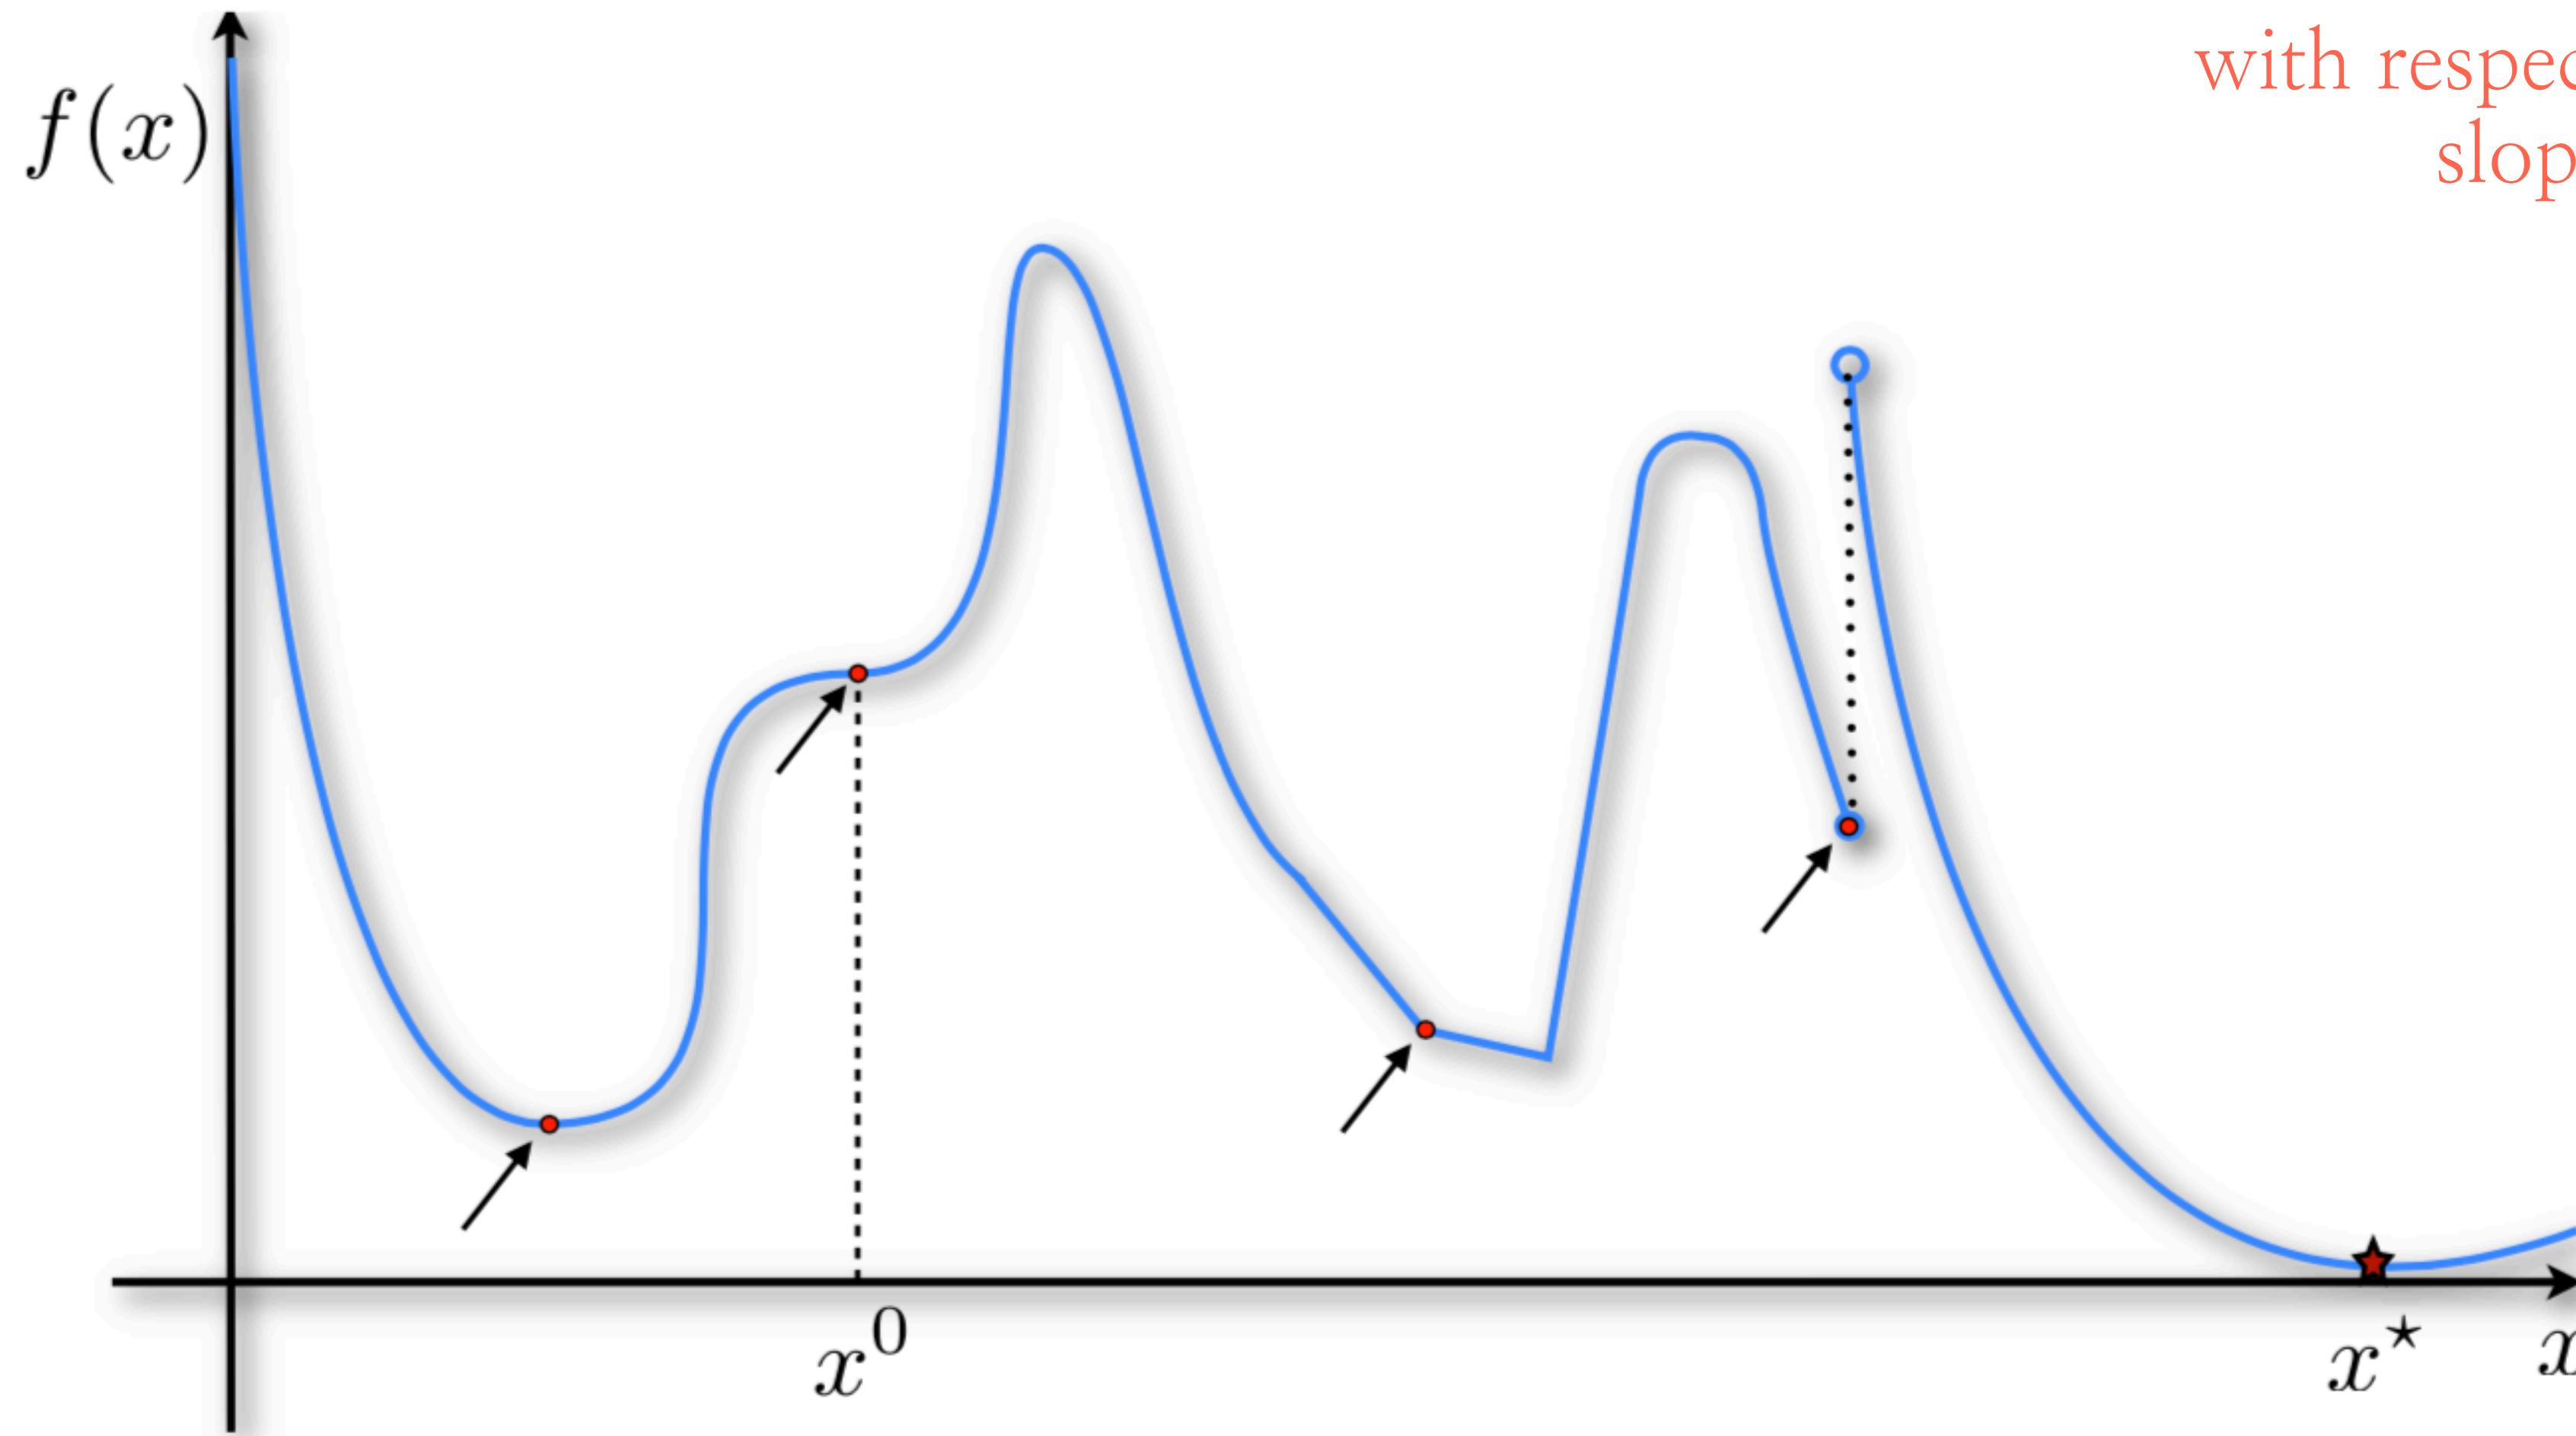
$$\frac{1}{2} \langle \nabla^2 f(x)(y - x), y - x \rangle \leq \frac{1}{2} \|\nabla^2 f(x)(y - x)\|_2 \cdot \|y - x\|_2 \leq \|\nabla^2 f(x)\|_2 \|y - x\|_2^2$$

$$f(y) \leq f(x) + \langle \nabla f(x), y - x \rangle + \frac{L}{2} \|x - y\|_2^2$$

Agnostic optimization



Agnostic optimization



What do you observe at
local minima/maxima
with respect to their
slope?

Types of solutions

- Global minimizer x^* :

Types of solutions

- Global minimizer x^* : $f(x^*) \leq f(x), \forall x$

(If we could somehow check that, we are golden!)

Types of solutions

- Global minimizer x^* : $f(x^*) \leq f(x), \forall x$
(If we could somehow check that, we are golden!)
- Local minimizer \hat{x} :

Types of solutions

- Global minimizer x^* : $f(x^*) \leq f(x), \forall x$
(If we could somehow check that, we are golden!)
- Local minimizer \hat{x} : $f(\hat{x}) \leq f(x), \forall x \in \mathcal{N}_{\hat{x}}$

Types of solutions

- Global minimizer x^* : $f(x^*) \leq f(x), \forall x$
(If we could somehow check that, we are golden!)
- Local minimizer \hat{x} : $f(\hat{x}) \leq f(x), \forall x \in \mathcal{N}_{\hat{x}}$
- What is the meaning of strict inequality vs. inequality?

Types of solutions

- Global minimizer x^* : $f(x^*) \leq f(x), \forall x$
(If we could somehow check that, we are golden!)
- Local minimizer \hat{x} : $f(\hat{x}) \leq f(x), \forall x \in \mathcal{N}_{\hat{x}}$
- What is the meaning of strict inequality vs. inequality?
- How do we recognize that a solution we have is a local (global) solution?

Types of solutions

- Global minimizer x^* : $f(x^*) \leq f(x), \forall x$
(If we could somehow check that, we are golden!)
- Local minimizer \hat{x} : $f(\hat{x}) \leq f(x), \forall x \in \mathcal{N}_{\hat{x}}$
- What is the meaning of strict inequality vs. inequality?
- How do we recognize that a solution we have is a local (global) solution?
 - 1st-order optimality condition: $\nabla f(\hat{x}) = 0$

Types of solutions

- Global minimizer x^* : $f(x^*) \leq f(x), \forall x$
(If we could somehow check that, we are golden!)
- Local minimizer \hat{x} : $f(\hat{x}) \leq f(x), \forall x \in \mathcal{N}_{\hat{x}}$
- What is the meaning of strict inequality vs. inequality?
- How do we recognize that a solution we have is a local (global) solution?
 - 1st-order optimality condition: $\nabla f(\hat{x}) = 0$
 - 2nd-order optimality condition: $\nabla f(\hat{x}) = 0$ and $\nabla^2 f(\hat{x}) \succeq 0$

Types of solutions

- Global minimizer x^* : $f(x^*) \leq f(x), \forall x$
(If we could somehow check that, we are golden!)
- Local minimizer \hat{x} : $f(\hat{x}) \leq f(x), \forall x \in \mathcal{N}_{\hat{x}}$
- What is the meaning of strict inequality vs. inequality?
- How do we recognize that a solution we have is a local (global) solution?

Necessary

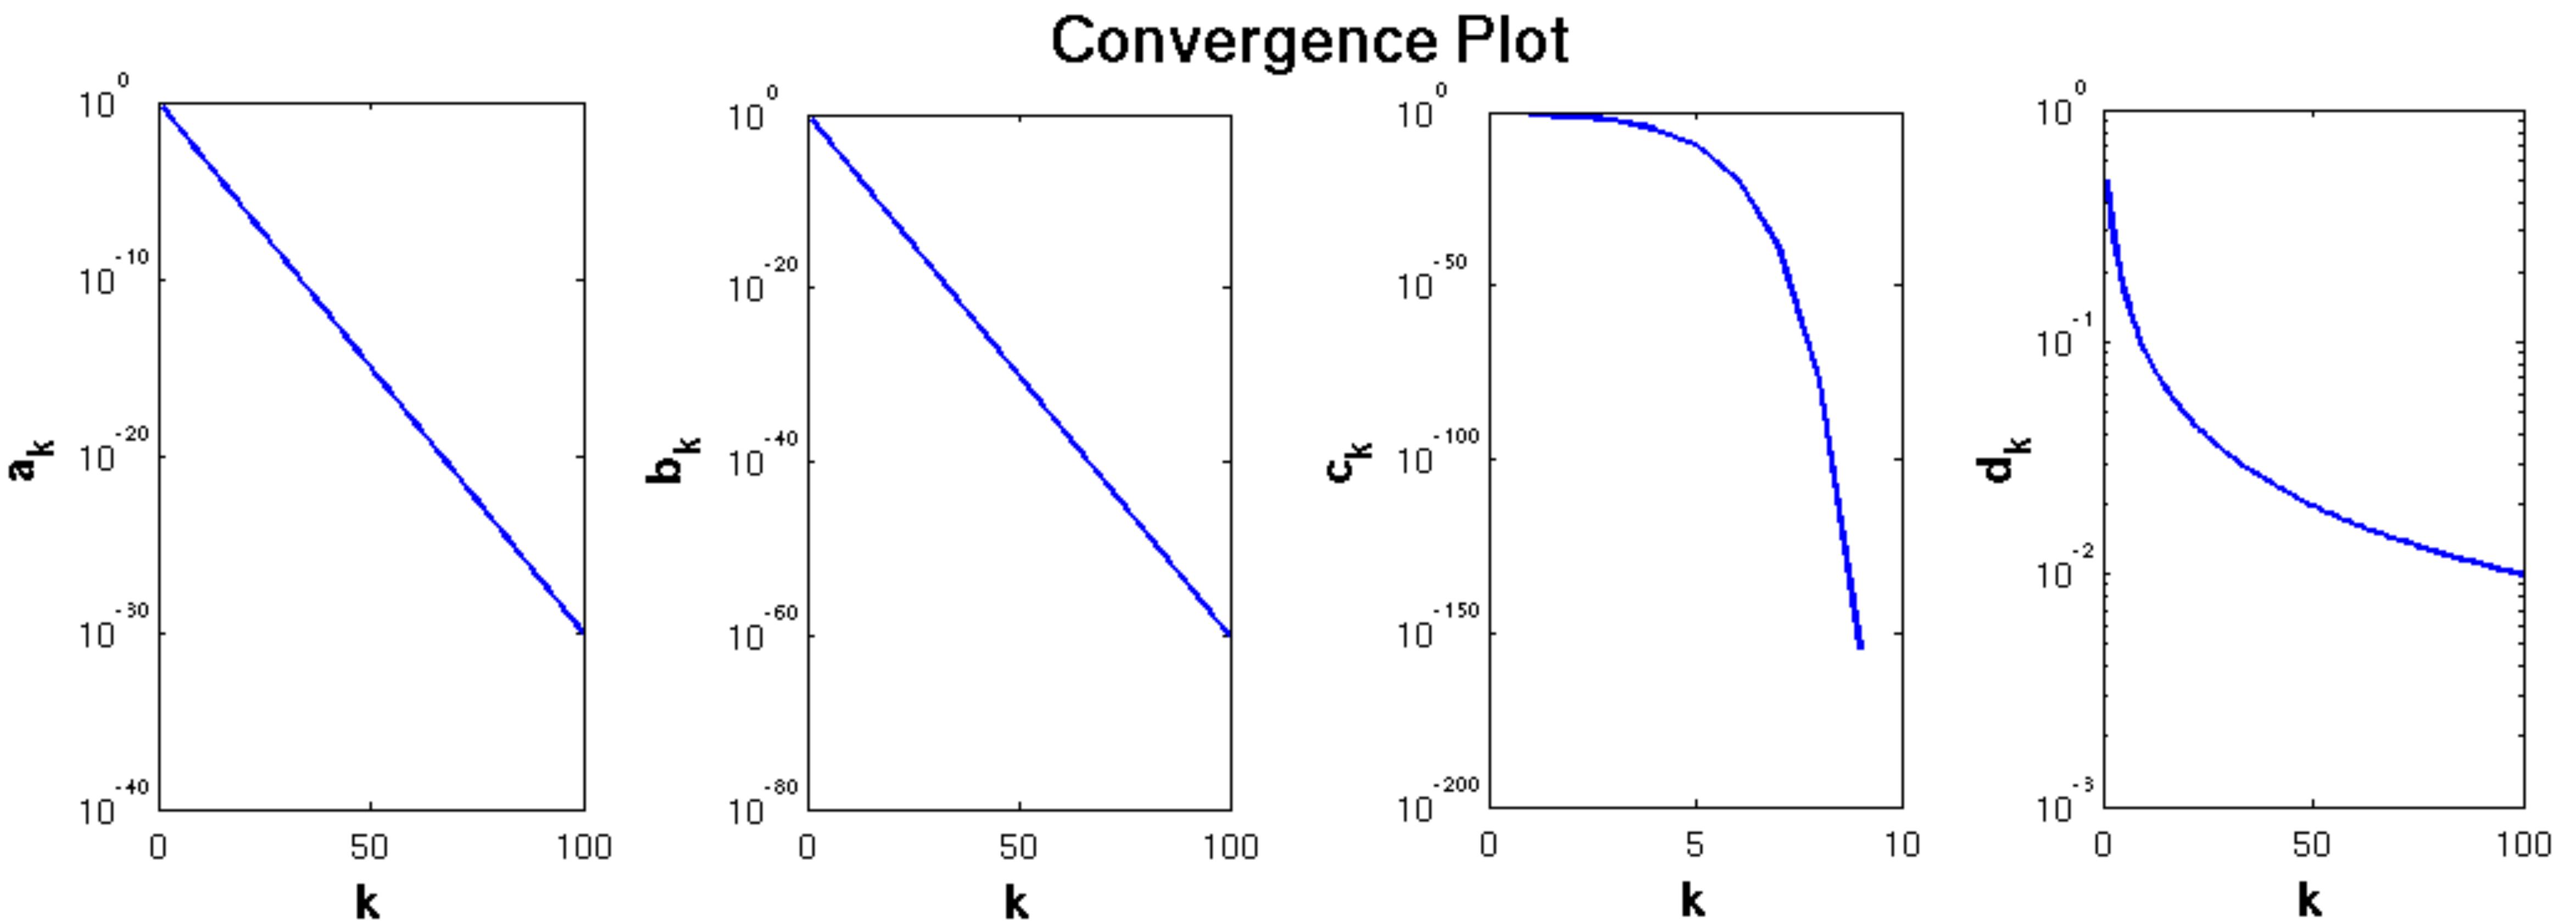
- 1st-order optimality condition: $\nabla f(\hat{x}) = 0$
- 2nd-order optimality condition: $\nabla f(\hat{x}) = 0$ and $\nabla^2 f(\hat{x}) \succeq 0$

First convergence result

Whiteboard

Convergence rates 101

(Source: Wikipedia)



$$O(\log 1/\varepsilon)$$

$$q^k, \quad q \in (0, 1)$$

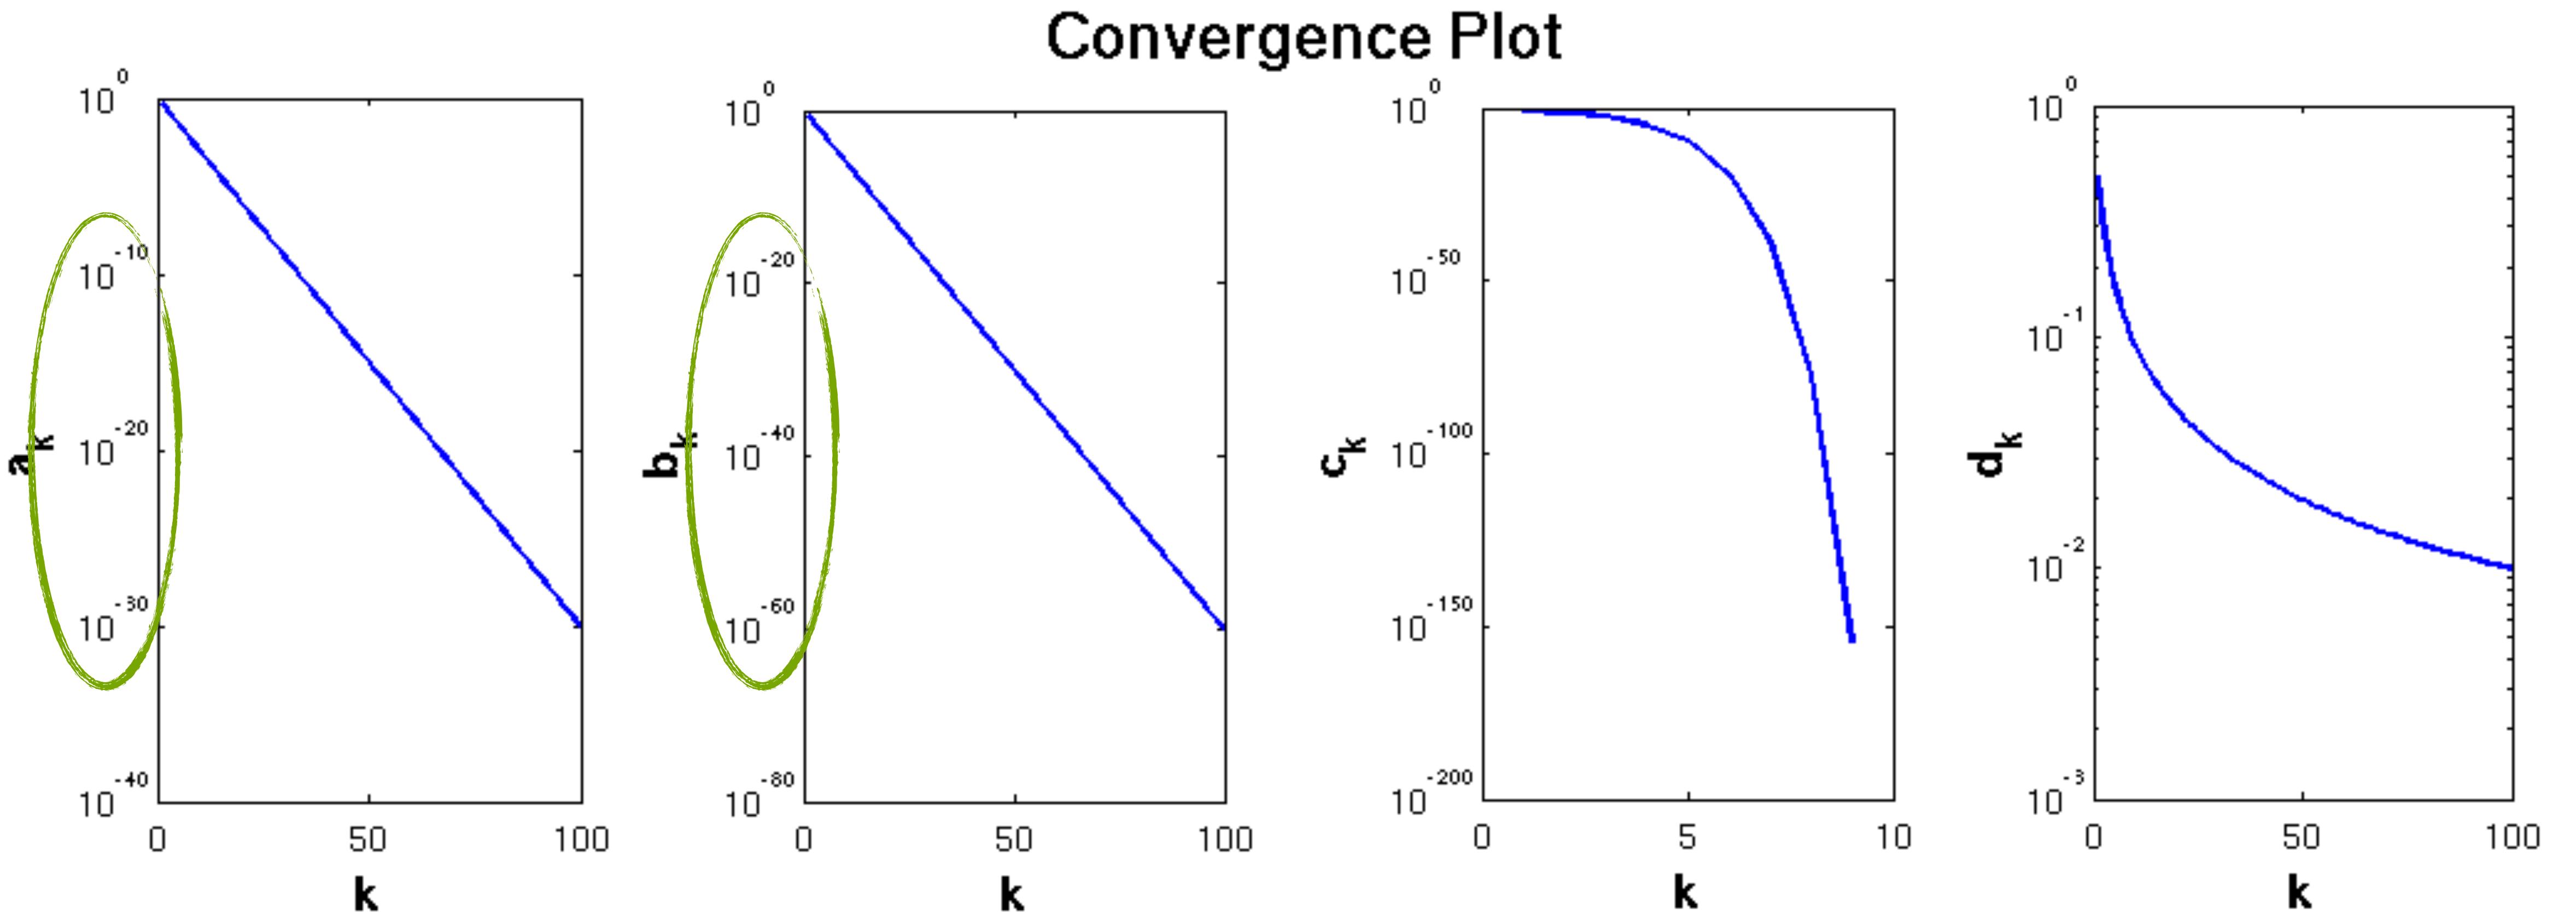
$$O(\log \log(1/\varepsilon))$$

$$O(1/\varepsilon^2), \quad O(1/\varepsilon), \quad O(1/\sqrt{\varepsilon})$$

$$O(1/k^2), \quad O(1/k), \quad O(\sqrt{k})$$

Convergence rates 101

(Source: Wikipedia)



$$O(\log 1/\varepsilon)$$

$$q^k, \quad q \in (0, 1)$$

$$O(\log \log(1/\varepsilon))$$

$$O(1/\varepsilon^2), \quad O(1/\varepsilon), \quad O(1/\sqrt{\varepsilon})$$

$$O(1/k^2), \quad O(1/k), \quad O(\sqrt(k))$$

First convergence result

$$\min_{x \in \mathbb{R}^p} f(x)$$

“Assume the objective is has Lipschitz continuous gradients. Then, gradient descent:

$$x_{t+1} = x_t - \eta \nabla f(x_t)$$

with step size

$$\eta = \frac{1}{L}$$

converges sublinearly to a stationary point; i.e.,

$$\min_t \|\nabla f(x_t)\|_2 \leq \sqrt{\frac{2L}{T+1}} \cdot (f(x_0) - f(x^\star))^{1/2} = O\left(\frac{1}{\sqrt{T}}\right)$$

First convergence result

- “But, which functions satisfy Lipschitz gradient continuity?”

First convergence result

- “But, which functions satisfy Lipschitz gradient continuity?”
 - Least-squares objectives: $f(x) = \frac{1}{2} \|Ax - b\|_2^2$

$$\|\nabla f(x) - \nabla f(y)\|_2 \leq \|A^\top A\|_2 \cdot \|x - y\|_2$$

First convergence result

- “But, which functions satisfy Lipschitz gradient continuity?”

- Least-squares objectives: $f(x) = \frac{1}{2} \|Ax - b\|_2^2$

$$\|\nabla f(x) - \nabla f(y)\|_2 \leq \|A^\top A\|_2 \cdot \|x - y\|_2$$

- Logistic regression objectives: $f(x) = \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-y_i \alpha_i^\top x))$

Whiteboard

First convergence result

- “But, which functions satisfy Lipschitz gradient continuity?”

- Least-squares objectives: $f(x) = \frac{1}{2} \|Ax - b\|_2^2$

$$\|\nabla f(x) - \nabla f(y)\|_2 \leq \|A^\top A\|_2 \cdot \|x - y\|_2$$

- Logistic regression objectives: $f(x) = \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-y_i \alpha_i^\top x))$



Whiteboard

“But these are actually convex!”

First convergence result

- “But, which functions satisfy Lipschitz gradient continuity?”

- Least-squares objectives: $f(x) = \frac{1}{2} \|Ax - b\|_2^2$

$$\|\nabla f(x) - \nabla f(y)\|_2 \leq \|A^\top A\|_2 \cdot \|x - y\|_2$$

- Logistic regression objectives: $f(x) = \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-y_i \alpha_i^\top x))$



“But these are actually convex!”

Whiteboard

- Non-convex objective: $f(x) = x^2 + 3 \sin^2(x)$

Demo

Conclusion

- We have set up background of smooth optimization
- We have provided the first convergence rate result, and defined different convergence rates that could be attainable

Conclusion

- We have set up background of smooth optimization
- We have provided the first convergence rate result, and defined different convergence rates that could be attainable

Next lecture

- Brief introduction to convex optimization and related topics