

COMP 414/514:
Optimization – Algorithms, Complexity
and Approximations

Lecture 7

Overview

- In the last lecture, we:
 - Talked about how acceleration leads to a better convergence rate
 - Worked in practice and theory with accelerated gradient descent variants
 - Discussed the limits and convergence rates of accelerated gradient descent
- Often, gradient descent is not sufficient in practice. In this lecture, we will:
 - Discuss **alternatives to batch gradient descent: stochastic gradient descent**
 - Discuss **alternatives to batch gradient descent: coordinate descent**
 - Discuss recent advances on these topics

Acceleration #2: Cut-off complexity per iteration

- Common situation in machine learning/signal processing

$$f(x) = \frac{1}{n} \sum_{i=1}^n f_i(x) \quad (\text{Empirical risk minimization})$$

where each $f_i(x)$ depends on, let's say, different part of input data.

- Examples:

- Least squares: $f_i(x) = \frac{1}{2} (y_i - \alpha_i^\top x)^2$

- Logistic regression: $f_i(x) = \log(1 + \exp(-y_i \alpha_i^\top x))$

- Dimensions to worry about: $x \in \mathbb{R}^p$, number of samples n

Acceleration #2: Cut-off complexity per iteration

- Stochastic gradient descent (SGD)

$$x_{t+1} = x_t - \eta_t \nabla f(x_t) \quad \longrightarrow \quad x_{t+1} = x_t - \eta_t \nabla f_{i_t}(x_t), \quad i_t \in [n]$$

where per iteration we select randomly $i_t \in [n]$.

- “Why do we want to do this?”

What is the complexity of
computing full gradient?

$$O(np)$$


(Assume least-squares objective)

What is the complexity of
computing a single gradient?

$$O(p)$$

- When is $n \gg p$? Big-data regime!
 - There is redundancy in data
 - Far from the optimal, exact gradients might have small returns

Acceleration #2: Cut-off complexity per iteration



Acceleration #2: Cut-off complexity per iteration

– Some notes on SGD (before we proceed)

1. It is a stochastic process that depends on a random sequence $i_t \in [n]$

(This means you will see some probability involved in theory)

2. While $-\nabla f(x_t)$ is a descent direction, $-\nabla f_{i_t}(x_t)$ might not be

(This means that some tools from deterministic optimization do not hold here)

3. The above lead to the intuition that if we have a descent direction in **expectation**, we probably will perform just fine


(This means that we will work with expectations w.r.t. the random sequence)

Guarantees of SGD

Whiteboard

Guarantees of SGD

Pic. from: “Optimization Methods for Large-Scale Machine Learning”



- Start with large step size; decrease it when SGD “stalls”

Guarantees of SGD

Whiteboard

Intuition behind preference for SGD


- Overall, for strongly convex and smooth functions

	iteration complexity	per-iteration cost	total comput. cost
batch GD	$\log \frac{1}{\varepsilon}$	n	$n \log \frac{1}{\varepsilon}$
SGD	$\frac{1}{\varepsilon}$	1	$\frac{1}{\varepsilon}$

- The real comparison is between $n \log \frac{1}{\varepsilon}$? $\frac{1}{\varepsilon}$
- In the big data regime, n can be huge!
- Gradient descent uses full dataset per iteration; there might be **redundancies**
- It actually works great in practice!

Intuition behind preference for SGD

Pic. from: “Optimization Methods for Large-Scale Machine Learning”



Variants of SGD

- Mini-batch SGD: instead of picking one sample, pick multiple

$$x_{t+1} = x_t - \eta_t \nabla f_{\mathcal{I}_t}(x_t) = x_t - \eta_t \cdot \sum_{j=1}^{|\mathcal{I}_t|} \nabla f_j(x_t)$$

- Still less time than computing the full gradient
 - Converges to a smaller ball around optimum: trade-off
-
- SGD with importance sampling: select “carefully” the next sample
 - Select $i_t \in [n]$ according to distribution $p \in [0, 1]^n$, $\sum_i p_i = 1$
 - Main question: can we compute a good probability distribution without too much effort?

Variants of SGD

- Stochastic variance-reduced gradient (SVRG)

$$x_{t+1} = x_t - \eta_t \left(\nabla f_{i_t}(x_t) - \underbrace{(\nabla f_{i_t}(\tilde{x}_q) - \nabla f(\tilde{x}_q))}_{\text{Bias in gradient estimate Correction term}} \right)$$

Bias in gradient estimate
Correction term

- Observe that: $\mathbb{E}[\nabla f_{i_t}(\cdot)] = \nabla f(\cdot)$; then

$$\mathbb{E} [\nabla f_{i_t}(x_t) - \nabla f_{i_t}(\tilde{x}_q) + \nabla f(\tilde{x}_q)] = \nabla f(x_t)$$

Unbiased estimator!
We expect smaller variance

- Theoretical guarantees:

$$\mathbb{E} [f(x_{t+1}) - f(x^*)] \leq \rho \cdot \mathbb{E} [f(x_t) - f(x^*)] , \quad \rho = O \left(\frac{1}{1-2\eta L} \cdot \left(\frac{1}{m\eta} + 2L\eta \right) \right) < 1$$

- Main drawback: Full gradient, but overall complexity $O \left((n + \kappa) \log \frac{1}{\varepsilon} \right)$

Performance of SGD


Demo

Why SGD is so important in machine learning?

(or some of the reasons)

- We ask batch optimization methods to find a model that minimizes an objective other than what they try to optimize
- SGD explores better the landscape

(more to come in future lectures)



where batch
methods
thrive

Pic. from: "Optimization Methods
for Large-Scale Machine Learning"

Acceleration #3: Coordinate descent methods

- Stochastic gradient descent (SGD) selects mini batches of training data; what if we subselect variables to update per iteration?

$$x_{t+1} = x_t - \eta_t \nabla f(x_t) \quad \longrightarrow \quad (x_{t+1})_{i_t} = (x_t)_{i_t} - \eta_t \nabla_{i_t} f(x_t),$$

where per iteration we select randomly $i_t \in [p]$.

- “Why do we want to do this?”

What is the complexity of
computing full gradient?

$$O(np)$$

(Assume least-squares objective)

What is the complexity of
computing a gradient for a
single variable?

$$O(n)$$

- When is $n \ll p$? High-dimensional case
 - There is not enough data
 - We will see that it provides solutions to distributed systems